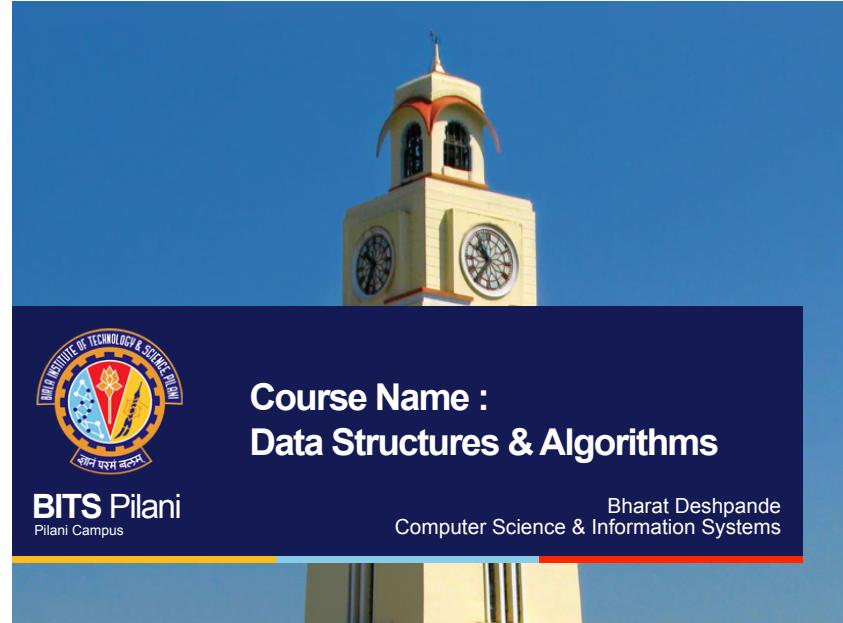


<p>Birla Institute of Technology & Science, Pilani Working and Learning Programmes Division First Semester 2014-2015</p> <p>Comprehensive Examination (ECA 3 Make-II)</p>
<p>Course No. : SS 2G519 Course Title : DATA STRUCTURES & ALGORITHMS DESIGN Nature of Exam : Open Book Weightage : 50% Duration : 3 Hours Date of Exam : 22/11/2014 (FN) Note:</p>
<p>1. Please follow all the <i>Instructions to Candidates</i> given on the cover page of the answer book. 2. All parts of a question should be answered consecutively. Each answer should start from a fresh page. 3. Assumptions made, if any, should be stated clearly at the beginning of your answer.</p>
<p>Q.1. Insert the elements $A = \{33, 22, 11, 43, 53, 63, 73, 96, 95, 94, 83, 82\}$ sequentially into an empty AVL tree. Write all intermediate steps clearly. [10]</p>
<p>Q.2. $E(G) = \{v1, v2, v3, v4, v2\bar{v}3, v3\bar{v}4, v3\bar{v}5, v5\bar{v}6, v5\bar{v}7, v6\bar{v}7\}$. G is an undirected graph with vertex set $V(G) = \{v_1, v_2, v_3, v_4, v_5, v_6, v_7\}$ and edge Set $E(G)$. Write all intermediate steps clearly and write the stack, active vertex at every iteration. Also make a table of depth first search index of the vertices. [3 + 3 + 6 + 10 = 22]</p>
<p>Q.3. Find the minimum spanning tree of the weighted graph. Starting at vertex v_1 using Prim's algorithm and find the weight of the minimum spanning tree. Write all intermediate steps clearly. Vertex set $V(G) = \{a, b, c, d, e, f, g, h\}$ Edge Set $E(G) = \{ab, ac, bc, bd, cd, ce, de, fg, eg, dg, gh, fh\}$ Weights = {ab = 2, ac = 1, bc = 1, bd = 3, cd = 5, ce = 7, de = 4, fg = 3, eg = 8, dg = 9, gh = 10, fh = 12}</p>
<p>Q.4. G is an undirected graph with vertex set $V(G) = \{a, b, c, d, e, f, g, h\}$ and edge Set $E(G) = \{ab, ac, bc, bd, ce, df, ef, fg, dg, gh, fh\}$. Is the graph an Euler graph? Justify your answer. (a) If the graph is an Euler graph, find the Euler circuit using Fleury's algorithm. (b) Find a hamiltonian circuit in the graph. (c) Find a hamiltonian circuit in the graph.</p>



Bina Institute of Technology & Sciences, Pilani															
Work-Integrated Learning Programmes Division															
First Semester 2014-2015															
Comprehensive Examination															
(EC-3 Regular)															
SS ZG519 DATA STRUCTURES & ALGORITHMS DESIGN															
<table border="1"> <tr> <td>No. of Pages = 1</td> </tr> <tr> <td>No. of Questions = 5</td> </tr> </table>		No. of Pages = 1	No. of Questions = 5												
No. of Pages = 1															
No. of Questions = 5															
Course No.	SS ZG519														
Course Title	DATA STRUCTURES & ALGORITHMS DESIGN														
Nature of Exam	Open Book														
Weightage	50%														
Duration	3 Hours														
Date of Exam	08/1/2014 (FN)														
Note:															
1. Please follow all the Instructions to Candidates given on the cover page of the answer book.															
2. All parts of a question should be answered consecutively. Each answer should start from a fresh page.															
3. Assumptions made if any, should be stated clearly at the beginning of your answer.															
Q.1. Illustrate merge sort to sort the array A = {54, 65, 9, 72, 81, 16, 23, 37, 49} in ascending order. Write all intermediate steps clearly. [8]															
Q.2. Build a max heap out of the array A = {8, 20, 9, 4, 15, 10, 7, 22, 3, 12}; Write all intermediate steps clearly. [10]															
Q.3. Construct a Huffman tree and a Huffman code for the following data. Clearly show the intermediate steps in the construction of the Huffman tree. [10]															
<table border="1"> <thead> <tr> <th>Character</th> <th>A</th> <th>B</th> <th>C</th> <th>D</th> <th>E</th> <th>F</th> </tr> </thead> <tbody> <tr> <td>Frequency</td> <td>0.28</td> <td>0.21</td> <td>0.67</td> <td>0.4</td> <td>0.22</td> <td>0.27</td> </tr> </tbody> </table>		Character	A	B	C	D	E	F	Frequency	0.28	0.21	0.67	0.4	0.22	0.27
Character	A	B	C	D	E	F									
Frequency	0.28	0.21	0.67	0.4	0.22	0.27									
Q.4. G is a an undirected graph with vertex set V(G) = { v1, v2, v3, v4, v5, v6, v7 } and edge Set E(G) = { v1v2, v1v4, v1v3, v2v4, v2v3, v3v4, v3v5, v4v5, v5v6, v5v7, v6v7 }. [10]															
Draw the graph with the given vertex set and edge set.															
<p>(a) Apply breadth first search algorithm to find the breadth first search tree of the graph G. Write all intermediate steps clearly and write the queue at every iteration.</p>															
<p>(b) Find the minimum spanning tree of the weighted graph G using Kruskal's algorithm and find the weight of the minimum spanning tree. Write all intermediate steps clearly.</p>															
Q.5. Find the minimum spanning tree of the weighted graph G using Kruskal's algorithm and find the weight of the minimum spanning tree. Write all intermediate steps clearly. [10]															
Vertices Set V(G) = {a, b, c, d, e, f, g}															
Edge Set E(G) = {ab, ac, bc, bd, cd, cf, de, ef, fg, eg, dg}															
Weights = { ab = 2, ac = 4, bc = 1, bd = 3, cd = 5, cf = 4, de = 7, ef = 6, fg = 3, eg = 8, dg = 9 }															



What is a program?

- **Algorithm**

An algorithm is a step-by-step procedure for solving a problem in a finite amount of time.

- **Data Structures**

Is a systematic way of organizing and accessing data, so that data can be used efficiently.

Algorithms + Data Structures = Program

Algorithmic problem



For eg: Sorting of integers

Input Instance : 8,4,5,2,10
Output Instance as a permutation of input : 2,4,5,8,10

What is good algorithm?



- Resources Used
 - Running time
 - Space used
- Resource Usage
 - Measured proportional to (input) size

Algorithmic Solution



- Algorithm describes actions on the input instance.
- Infinitely many correct algorithm for the same problem.

Infinite number of input instances satisfying the specification.

Two key points: Repeatable argument & Correctness

Measuring the running time



- Write a program implementing the algorithm.
- Run the program with inputs of varying size and composition.
- Use a method like `System.currentTimeMillis()` to get an accurate measure of the actual running time.

Limitations of experimental studies



- Implementation is a must.
- Execution is possible on limited set of inputs.
- If we need to compare two algorithms we need to use the same environment (like hardware, software etc)

Analytical model to analyze algorithm



- Algorithm should be analyzed by using general methodology.
- This approach uses:
 - High level description of the algorithm.
 - Takes into account all possible inputs.
 - Allows one to evaluate the efficiency of any algorithm in a way that is independent of the hardware and the software environment.

7

8
BITS Pilani, Pilani Campus

Pseudo-code



- A mixture of natural language and high level programming concepts that describes the main ideas behind a generic implementation of a data structure and algorithms.

Algorithm *arrayMax(A, n)*

Input: An array *A* of *n* integers

Output: The maximum element of *A*

```
currentMax ← A[0]
for i ← 1 to n - 1 do
    if A[i] > currentMax then currentMax ← A[i]
return currentMax
```

Pseudo-code



- Is structured than usual prose but less formal than a programming language.
- Expressions
 - Use standard mathematical symbols to describe numeric and Boolean expressions.
 - Uses \leftarrow for assignment.
 - Use = for the equality relationship.
- Method declaration
 - Algorithm name(param1,param2...)

9

10
BITS Pilani, Pilani Campus



Assumptions



Individual statement considered as “unit” time

- Not applicable for function calls and loops

Individual variable considered as “unit” storage

Often referred to as “algorithmic complexity”

11

BITS Pilani, Pilani Campus

Complexity Example [2]



Example 2 (a and N are input)

```
j = 0;
while (j < N) do
    a[j] = a[j] * a[j];
    b[j] = a[j] + j;
    j = j + 1;
endwhile;
// 3N + 1 units of time and N+1 units of storage
// time units prop. to N and storage prop. to N
```

13

BITS Pilani, Pilani Campus

Complexity Example [1]

Example 1 (Y and Z are input)

X = Y * Z;

X = Y * X + Z;

// 2 units of time and 1 unit of storage

// Constant Unit of time and Constant Unit of storage

12

BITS Pilani, Pilani Campus

Complexity Example [3]



Example 3 (a and N are input)

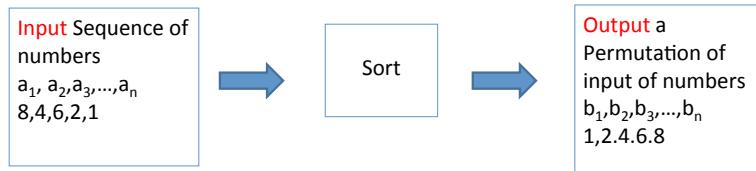
```
j = 0;
while (j < N) do
    k = 0;
    while (k < N) do
        a[k] = a[j] + a[k];
        k = k + 1;
    endwhile;
    b[j] = a[j] + j;
    j = j + 1;
endwhile;
//??? units of time and ??? units of storage
// time prop. to N2 and storage prop. to N
```

14

BITS Pilani, Pilani Campus



Example of sorting



Correctness(Requirement for the output)

For any input algorithm halts with the output:

- $b_1 < b_2 < b_3 < \dots < b_n$
- $b_1, b_2, b_3, \dots, b_n$ is a permutation of $a_1, a_2, a_3, \dots, a_n$

Running time of algorithm depends on

- Number of elements n.
- How (partially)sorted they are.

15

BITS Pilani, Pilani Campus

Order Notation

- **Purpose**
 - Capture proportionality
 - Machine independent measurement
 - Asymptotic growth
(i.e. large values of input size N)

16

BITS Pilani, Pilani Campus

Motivation for Order Notation



Examples

- $100 * \log_2 N < N$ for $N > 1000$
- $70 * N + 3000 < N^2$ for $N > 100$
- $10^5 * N^2 + 10^6 * N < 2^N$ for $N > 26$

Asymptotic Analysis



- Goal: To simplify analysis of running time of algorithm .eg $3n^2 = n^2$.
- Capturing the essence: how the running time of the algorithm increases with the size of the input in the limit.

17

BITS Pilani, Pilani Campus

18

BITS Pilani, Pilani Campus



Asymptotic Notation



- The big O notation

Definition

Let f and g be functions from the set of integers to the set of real numbers. We say that $f(x)$ is in $O(g(x))$ if there are constants $C >$ and k such that $|f(x)| \leq C |g(x)|$, whenever $x \geq k$.

- This is read as $f(x)$ is **big-oh** of $g(x)$

Note: Pair of C and k is never unique.

19

BITS Pilani, Pilani Campus

Linear Search



```
function search(X, A, N)
j = 0;
while (j < N)
    if (A[j] == X) return j;
    j++;
endwhile;
return "Not-found";
```

21

BITS Pilani, Pilani Campus

Order Notation

Examples

$$g(n) = 17 * N + 5$$

$$\lim_{n \rightarrow \infty} g(n) / f(n) = c \\ \lim_{n \rightarrow \infty} (17 * N + 5) / N = 17. \text{ The asymptotic complexity is } O(N)$$

$$g(n) = 5 * N^3 + 10 * N^2 + 3$$

$$\lim_{n \rightarrow \infty} (5 * N^3 + 10 * N^2 + 3) / N^3 = 5. \text{ The asymptotic complexity is } O(N^3)$$

$$g(n) = C1 * N^k + C2 * N^{k-1} + \dots + Ck * N + C$$

$$\lim_{n \rightarrow \infty} (C1 * N^k + C2 * N^{k-1} + \dots + Ck * N + C) / N^k = C1. \\ \text{The asymptotic complexity is } O(N^k)$$

$$2^N + 4 * N^3 + 16 \text{ is } O(2^N)$$

$$5 * N * \log(N) + 3 * N \text{ is } O(N * \log(N))$$

$$1789 \text{ is } O(1)$$

20

BITS Pilani, Pilani Campus

Linear Search - Complexity



Time Complexity

“if” statement introduces possibilities

- Best-case: $O(1)$
- Worst case: $O(N)$
- Average case: ???

22

BITS Pilani, Pilani Campus

Binary Search Algorithm



Assume: Sorted Sequence of numbers
 low = 1; high = N;
 while (low <= high) do
 mid = (low + high) /2;
 if (A[mid] == x) return x;
 else if (A[mid] < x) low = mid +1;
 else high = mid - 1;
 endwhile;
 return Not-Found;

23

BITS Pilani, Pilani Campus

Binary Search - Complexity



- Worst case:
 - K steps such that $2^K = N$
 - i.e. $\log_2 N$ steps is $O(\log(N))$

25

BITS Pilani, Pilani Campus

Binary Search - Complexity

- Best Case
 - $O(1)$
- Worst case:
 - Loop executes until **low <= high**
 - Size halved in each iteration
 - $N, N/2, N/4, \dots 1$
 - How many steps ?

24

BITS Pilani, Pilani Campus



Course Name :
Data Structures & Algorithms

BITS Pilani
Pilani Campus

Bharat Deshpande
Computer Science & Information Systems

Assumptions

Individual statement considered as “unit” time

- Not applicable for function calls and loops

Individual variable considered as “unit” storage

Often referred to as “algorithmic complexity”

Complexity Example [1]

Example 1 (Y and Z are input)

```
X = Y * Z;  
X = Y * X + Z;  
// 2 units of time and 1 unit of storage  
// Constant Unit of time and Constant Unit of  
storage
```

2

3

Complexity Example [2]

Example 2 (a and N are input)

```
j = 0;  
while (j < N) do  
    a[j] = a[j] * a[j];  
    b[j] = a[j] + j;  
    j = j + 1;  
endwhile;  
// 3N + 1 units of time and N+1 units of storage  
// time units prop. to N and storage prop. to N
```

Complexity Example [3]

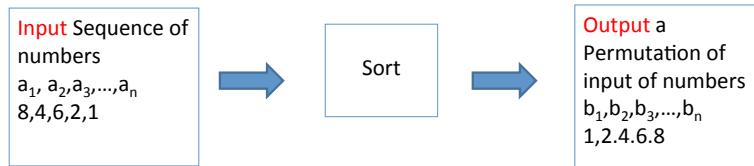
Example 3 (a and N are input)

```
j = 0;  
while (j < N) do  
    k = 0;  
    while (k < N) do  
        a[k] = a[j] + a[k];  
        k = k + 1;  
    endwhile;  
    b[j] = a[j] + j;  
    j = j + 1;  
endwhile;  
//??? units of time and ??? units of storage  
// time prop. to N2 and storage prop. to N
```

4

5

Example of sorting



Correctness(Requirement for the output)

For any input algorithm halts with the output:

- $b_1 < b_2 < b_3 < \dots < b_n$
- $b_1, b_2, b_3, \dots, b_n$ is a permutation of $a_1, a_2, a_3, \dots, a_n$

Running time of algorithm depends on

- Number of elements n.
- How (partially)sorted they are.

Order Notation

- Purpose
 - Capture proportionality
 - Machine independent measurement
 - Asymptotic growth (i.e. large values of input size N)

6

7

Motivation for Order Notation

Examples

- $100 * \log_2 N < N$ for $N > 1000$
- $70 * N + 3000 < N^2$ for $N > 100$
- $10^5 * N^2 + 10^6 * N < 2^N$ for $N > 26$

Asymptotic Analysis

- Goal: To simplify analysis of running time of algorithm .eg $3n^2 = n^2$.
- Capturing the essence: how the running time of the algorithm increases with the size of the input in the limit.

8

9

Asymptotic Notation

- The big O notation

Definition

Let f and g be functions from the set of integers to the set of real numbers. We say that $f(x)$ is in $O(g(x))$ if there are constants $C >$ and k such that $|f(x)| \leq C|g(x)|$, whenever $x \geq k$.

- This is read as $f(x)$ is **big-oh** of $g(x)$

Note: Pair of C and k is never unique.

10

Order Notation

Examples

$$g(n) = 17n + 5$$

$$\lim_{n \rightarrow \infty} g(n) / f(n) = c \\ \lim_{n \rightarrow \infty} (17n + 5)/n = 17. \text{ The asymptotic complexity is } O(N)$$

$$g(n) = 5N^3 + 10N^2 + 3$$

$$\lim_{n \rightarrow \infty} (5N^3 + 10N^2 + 3) / N^3 = 5. \text{ The asymptotic complexity is } O(N^3)$$

$$g(n) = C_1N^k + C_2N^{k-1} + \dots + C_kN + C$$

$$\lim_{n \rightarrow \infty} (C_1N^k + C_2N^{k-1} + \dots + C_kN + C) / N^k = C_1. \\ \text{The asymptotic complexity is } O(N^k)$$

$$2^N + 4N^3 + 16 \text{ is } O(2^N)$$

$$5N \log(N) + 3N \text{ is } O(N \log(N))$$

$$1789 \text{ is } O(1)$$

11

Linear Search

```
function search(X, A, N)
j = 0;
while (j < N)
    if (A[j] == X) return j;
    j++;
endwhile;
return "Not-found";
```

12

Linear Search - Complexity

Time Complexity

“if” statement introduces possibilities

- Best-case: $O(1)$
- Worst case: $O(N)$
- Average case: ???

13

Binary Search Algorithm

Assume: Sorted Sequence of numbers
low = 1; high = N;
while (low <= high) do
 mid = (low + high) /2;
 if (A[mid] == x) return x;
 else if (A[mid] < x) low = mid +1;
 else high = mid - 1;
endwhile;
return Not-Found;

14

Binary Search - Complexity

- Best Case
 - $O(1)$
- Worst case:
 - Loop executes until $low \leq high$
 - Size halved in each iteration
 - $N, N/2, N/4, \dots 1$
 - How many steps ?

15



Binary Search - Complexity

- Worst case:
 - K steps such that $2^K = N$
i.e. $\log_2 N$ steps is $O(\log(N))$

16

Abstract Data Type

- Abstract Data Type
- In computer science, an **abstract data type** (**ADT**) is a mathematical model for a certain class of data structures that have similar behavior.

17

BITS Pilani, Pilani Campus



Abstract Data Types (ADTs)



- A method for achieving abstraction for data structures and algorithms
- ADT = model + operations
- Describes what each operation does, but not how it does it
- An ADT is independent of its implementation

Abstract Data Types

- Typical operations on data
 - Add data to a data collection
 - Remove data from a data collection
 - Ask questions about the data in a data collection

BITS Pilani, Pilani Campus

Abstract Data Types



- Data abstraction
 - Asks you to think *what* you can do to a collection of data independently of *how* you do it
 - Allows you to develop each data structure in relative isolation from the rest of the solution
 - A natural extension of procedural abstraction

Examples



- Simple ADTs
 - Stack
 - Queue
 - Vector
 - Lists
 - Sequences
 - Iterators
- All these are called **Linear Data Structures**

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

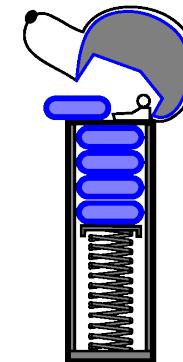
Stacks

- A **stack** is a container of objects that are inserted and removed according to the last-in-first-out (**LIFO**) principle.
- Objects can be inserted at any time, but only the last (the most-recently inserted) object can be removed.
- Inserting an item is known as “**pushing**” onto the stack. “**Popping**” off the stack is synonymous with removing an item.



Stacks

- A coin dispenser as an analogy:



BITS Pilani, Pilani Campus

Stacks: An Array Implementation



- Create a stack using an array by specifying a maximum size N for our stack.
- The stack consists of an N -element array S and an integer variable t , the index of the top element in array S .



- Array indices start at 0, so we initialize t to -1

BITS Pilani, Pilani Campus



Stacks: An Array Implementation

- **Pseudo code**

```
Algorithm size()
if size() == N then
    return Error
    t = t + 1
    S[t] = o
```

```
Algorithm isEmpty()
return (t < 0)
```

```
Algorithm top()
if isEmpty() then
    return Error
    return S[t]
```

```
Algorithm push(o)
if size() == N then
    return Error
    t = t + 1
    S[t] = o

Algorithm pop()
if isEmpty() then
    return Error
    t = t - 1
    return S[t+1]
```



BITS Pilani, Pilani Campus

Stacks: An Array Implementation



The array implementation is simple and efficient (methods performed in $O(1)$).

Disadvantage

There is an upper bound, N , on the size of the stack.

The arbitrary value N may be too small for a given application OR
a waste of memory.

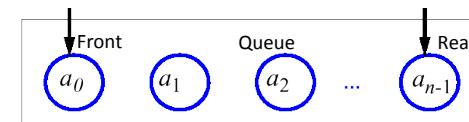
Queues



- The **queue** supports three fundamental methods:
 - **New():ADT** – Creates an empty queue
 - **Enqueue(S:ADT, o:element):ADT** - Inserts object o at the rear of the queue
 - **Dequeue(S:ADT):ADT** - Removes the object from the front of the queue; an error occurs if the queue is empty
 - **Front(S:ADT):element** - Returns, but does not remove, the front element; an error occurs if the queue is empty

Queues

- A queue differs from a stack in that its insertion and removal routines follows the **first-in-first-out** (FIFO) principle.
- Elements may be inserted at any time, but only the element which has been in the queue the longest may be removed.
- Elements are inserted at the **rear** (enqueued) and removed from the **front** (dequeued)



Queues: An Array Implementation

- Create a queue using an array in a circular fashion
- A maximum size N is specified.
- The queue consists of an N -element array Q and two integer variables:
 - f , index of the front element (head – for dequeue)
 - r , index of the element after the rear one (tail – for enqueue)
 - Initially, $f=r=0$ and the queue is empty if $f=r$



Queues



Disadvantage

Repeatedly enqueue and dequeue a single element N times.

Finally, $f=r=N$.

- No more elements can be added to the queue,

though there is space in the queue.

Solution

Let f and r wraparound the end of queue.

30
BITS Pilani, Pilani Campus

Queues: An Array Implementation



Pseudo code

```
Algorithm size()
return  $(N-f+r) \bmod N$ 

Algorithm isEmpty()
return  $(f=r)$ 

Algorithm front()
if isEmpty() then
    return Error
return  $Q[f]$ 
```

```
Algorithm dequeue()
if isEmpty() then
    return Error
 $Q[f]=\text{null}$ 
 $f=(f+1) \bmod N$ 

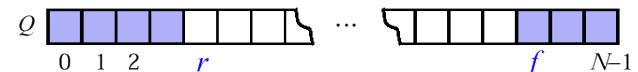
Algorithm enqueue(o)
if size =  $N - 1$  then
    return Error
 $Q[r]=o$ 
 $r=(r + 1) \bmod N$ 
```

BITS Pilani, Pilani Campus

Queues: An Array Implementation



“wrapped around” configuration



- Each time r or f is incremented, compute this increment as $(r+1) \bmod N$ or $(f+1) \bmod N$

BITS Pilani, Pilani Campus

Arrays: pluses and minuses

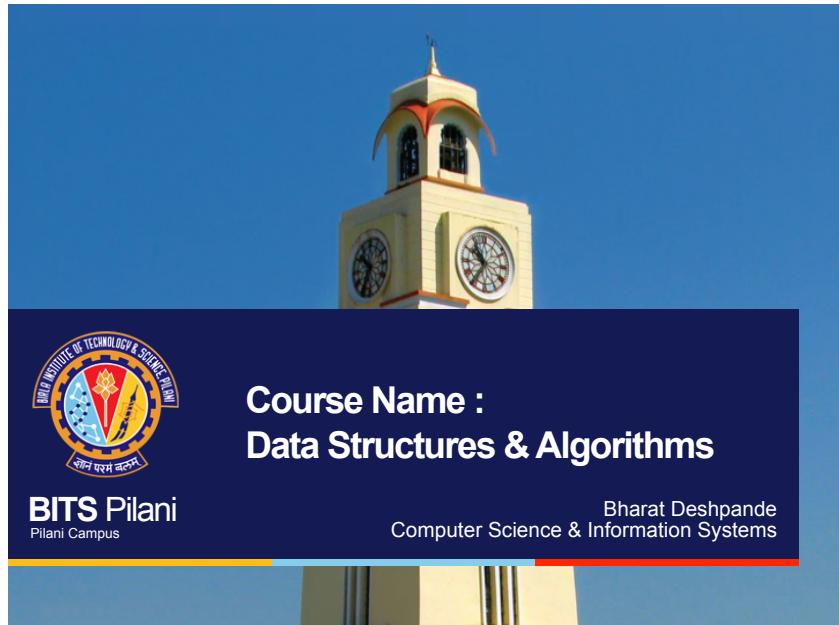


+ Fast element access.

-- Impossible to resize.

- Many applications require resizing!
- Required size not always immediately available.

33
BITS Pilani, Pilani Campus



The slide features the BITSPilani logo on the left, which includes a circular emblem with a torch and the text 'BITS PILANI' and 'राजस्थान विश्वविद्यालय'. To the right of the logo is a yellow clock tower against a clear blue sky. Below the logo, the course information is displayed:

Course Name :
Data Structures & Algorithms

BITS Pilani
 Pilani Campus

Bharat Deshpande
 Computer Science & Information Systems

Order Notation

- Purpose
 - Capture proportionality
 - Machine independent measurement
 - Asymptotic growth
(i.e. large values of input size N)

2

Asymptotic Analysis

- Goal: To simplify analysis of running time of algorithm .eg $3n^2=n^2$.
- Capturing the essence: how the running time of the algorithm increases with the size of the input in the limit.

Asymptotic Notation

- The big O notation
Definition
 Let f and g be functions from the set of integers to the set of real numbers. We say that $f(x)$ is in $O(g(x))$ if there are constants $C >$ and k such that $|f(x)| \leq C |g(x)|$, whenever $x \geq k$.
 • This is read as $f(x)$ is **big-oh** of $g(x)$
Note: Pair of C and k is never unique.

3

4

Order Notation

Examples

$$g(n) = 17N + 5$$

$$\lim_{n \rightarrow \infty} g(n) / f(n) = c$$

$(17N + 5)/N = 17$. The asymptotic complexity is $O(N)$

$$g(n) = 5N^3 + 10N^2 + 3$$

$$\lim_{n \rightarrow \infty} (5N^3 + 10N^2 + 3) / N^3 = 5$$
. The asymptotic complexity is $O(N^3)$

$$g(n) = C_1N^k + C_2N^{k-1} + \dots + C_kN + C$$

$$\lim_{n \rightarrow \infty} (C_1N^k + C_2N^{k-1} + \dots + C_kN + C) / N^k = C_1$$

The asymptotic complexity is $O(N^k)$

$$2^N + 4N^3 + 16 \text{ is } O(2^N)$$

$$5N^2\log(N) + 3N \text{ is } O(N^2\log(N))$$

$$1789 \text{ is } O(1)$$

Linear Search

function search(X, A, N)

j = 0;

while (j < N)

if (A[j] == X) return j;

j++;

endwhile;

return "Not-found";

5

6

Linear Search - Complexity

Time Complexity

“if” statement introduces possibilities

- Best-case: $O(1)$
- Worst case: $O(N)$
- Average case: ???

Binary Search Algorithm

Assume: Sorted Sequence of numbers

low = 1; high = N;

while (low <= high) do

mid = (low + high) /2;

if (A[mid] == x) return x;

else if (A[mid] < x) low = mid + 1;

else high = mid – 1;

endwhile;

return Not-Found;

7

8

Binary Search - Complexity

- Best Case
 - $O(1)$
- Worst case:
 - Loop executes until **low <= high**
 - Size halved in each iteration
 - $N, N/2, N/4, \dots 1$
 - How many steps ?

Binary Search - Complexity

- Worst case:
 - K steps such that $2^K = N$
 - i.e. $\log_2 N$ steps is $O(\log(N))$

9

10

Abstract Data Type



- Abstract Data Type

In computer science, an **abstract data type (ADT)** is a mathematical model for a certain class of data structures that have similar behavior.

Abstract Data Types (ADTs)



- A method for achieving abstraction for data structures and algorithms
- ADT = model + operations
- Describes what each operation does, but not how it does it
- An ADT is independent of its implementation

11

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus



Abstract Data Types



- Typical operations on data
 - Add data to a data collection
 - Remove data from a data collection
 - Ask questions about the data in a data collection

Abstract Data Types

- Data abstraction
 - Asks you to think *what* you can do to a collection of data independently of *how* you do it
 - Allows you to develop each data structure in relative isolation from the rest of the solution
 - A natural extension of procedural abstraction

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

Examples



- Simple ADTs
 - Stack
 - Queue
 - Vector
 - Lists
 - Sequences
 - Iterators
- All these are called **Linear Data Structures**

BITS Pilani, Pilani Campus



Stacks

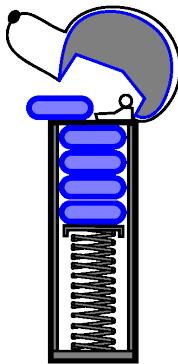
- A **stack** is a container of objects that are inserted and removed according to the last-in-first-out (**LIFO**) principle.
- Objects can be inserted at any time, but only the last (the most-recently inserted) object can be removed.
- Inserting an item is known as “**pushing**” onto the stack. “**Popping**” off the stack is synonymous with removing an item.

BITS Pilani, Pilani Campus

Stacks



- A coin dispenser as an analogy:



BITS Pilani, Pilani Campus

Stacks: An Array Implementation



Pseudo code

Algorithm **size()**

return $t+1$

Algorithm **isEmpty()**

return ($t < 0$)

Algorithm **top()**

if isEmpty() **then**

return Error

return $S[t]$

```
Algorithm push(o)
    if size() == N then
        return Error
     $t = t + 1$ 
     $S[t] = o$ 

Algorithm isEmpty()
    return ( $t < 0$ )

Algorithm top()
    if isEmpty() then
        return Error
     $t = t - 1$ 
    return  $S[t + 1]$ 
```



BITS Pilani, Pilani Campus

Stacks: An Array Implementation



- Create a stack using an array by specifying a maximum size N for our stack.
- The stack consists of an N -element array S and an integer variable t , the index of the top element in array S .



- Array indices start at 0, so we initialize t to -1

BITS Pilani, Pilani Campus

Stacks: An Array Implementation



The array implementation is simple and efficient (methods performed in $O(1)$).

Disadvantage

There is an upper bound, N , on the size of the stack.

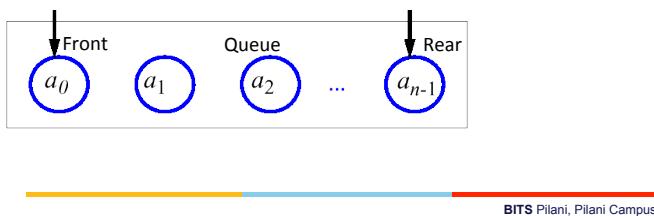
The arbitrary value N may be too small for a given application OR

a waste of memory.

BITS Pilani, Pilani Campus

Queues

- A queue differs from a stack in that its insertion and removal routines follows the **first-in-first-out** (FIFO) principle.
- Elements may be inserted at any time, but only the element which has been in the queue the longest may be removed.
- Elements are inserted at the **rear** (enqueued) and removed from the **front** (dequeued)



BITs Pilani, Pilani Campus

Queues: An Array Implementation

- Create a queue using an array in a circular fashion
- A maximum size N is specified.
- The queue consists of an N -element array Q and two integer variables:
 - f , index of the front element (head – for dequeue)
 - r , index of the element after the rear one (tail – for enqueue)
 - Initially, $f=r=0$ and the queue is empty if $f=r$



BITs Pilani, Pilani Campus

Queues

- The **queue** supports three fundamental methods:
 - New():ADT** – Creates an empty queue
 - Enqueue(S:ADT, o:element):ADT** - Inserts object o at the rear of the queue
 - Dequeue(S:ADT):ADT** - Removes the object from the front of the queue; an error occurs if the queue is empty
 - Front(S:ADT):element** - Returns, but does not remove, the front element; an error occurs if the queue is empty

BITs Pilani, Pilani Campus

Queues

Disadvantage

Repeatedly enqueue and dequeue a single element N times.

Finally, $f=r=N$.

- No more elements can be added to the queue,

though there is space in the queue.

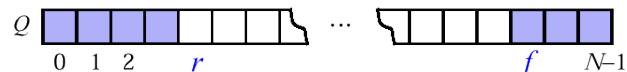
Solution

Let f and r wraparound the end of queue.



Queues: An Array Implementation

“wrapped around” configuration



- Each time r or f is incremented, compute this increment as $(r+1) \bmod N$ or $(f+1) \bmod N$

BITs Pilani, Pilani Campus

BITs Pilani, Pilani Campus

Arrays: pluses and minuses



- + Fast element access.
- Impossible to resize.
- Many applications require resizing!
- Required size not always immediately available.

27
BITs Pilani, Pilani Campus

Queues: An Array Implementation

• Pseudo code

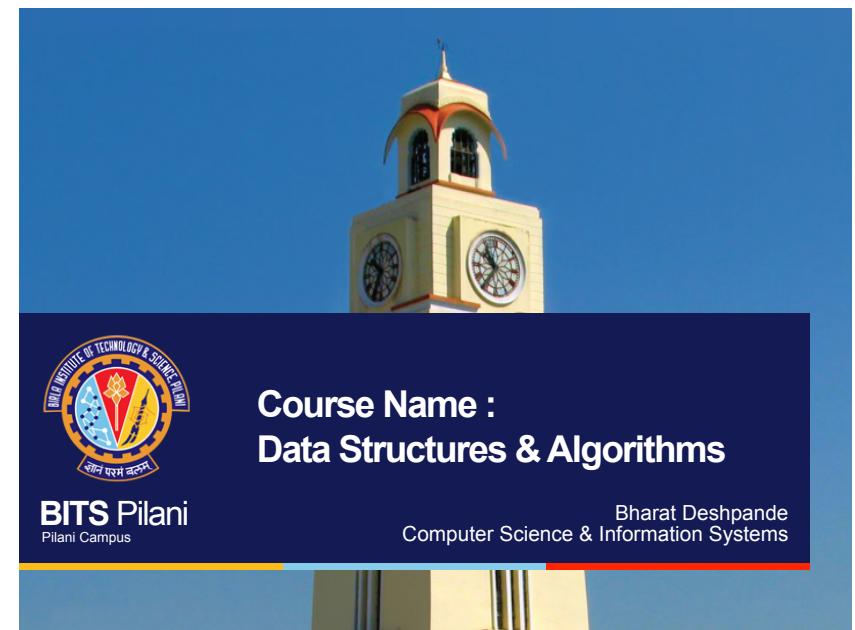
```
Algorithm size()
  return  $(N-f+r) \bmod N$ 
```

```
Algorithm isEmpty()
  return  $(f=r)$ 
```

```
Algorithm front()
  if isEmpty() then
    return Error
  return  $Q[f]$ 
```

```
Algorithm dequeue()
  if isEmpty() then
    return Error
   $Q[f]=\text{null}$ 
   $f=(f+1) \bmod N$ 

Algorithm enqueue(o)
  if size =  $N - 1$  then
    return Error
   $Q[r]=o$ 
   $r=(r + 1) \bmod N$ 
```



Examples



- **Simple ADTs**

- Stack
- Queue
- Vector
- Lists
- Sequences
- Iterators

All these are called **Linear Data Structures**

Stacks

- A **stack** is a container of objects that are inserted and removed according to the last-in-first-out (**LIFO**) principle.
- Objects can be inserted at any time, but only the last (the most-recently inserted) object can be removed.
- Inserting an item is known as “**pushing**” onto the stack. “**Popping**” off the stack is synonymous with removing an item.

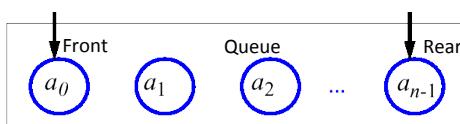
BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

Queues



- A queue differs from a stack in that its insertion and removal routines follows the **first-in-first-out** (FIFO) principle.
- Elements may be inserted at any time, but only the element which has been in the queue the longest may be removed.
- Elements are inserted at the **rear** (enqueued) and removed from the **front** (dequeued)



BITS Pilani, Pilani Campus

Queues

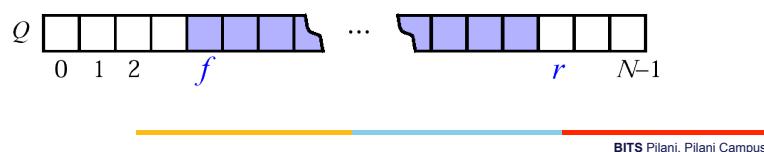


- The **queue** supports three fundamental methods:
 - **New():ADT** – Creates an empty queue
 - **Enqueue(S:ADT, o:element):ADT** - Inserts object o at the rear of the queue
 - **Dequeue(S:ADT):ADT** - Removes the object from the front of the queue; an error occurs if the queue is empty
 - **Front(S:ADT):element** - Returns, but does not remove, the front element; an error occurs if the queue is empty

BITS Pilani, Pilani Campus

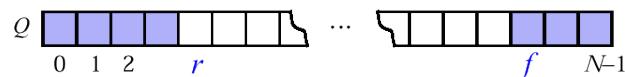
Queues: An Array Implementation

- Create a queue using an array
- A maximum size N is specified.
- The queue consists of an N -element array Q and two integer variables:
 - f , index of the front element (head – for dequeue)
 - r , index of the element after the rear one (tail – for enqueue)
 - Initially, $f=r=0$ and the queue is empty if $f=r$



Queues: An Array Implementation

“wrapped around” configuration



- Each time r or f is incremented, compute this increment as $(r+1)\bmod N$ or $(f+1)\bmod N$

Queues

Disadvantage

Repeatedly enqueue and dequeue a single element N times.

Finally, $f=r=N$.

- No more elements can be added to the queue,
- though there is space in the queue.

Solution

Let f and r wraparound the end of queue.

Queues: An Array Implementation

Pseudo code

```
Algorithm size()
  return  $(N-f+r) \bmod N$ 
```

```
Algorithm isEmpty()
  return  $(f=r)$ 
```

```
Algorithm front()
  if isEmpty() then
    return Error
  return  $Q[f]$ 
```

```
Algorithm dequeue()
  if isEmpty() then
    return Error
   $Q[f]=\text{null}$ 
   $f=(f+1) \bmod N$ 
```

```
Algorithm enqueue(o)
  if size =  $N - 1$  then
    return Error
   $Q[r]=o$ 
   $r=(r + 1) \bmod N$ 
```



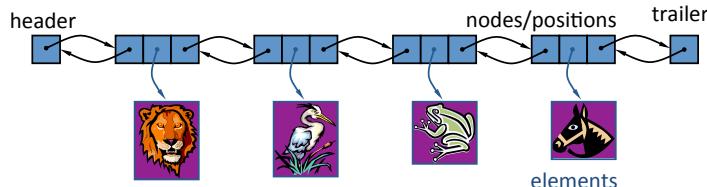
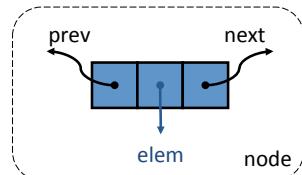
Arrays: pluses and minuses

- + Fast element access.
- Impossible to resize.
- Many applications require resizing!
- Required size not always immediately available.

10
BITS Pilani, Pilani Campus

Doubly Linked List

- A doubly linked list is often more convenient!
- Nodes store:
 - element
 - link to the previous node
 - link to the next node
- Special trailer and header nodes

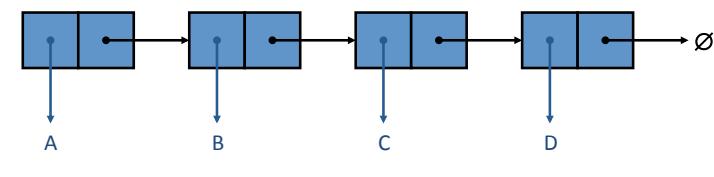
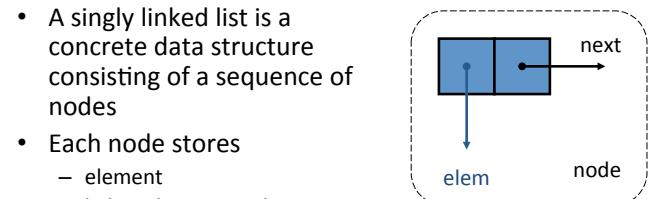


12
BITS Pilani, Pilani Campus



Singly Linked Lists

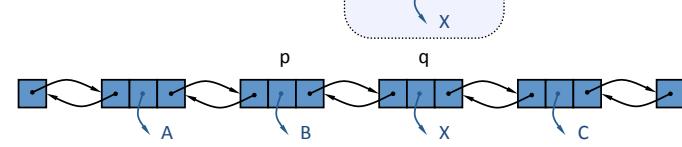
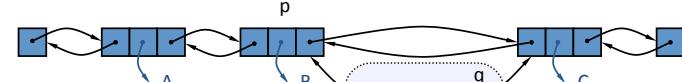
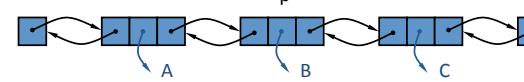
- A singly linked list is a concrete data structure consisting of a sequence of nodes
- Each node stores
 - element
 - link to the next node



11
BITS Pilani, Pilani Campus

Insertion

- We visualize operation `insertAfter(p, X)`, which returns position `q`



13
BITS Pilani, Pilani Campus



Insertion Algorithm

Algorithm insertAfter(p, e):

```
Create a new node  $v$ 
 $v.element \leftarrow e$ 
 $v.prev \leftarrow p$  {link  $v$  to its predecessor}
 $v.next \leftarrow p.next$  {link  $v$  to its successor}
 $(p.next).prev \leftarrow v$  {link  $p$ 's old successor to  $v$ }
 $p.next \leftarrow v$  {link  $p$  to its new successor,  $v$ }
return  $v$  {the position for the element  $e$ }
```

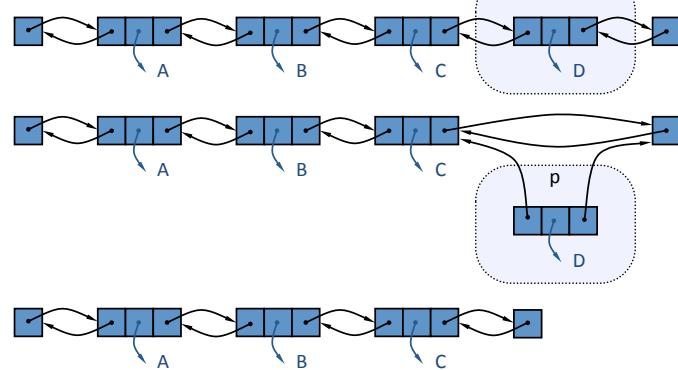
14

BITS Pilani, Pilani Campus



Deletion

- We visualize $\text{remove}(p)$, where $p == \text{last}()$



15

BITS Pilani, Pilani Campus

Deletion Algorithm



Algorithm remove(p):

```
 $t \leftarrow p.element$  {a temporary variable to hold the return value}
 $(p.prev).next \leftarrow p.next$  {linking out  $p$ }
 $(p.next).prev \leftarrow p.prev$ 
 $p.prev \leftarrow \text{null}$  {invalidating the position  $p$ }
 $p.next \leftarrow \text{null}$ 
return  $t$ 
```

16

BITS Pilani, Pilani Campus



Worst-case running time

- In a doubly linked list
 - + insertion at head or tail is in $O(1)$
 - + deletion at either end is in $O(1)$
 - element access is still in $O(n)$

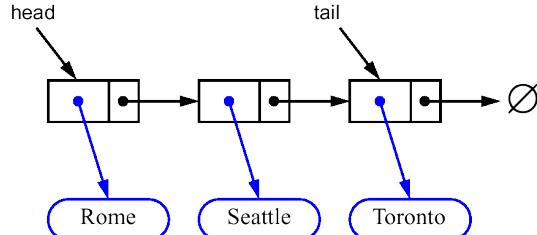
17

BITS Pilani, Pilani Campus

Stacks: Singly Linked List implementation



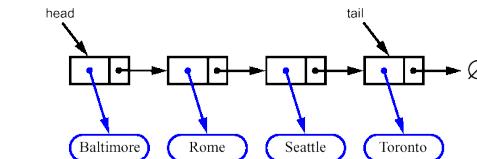
- Nodes (*data, pointer*) connected in a chain by links



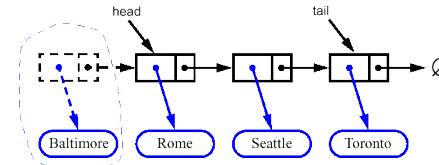
- the head or the tail of the list could serve as the top of the stack

BITs Pilani, Pilani Campus

Queues: Linked List Implementation



- Dequeue - advance head reference



BITs Pilani, Pilani Campus

Linear data structures



- Here are some of the data structures we have studied so far:
 - Arrays
 - Singly-linked lists and doubly-linked lists
 - Stacks and queues
 - These all have the property that their elements can be adequately displayed in a straight line
- How to obtain data structures for data that have nonlinear relationships**

BITs Pilani, Pilani Campus

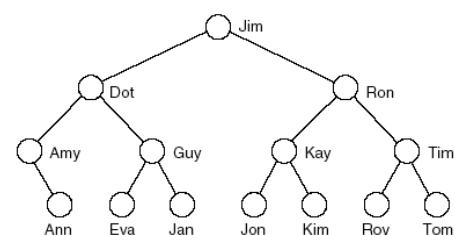
Tree



Tree represents hierarchy.

Examples of trees:

- Directory tree
- Family tree
- Company organization chart
- Table of contents



- structure resembles branches of a “tree”, hence the name.

BITs Pilani, Pilani Campus



Course Name :
Data Structures & Algorithms

BITSPilani
Pilani Campus

Bharat Deshpande
Computer Science & Information Systems



Abstract Data Type

- Abstract Data Type

In computer science, an **abstract data type** (**ADT**) is a mathematical model for a certain class of data structures that have similar behavior.

2

BITSPilani, Pilani Campus



Abstract Data Types (ADTs)



- A method for achieving abstraction for data structures and algorithms
- ADT = model + operations
- Describes what each operation does, but not how it does it
- An ADT is independent of its implementation

Abstract Data Types

- Typical operations on data
 - Add data to a data collection
 - Remove data from a data collection
 - Ask questions about the data in a data collection

BITSPilani, Pilani Campus

Abstract Data Types



- Data abstraction
 - Asks you to think *what* you can do to a collection of data independently of *how* you do it
 - Allows you to develop each data structure in relative isolation from the rest of the solution
 - A natural extension of procedural abstraction

Examples

- Simple ADTs

- Stack
- Queue
- Vector
- Lists
- Sequences
- Iterators

All these are called **Linear Data Structures**

BITs Pilani, Pilani Campus

Stacks

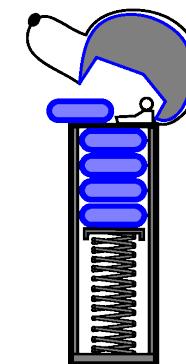


- A **stack** is a container of objects that are inserted and removed according to the last-in-first-out (**LIFO**) principle.
- Objects can be inserted at any time, but only the last (the most-recently inserted) object can be removed.
- Inserting an item is known as “**pushing**” onto the stack. “**Popping**” off the stack is synonymous with removing an item.

Stacks



- A coin dispenser as an analogy:



BITs Pilani, Pilani Campus

BITs Pilani, Pilani Campus



Stacks: An Array Implementation



- Create a stack using an array by specifying a maximum size N for our stack.
- The stack consists of an N -element array S and an integer variable t , the index of the top element in array S .



- Array indices start at 0, so we initialize t to -1

BITS Pilani, Pilani Campus

Stacks: An Array Implementation



The array implementation is simple and efficient (methods performed in $O(1)$).

Disadvantage

There is an upper bound, N , on the size of the stack.

The arbitrary value N may be too small for a given application OR
a waste of memory.

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

Stacks: An Array Implementation



Pseudo code

```
Algorithm size()  
    if size() == N then
```

```
        return Error  
    t = t + 1
```

```
    S[t] = o
```

```
Algorithm isEmpty()  
    return (t < 0)
```

```
Algorithm pop()  
    if isEmpty() then
```

```
        return Error  
    t = t - 1
```

```
    return S[t + 1]
```

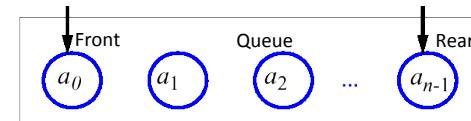


BITS Pilani, Pilani Campus

Queues



- A queue differs from a stack in that its insertion and removal routines follows the **first-in-first-out** (FIFO) principle.
- Elements may be inserted at any time, but only the element which has been in the queue the longest may be removed.
- Elements are inserted at the **rear** (enqueued) and removed from the **front** (dequeued)



BITS Pilani, Pilani Campus



Queues



- The **queue** supports three fundamental methods:
 - New():ADT** – Creates an empty queue
 - Enqueue(S:ADT, o:element):ADT** - Inserts object o at the rear of the queue
 - Dequeue(S:ADT):ADT** - Removes the object from the front of the queue; an error occurs if the queue is empty
 - Front(S:ADT):element** - Returns, but does not remove, the front element; an error occurs if the queue is empty

BITs Pilani, Pilani Campus

Queues



Disadvantage

Repeatedly enqueue and dequeue a single element N times.

Finally, $f=r=N$.

- No more elements can be added to the queue,

though there is space in the queue.

Solution

Let f and r wraparound the end of queue.

15
BITs Pilani, Pilani Campus

Queues: An Array Implementation

- Create a queue using an array
- A maximum size N is specified.
- The queue consists of an N -element array Q and two integer variables:
 - f , index of the front element (head – for dequeue)
 - r , index of the element after the rear one (tail – for enqueue)
- Initially, $f=r=0$ and the queue is empty if $f=r$

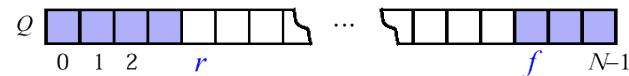


BITs Pilani, Pilani Campus

Queues: An Array Implementation



“wrapped around” configuration



- Each time r or f is incremented, compute this increment as $(r+1)\text{mod}N$ or $(f+1)\text{mod}N$

BITs Pilani, Pilani Campus

Queues: An Array Implementation



- Pseudo code

```
Algorithm size()
return ( $N-f+r$ ) mod  $N$ 
```

```
Algorithm isEmpty()
return ( $f=r$ )
```

```
Algorithm front()
if isEmpty() then
    return Error
return  $Q[f]$ 
```

```
Algorithm dequeue()
```

```
if isEmpty() then
    return Error
```

```
 $Q[f]=\text{null}$ 
```

```
 $f=(f+1) \bmod N$ 
```

```
Algorithm enqueue(o)
```

```
if size =  $N - 1$  then
    return Error
```

```
 $Q[r]=o$ 
```

```
 $r=(r + 1) \bmod N$ 
```

BITs Pilani, Pilani Campus

Arrays: pluses and minuses



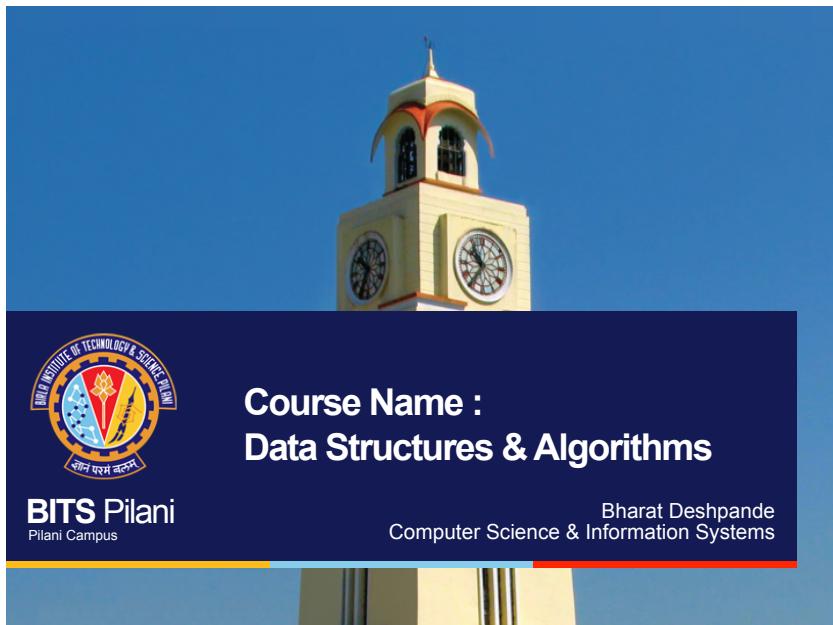
- + Fast element access.

- Impossible to resize.

- Many applications require resizing!
- Required size not always immediately available.

18

BITs Pilani, Pilani Campus



The slide features a large image of the BITs Pilani clock tower against a clear blue sky. Below the image is a dark blue rectangular area containing course information. On the left, the BITs Pilani logo is displayed, which includes a circular emblem with a torch and the text "BITS INSTITUTE OF TECHNOLOGY & SCIENCE PILANI" and "ज्ञान परम वस्तु". To the right of the logo, the course name "Course Name : Data Structures & Algorithms" is written in white. At the bottom right of the dark blue area, the text "Bharat Deshpande Computer Science & Information Systems" is visible. The overall background of the slide is light blue.

Linear data structures



- Here are some of the data structures we have studied so far:
 - Arrays
 - Singly-linked lists and doubly-linked lists
 - Stacks and queues
 - These all have the property that their elements can be adequately displayed in a straight line
- **How to obtain data structures for data that have nonlinear relationships**

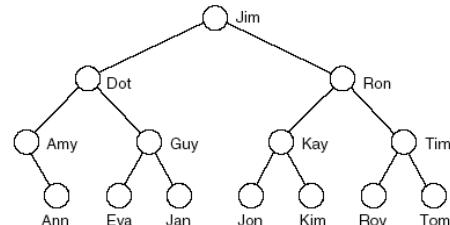
BITs Pilani, Pilani Campus

Tree

Tree represents hierarchy.

Examples of trees:

- Directory tree
- Family tree
- Company organization chart
- Table of contents



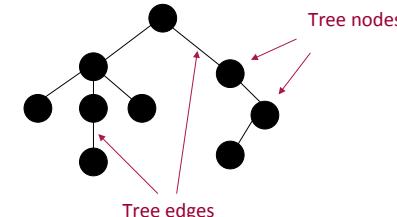
-structure resembles branches of a “tree”, hence the name.

BITs Pilani, Pilani Campus

Trees

Trees have **nodes**. They also have **edges** that connect the nodes.

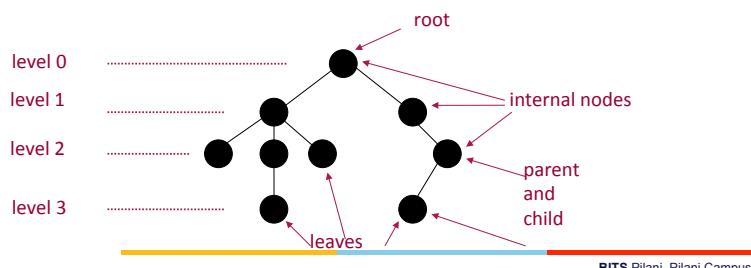
- Between two nodes there is *always only one path*.



BITs Pilani, Pilani Campus

Trees: More Definitions

- Trees that we consider are rooted. Once the **root** is defined (by the user) all nodes have a specific **level**.
- Trees have **internal nodes** and **leaves**. Every node (except the root) has a **parent** and it also has zero or more **children**.



BITs Pilani, Pilani Campus

Tree Terminology (1)

A **vertex (or node)** is an object that can have a name and can carry other associated information

- The first or top node in a tree is called the **root node**.
- An **edge** is a connection between two vertices
- A **path** in a tree is a list of distinct vertices in which successive vertices are connected by edges in the tree.
- The defining property of a tree is that there is precisely one path connecting any two nodes.
- A disjoint set of trees is called a **forest**
- Nodes with no children are **leaves, terminal or external nodes**

BITs Pilani, Pilani Campus

Tree Terminology (2)



Child of a node u :- Any node reachable from u by 1 edge.

Parent node :- If b is a child of a , then a is the parent of b .

- All nodes except root have exactly one parent.

Subtree:-any node of a tree, with all of its descendants.

Depth of a node :

- Depth of root node is 0.

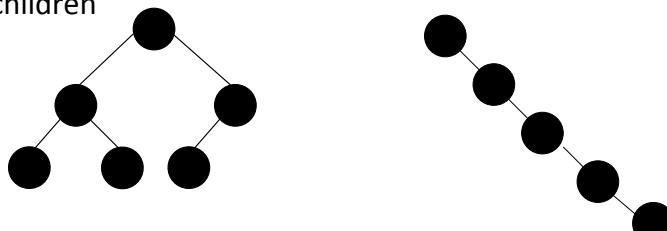
-Depth of any other node is 1 greater than depth of its parent.

Binary Trees



Definition: A binary tree is either empty or it consists of a root together with two binary trees called the left subtree and the right subtree.

A **binary tree** is a tree in which each node has *atmost* 2 children



BITs Pilani, Pilani Campus

Tree Terminology (3)



The size of a tree is the number of nodes in it

Height : Maximum of all depths.

Each node except the root has exactly one node above it in the tree, (i.e. it's parent), and we extend the family analogy talking of children, siblings, or grandparents
Nodes that share parents are called **siblings**.

BITs Pilani, Pilani Campus

Properties of Binary trees



Proper Binary Tree

- Each internal node has exactly two children

Let

n - number of nodes

n_e – number of external nodes or leaves

n_i – number of internal nodes

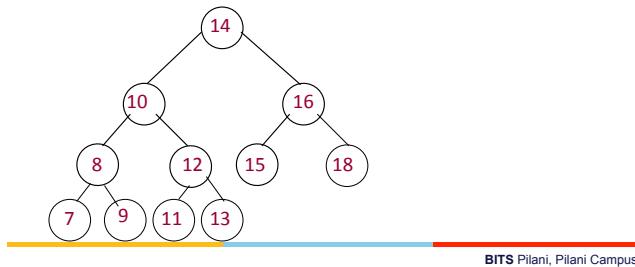
h – height of T , Then the following holds:

- $h+1 \leq n_e \leq 2^h$
- $h \leq n_i \leq 2^h - 1$
- $h + 1 \leq n \leq 2^{h+1} - 1$
- $\log(n+1) - 1 \leq h \leq (n - 1)/2$

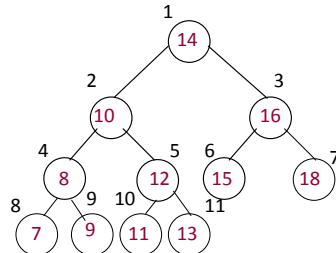
BITs Pilani, Pilani Campus

Complete Binary Trees

- Nodes in trees can contain **keys** (letters, numbers, etc)
- Complete binary tree:** A binary tree in which every level, except possibly the deepest, is completely filled. At depth n , the height of the tree, all nodes must be as far left as possible.



Complete Binary Trees: Array Representation



Complete Binary Trees: Array Representation

Complete Binary Trees can be represented in memory with the use of an array A so that all nodes can be accessed in $O(1)$ time:

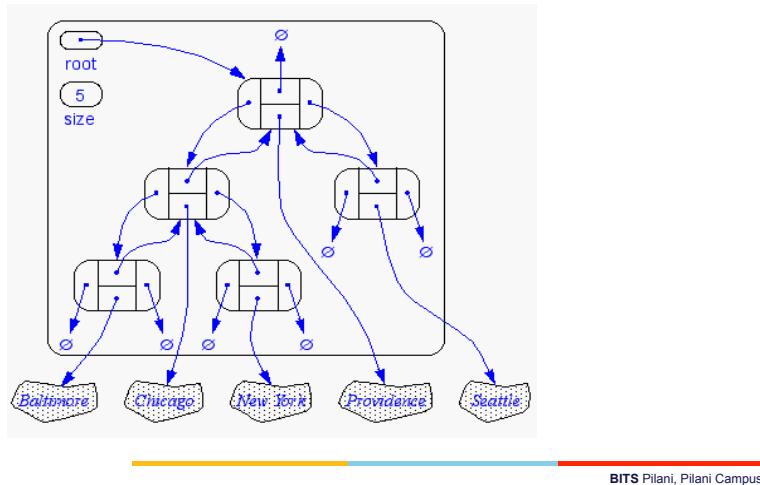
- Label nodes sequentially top-to-bottom and left-to-right
- Left child of $A[i]$ is at position $A[2i]$
- Right child of $A[i]$ is at position $A[2i + 1]$
- Parent of $A[i]$ is at $A[i/2]$

Binary Trees: Linked List Representation

A **Binary tree** is a linked data structure. Each node contains data (including a key and satellite data), and pointers left, right and p .

- Left points to the left child of the node.**
- Right points to the right child of the node.**
- p points to the parent of the node.**
- If a child is missing, the pointer is NIL.
- If a parent is missing, p is NIL.
- The **root** of the tree is the only node for which p is NIL.
- Nodes for which both left and right are NIL are **leaves**.

Binary Trees: Linked List Representation



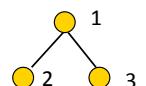
BITS Pilani, Pilani Campus

Binary Tree Traversals

- A binary tree is defined recursively: it consists of a root, a left subtree and a right subtree
- To **traverse** (or walk) the binary tree is to visit each node in the binary tree exactly once.
- Tree traversals are naturally recursive.
- Since a binary tree has three parts, there are six possible ways to traverse the binary tree:
 - root, left, right : preorder (root, right, left)
 - left, root, right: inorder (right, root, left)
 - left, right, root: postorder (right, left, root)

BITS Pilani, Pilani Campus

Binary Tree Traversals



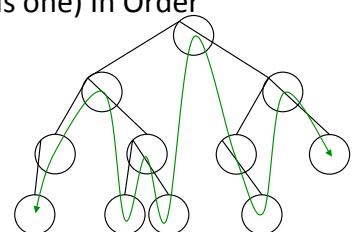
preorder: 1 2 3
inorder: 2 1 3
postorder: 2 3 1

BITS Pilani, Pilani Campus

Tree Traversal: InOrder

In-order traversal

- Visit left subtree (if there is one) In Order
- print the key of the current node
- Visit right subtree (if there is one) In Order



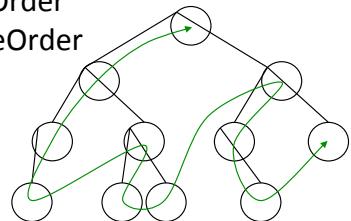
BITS Pilani, Pilani Campus

Tree Traversal: PreOrder

Another common traversal is **PreOrder**.

It goes as deep as possible (visiting as it goes) then left to right

- print root
- Visit left subtree in PreOrder
- Visit right subtree in PreOrder

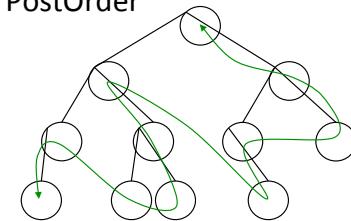


BITS Pilani, Pilani Campus

Tree Traversal: PostOrder

PostOrder traversal also goes as deep as possible, but only visits internal nodes during backtracking.
recursive:

- Visit left subtree in PostOrder
- Visit right subtree in PostOrder
- print root



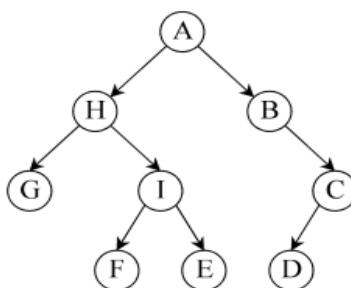
BITS Pilani, Pilani Campus

Tree Traversal: examples

Preorder (DLR) traversal yields: A, H, G, I, F, E, B, C, D

Postorder (LRD) traversal yields: G, F, E, I, H, D, C, B,

In-order (LDR) traversal yields: G, H, F, I, E, A, B, D, C



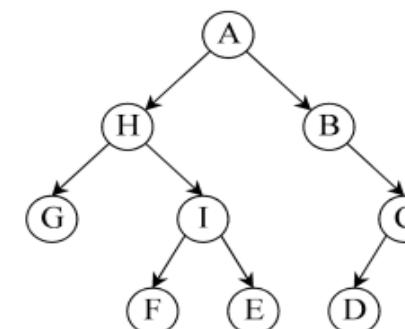
BITS Pilani, Pilani Campus

Tree Traversal: examples

* Preorder (DLR) traversal yields: A, H, G, I, F, E, B, C, D

* Postorder (LRD) traversal yields: G, F, E, I, H, D, C, B, A

* In-order (LDR) traversal yields: G, H, F, I, E, A, B, D, C



BITS Pilani, Pilani Campus



Course Name :
Data Structures & Algorithms

BITS Pilani
Pilani Campus

Bharat Deshpande
Computer Science & Information Systems



Properties of Binary trees

Proper Binary Tree

- Each internal node has exactly two children
- Let

n - number of nodes

n_e – number of external nodes or leaves

n_i – number of internal nodes

h – height of T , Then the following holds:

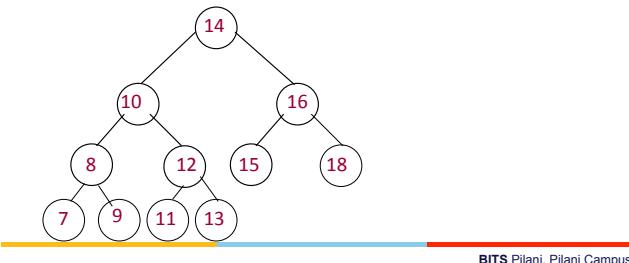
- $h+1 \leq n_e \leq 2^h$
- $h \leq n_i \leq 2^h - 1$
- $h + 1 \leq n \leq 2^{h+1} - 1$
- $\log(n+1) - 1 \leq h \leq (n - 1)/2$

BITS Pilani, Pilani Campus

Complete Binary Trees



- Nodes in trees can contain **keys** (letters, numbers, etc)
- **Complete binary tree**: A binary tree in which every level, except possibly the deepest, is completely filled. At depth n , the height of the tree, all nodes must be as far left as possible.



Complete Binary Trees: Array Representation

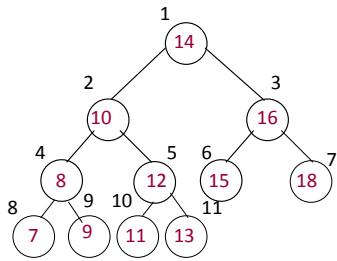


Complete Binary Trees can be represented in memory with the use of an array A so that all nodes can be accessed in $O(1)$ time:

- Label nodes sequentially top-to-bottom and left-to-right
- Left child of $A[i]$ is at position $A[2i]$
- Right child of $A[i]$ is at position $A[2i + 1]$
- Parent of $A[i]$ is at $A[i/2]$

BITS Pilani, Pilani Campus

Complete Binary Trees: Array Representation



BITS Pilani, Pilani Campus

Binary Trees: Linked List Representation



A **Binary tree** is a linked data structure. Each node contains data (including a key and satellite data), and pointers left, right and p.

• **Left** points to the left child of the node.

• **Right** points to the right child of the node.

• **p** points to the parent of the node.

• If a child is missing, the pointer is NIL.

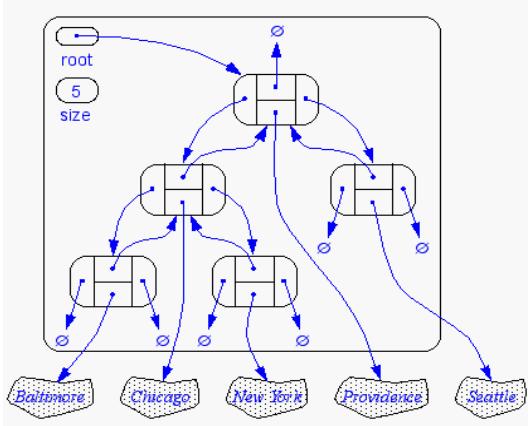
• If a parent is missing, p is NIL.

• The **root** of the tree is the only node for which p is NIL.

• Nodes for which both left and right are NIL are **leaves**.

BITS Pilani, Pilani Campus

Binary Trees: Linked List Representation



BITS Pilani, Pilani Campus

Binary Tree Traversals



– A binary tree is defined recursively: it consists of a root, a left subtree and a right subtree

– To **traverse (or walk)** the binary tree is to visit each node in the binary tree exactly once.

– Tree traversals are naturally recursive.

– Since a binary tree has three parts, there are six possible ways to traverse the binary tree:

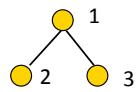
• root, left, right : preorder (root, right, left)

• left, root, right: inorder (right, root, left)

• left, right, root: postorder (right, left, root)

BITS Pilani, Pilani Campus

Binary Tree Traversals



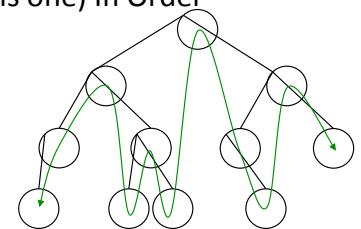
preorder: 1 2 3
inorder: 2 1 3
postorder: 2 3 1

BITS Pilani, Pilani Campus

Tree Traversal: InOrder

In-order traversal

- Visit left subtree (if there is one) In Order
- print the key of the current node
- Visit right subtree (if there is one) In Order



BITS Pilani, Pilani Campus

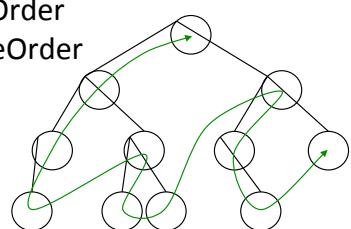
Tree Traversal: PreOrder



Another common traversal is **PreOrder**.

It goes as deep as possible (visiting as it goes) then left to right

- print root
- Visit left subtree in PreOrder
- Visit right subtree in PreOrder



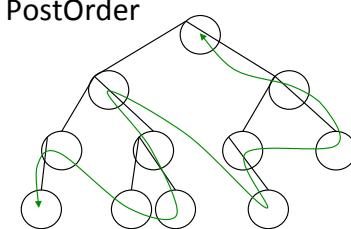
BITS Pilani, Pilani Campus

Tree Traversal: PostOrder



PostOrder traversal also goes as deep as possible, but only visits internal nodes during backtracking.
recursive:

- Visit left subtree in PostOrder
- Visit right subtree in PostOrder
- print root



BITS Pilani, Pilani Campus

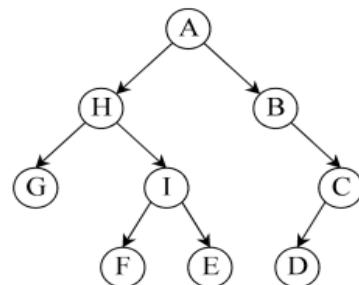
Tree Traversal: examples



Preorder (DLR) traversal yields: A, H, G, I, F, E, B, C, D

Postorder (LRD) traversal yields: G, F, E, I, H, D, C, B,

In-order (LDR) traversal yields: G, H, F, I, E, A, B, D, C



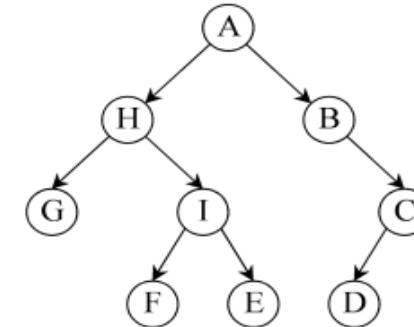
BITS Pilani, Pilani Campus

Tree Traversal: examples

* Preorder (DLR) traversal yields: A, H, G, I, F, E, B, C, D

* Postorder (LRD) traversal yields: G, F, E, I, H, D, C, B, A

* In-order (LDR) traversal yields: G, H, F, I, E, A, B, D, C



BITS Pilani, Pilani Campus

Sorting Problem



- Input : A sequence of n numbers

$\langle a_1, a_2, \dots, a_n \rangle$

- Output : A permutation (reordering)

$\langle a'_1, a'_2, \dots, a'_n \rangle$ of the input sequence such that $a'_1 \leq a'_2 \leq \dots \leq a'_n$

- Solutions : Many!

- First Solution : “Insertion Sort”

15
BITS Pilani, Pilani Campus

Insertion Sort



- **Big idea:**

- Inserting an element into a sorted list in the appropriate position retains the order.
- Works the way many people sort a hand of playing cards.
- Start with an empty left hand and the cards face down on the table.
- We remove one card from the table and insert it in the correct position in left hand.

16
BITS Pilani, Pilani Campus

Insertion Sort

- To find the correct position for a card, we compare it with each of the cards already in the hand, from right to left.

Important :

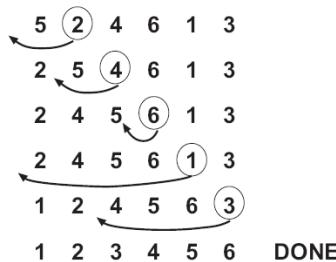
At all times the cards in the left hand are sorted, and these cards were originally the top cards of the pile on the table.

 17

 BITS Pilani, Pilani Campus

Insertion Sort (Con't)

Example:



 19

 BITS Pilani, Pilani Campus

Insertion Sort

Crucial Idea

- Start with a singleton list – sorted trivially.
- Repeatedly insert elements – one at a time – while keeping it sorted.
- Initially, x will need to be the second element and a[1] the ‘sorted part’.
- Sorted part is extended by first inserting the 2nd element, then the 3rd & so on.

 BITS Pilani, Pilani Campus

 18

Insertion Sort – Pseudo Code

```

InsertionSort(A, n) {
    for i = 2 to n {
        key = A[i]
        *insert A[i] into the sorted sequence A[1,...,i-1]
        j = i - 1;
        while (j > 0) and (A[j] > key) {
            A[j+1] = A[j]
            j = j - 1
        }
        A[j+1] = key
    }
}
  
```



 BITS Pilani, Pilani Campus



Insertion Sort - Analysis

- **Best Case Analysis**

The best case for insertion sort occurs when the list is already sorted.
In this case, insertion sort requires $n-1$ comparisons i.e., $O(n)$ complexity.

- **Worst Case Analysis**

for each value of i , what is the maximum number of key comparisons possible?
- Answer: $i - 1$

- Thus, the total time in the worst case is

$$\begin{aligned} T(n) &= 1+2+3+\dots+(n-1) \\ &= n(n-1)/2 \\ &= O(n^2) \end{aligned}$$

21

BITS Pilani, Pilani Campus

Insertion Sort – Average Case



- When we deal with entry i , how far back must we go to insert it?

Answer:

There are i possible positions: not moving at all, moving by one position up to moving by $i - 1$ positions.

- Given randomness, these are equally likely.

23
BITS Pilani, Pilani Campus

Insertion Sort - Analysis

- **Average Case Analysis**

- We assume that all permutations of the keys are equally likely as input.
- We also assume that the keys are distinct.
- We first determine how many key comparisons are done on average to insert one new element into the sorted segment.

22
BITS Pilani, Pilani Campus22
BITS Pilani, Pilani Campus

Insertion Sort – Average Case



Average no. of comparisons

$$= \frac{1}{i} \sum_{j=1}^{i-1} j + \frac{i-1}{i} = \frac{i-1}{2} + 1 - \frac{1}{i}$$

Total =

$$\sum_{i=1}^{n-1} \left(\frac{i-1}{2} + 1 - \frac{1}{i} \right) = \frac{(n-1)(n-2)}{4} + n - 1 - \sum_{i=1}^{n-1} \frac{1}{i}$$

24
BITS Pilani, Pilani Campus

Insertion Sort – Average Case



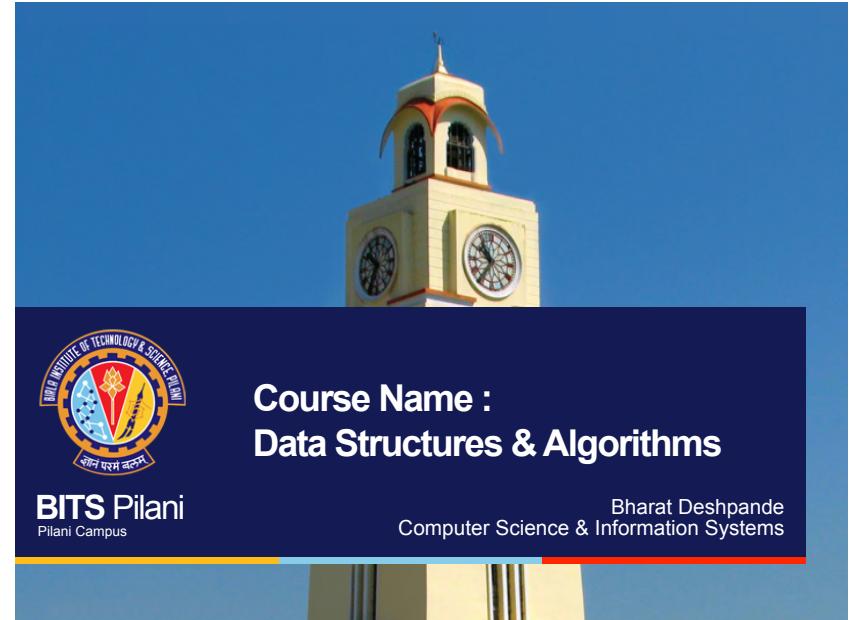
- Well known

$$\sum_{i=1}^n \frac{1}{i} \approx \ln n$$

- Thus, the total number of comparisons
= $O(n^2)$

25

BITS Pilani, Pilani Campus



Sorting Problem



- Input : A sequence of n numbers
 $\langle a_1, a_2, \dots, a_n \rangle$
- Output : A permutation (reordering)
 $\langle a'_1, a'_2, \dots, a'_n \rangle$ of the input sequence such that $a'_1 \leq a'_2 \leq \dots \leq a'_n$
- Solutions : Many!
- First Solution : “Insertion Sort”

2

BITS Pilani, Pilani Campus

Insertion Sort



- **Big idea:**
 - Inserting an element into a sorted list in the appropriate position retains the order.
 - Works the way many people sort a hand of playing cards.
 - Start with an empty left hand and the cards face down on the table.
 - We remove one card from the table and insert it in the correct position in left hand.

3

BITS Pilani, Pilani Campus

Insertion Sort

- To find the correct position for a card, we compare it with each of the cards already in the hand, from right to left.

Important :

At all times the cards in the left hand are sorted, and these cards were originally the top cards of the pile on the table.



BITS Pilani, Pilani Campus

Insertion Sort

Crucial Idea

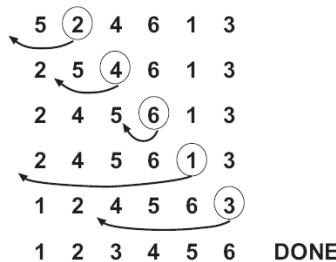
- Start with a singleton list – sorted trivially.
- Repeatedly insert elements – one at a time – while keeping it sorted.
- Initially, x will need to be the second element and a[1] the ‘sorted part’.
- Sorted part is extended by first inserting the 2nd element, then the 3rd & so on.



BITS Pilani, Pilani Campus

Insertion Sort (Con't)

Example:



6
BITS Pilani, Pilani Campus

Insertion Sort – Pseudo Code

```

InsertionSort(A, n) {
    for i = 2 to n {
        key = A[i]
        *insert A[i] into the sorted sequence A[1,...,i-1]
        j = i - 1;
        while (j > 0) and (A[j] > key) {
            A[j+1] = A[j]
            j = j - 1
        }
        A[j+1] = key
    }
}

```

7
BITS Pilani, Pilani Campus



Insertion Sort - Analysis

- **Best Case Analysis**

The best case for insertion sort occurs when the list is already sorted.
In this case, insertion sort requires $n-1$ comparisons i.e., $O(n)$ complexity.

- **Worst Case Analysis**

for each value of i , what is the maximum number of key comparisons possible?
- Answer: $i - 1$

- Thus, the total time in the worst case is

$$\begin{aligned} T(n) &= 1+2+3+\dots+(n-1) \\ &= n(n-1)/2 \\ &= O(n^2) \end{aligned}$$

8

BITS Pilani, Pilani Campus

9

BITS Pilani, Pilani Campus

Insertion Sort – Average Case



- When we deal with entry i , how far back must we go to insert it?

Answer:

There are i possible positions: not moving at all, moving by one position up to moving by $i - 1$ positions.

- Given randomness, these are equally likely.

10

BITS Pilani, Pilani Campus

11

BITS Pilani, Pilani Campus

Insertion Sort - Analysis

- **Average Case Analysis**

- We assume that all permutations of the keys are equally likely as input.
- We also assume that the keys are distinct.
- We first determine how many key comparisons are done on average to insert one new element into the sorted segment.

Average no. of comparisons

$$= \frac{1}{i} \sum_{j=1}^{i-1} j + \frac{i-1}{i} = \frac{i-1}{2} + 1 - \frac{1}{i}$$

Total =

$$\sum_{i=1}^{n-1} \left(\frac{i-1}{2} + 1 - \frac{1}{i} \right) = \frac{(n-1)(n-2)}{4} + n - 1 - \sum_{i=1}^{n-1} \frac{1}{i}$$

Insertion Sort – Average Case

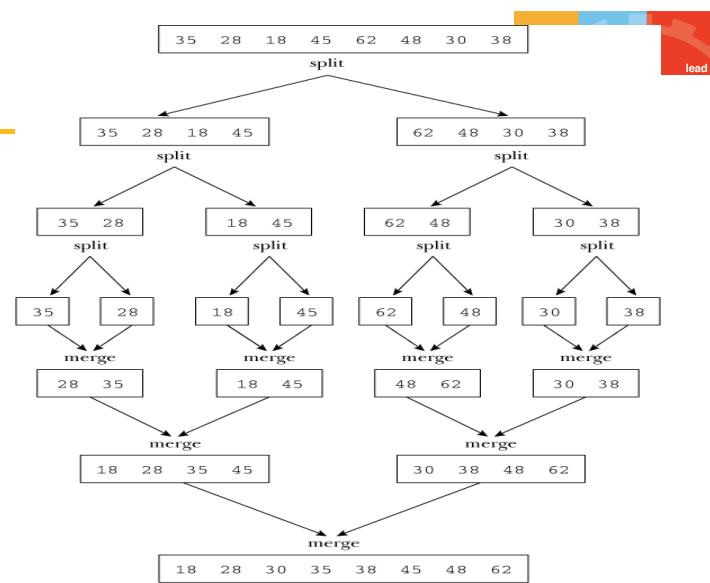


- Well known

$$\sum_{i=1}^n \frac{1}{i} \approx \ln n$$

- Thus, the total number of comparisons
= $O(n^2)$

12
BITS Pilani, Pilani Campus



BITS Pilani, Pilani Campus

Merge Sort

- Divide:** Divide the $n_{_element}$ sequence to be sorted into two subsequences of $n/2$ elements each.
- Conquer:** Sort the two subsequences recursively using merge sort.
- Combine:** Merge the two sorted subsequences to produce a sorted list.
- The general algorithm for the merge sort is as follows:
 - Find the mid-position of the list.
 - Merge sort the first sublist.
 - Merge sort the second sublist.
 - Merge the first sublist and the second sublist.

13
BITS Pilani, Pilani Campus

Merging two sorted Array



- Once the sublists are sorted, the next step in the merge sort algorithm is to merge the sorted sublists.
- Suppose L_1 and L_2 are two sorted lists as follows:
 - L_1 : 2, 7, 16, 35
 - L_2 : 5, 20, 25, 40, 50
- Merge L_1 and L_2 into a third list, say L_3 .
- The merge process is as follows:
repeatedly compare, using a loop, the elements of L_1 with the elements of L_2 and copy the smaller element into L_3 .

15

BITS Pilani, Pilani Campus

Example

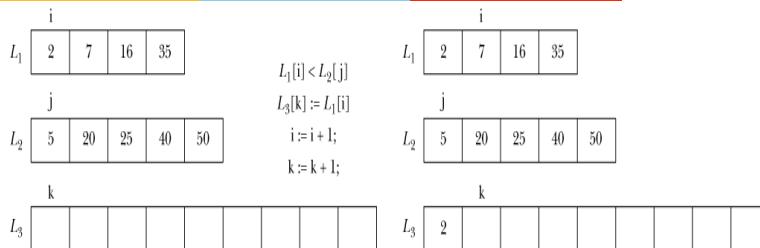


FIGURE 9.23 L_1 , L_2 , and L_3 before and after the first iteration

- First compare $L_1[1]$ with $L_2[1]$ and see that $L_1[1] < L_2[1]$, so copy $L_1[1]$ into $L_3[1]$
- Time :** If both the list has n elements each then the merging process takes $2n$ time in the worst case.



Merge Sort Complexity (Con't)



• Running time analysis (Con't):

$$\begin{aligned}
 T(1) &= 1 && \text{Initial condition} \\
 T(n) &= 2T\left(\frac{n}{2}\right) + n \\
 &= 2(2T\left(\frac{n}{4}\right) + \frac{n}{2}) + n \\
 &= 4T\left(\frac{n}{4}\right) + 2n \\
 &= 4(2T\left(\frac{n}{8}\right) + \frac{n}{4}) + 2n \\
 &= 8T\left(\frac{n}{8}\right) + 3n \\
 &\vdots \\
 &= 2^k T\left(\frac{n}{2^k}\right) + kn && \text{Since } n = 2^k, \text{ we have } k = \log_2 n \\
 &= nT(1) + n \log_2 n && \text{Since } T(1) = 1 \\
 &= n + n \log_2 n \\
 &= O(n \log n)
 \end{aligned}$$



Merge Sort (Complexity)



- Running time analysis:
 - $T(n)$: worst-case running time of merge sort to sort n numbers (assume n is a power of 2)

Running time analysis can be modeled as an recurrence equation:

$$T(1) = 1, \text{ if } n=1$$

$$T(n) = 2T(n/2) + n, \text{ if } n>1$$



Design Strategy



Divide and Conquer

- is a general algorithm design paradigm:
 - Divide:** divide the input data S in two or more disjoint subsets S_1, S_2, \dots
 - Recur:** solve the subproblems recursively
 - Conquer:** combine the solutions for S_1, S_2, \dots , into a solution for S
- The base case for the recursion are subproblems of constant size
- Analysis can be done using **recurrence equations**



Master Method (Appendix)

- Many divide-and-conquer recurrence equations have the form:

$$T(n) = \begin{cases} c & \text{if } n < d \\ aT(n/b) + f(n) & \text{if } n \geq d \end{cases}$$

The Master Theorem:

- if $f(n)$ is $O(n^{\log_b a - \varepsilon})$, then $T(n)$ is $\Theta(n^{\log_b a})$
- if $f(n)$ is $\Theta(n^{\log_b a} \log^k n)$, then $T(n)$ is $\Theta(n^{\log_b a} \log^{k+1} n)$
- if $f(n)$ is $\Omega(n^{\log_b a + \varepsilon})$, then $T(n)$ is $\Theta(f(n))$, provided $af(n/b) \leq \delta f(n)$ for some $\delta < 1$.

20

BITS Pilani, Pilani Campus

Heap Data Structure

- A **heap** is a complete binary tree with the following two properties:

– **Structural property:** all levels are full, except possibly the last one, which is filled from left to right

– **Order (heap) property:** for any node x
 $\text{Parent}(x) \geq x$

- From the heap property, it follows that:

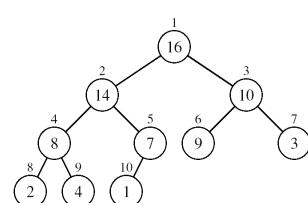
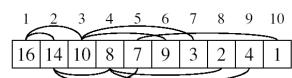
- “The root is the maximum element of the heap!”

21

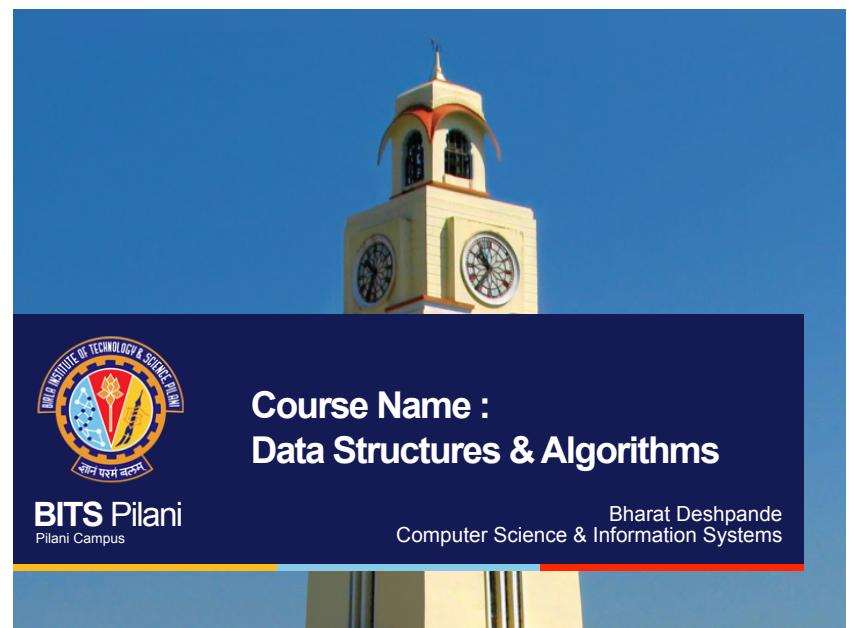
BITS Pilani, Pilani Campus

Heap represented as Array

- A heap can be stored as an array A .
- Root of tree is $A[1]$
- Left child of $A[i] = A[2i]$
- Right child of $A[i] = A[2i + 1]$
- Parent of $A[i] = A[\lfloor i/2 \rfloor]$



22
BITS Pilani, Pilani Campus



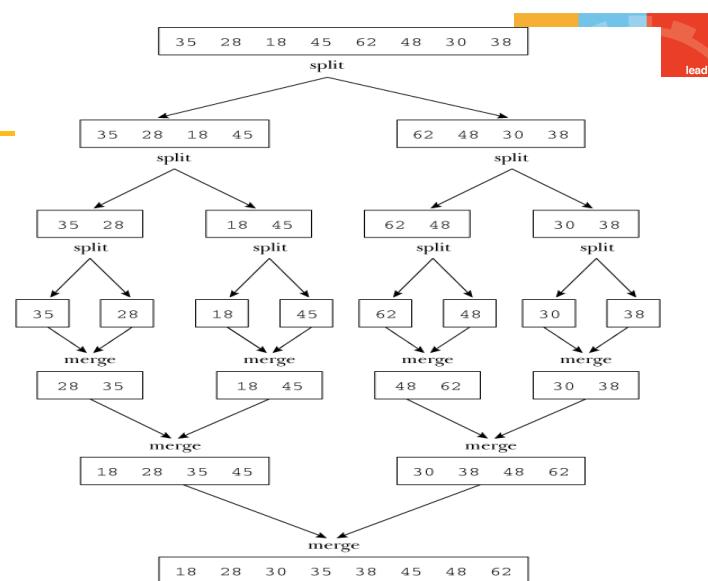
BITS Pilani
Pilani Campus

Bharat Deshpande
Computer Science & Information Systems

Sorting Problem

- Input : A sequence of n numbers
 $\langle a_1, a_2, \dots, a_n \rangle$
- Output : A permutation (reordering)
 $\langle a'_1, a'_2, \dots, a'_n \rangle$ of the input sequence such that $a'_1 \leq a'_2 \leq \dots \leq a'_n$
- Solutions : Many!
- First Solution : “Insertion Sort”

2
BITS Pilani, Pilani Campus



Merge Sort

- **Divide:** Divide the n_element sequence to be sorted into two subsequences of $n/2$ elements each.
- **Conquer:** Sort the two subsequences recursively using merge sort.
- **Combine:** Merge the two sorted subsequences to produce a sorted list.
- The general algorithm for the merge sort is as follows:
- If the list is of size greater than 1, then
 - a. Find the mid-position of the list.
 - b. Merge sort the first sublist.
 - c. Merge sort the second sublist.
 - d. Merge the first sublist and the second sublist.

3
BITS Pilani, Pilani Campus

Merging two sorted Array

- Once the sublists are sorted, the next step in the merge sort algorithm is to merge the sorted sublists.
- Suppose L_1 and L_2 are two sorted lists as follows:
 - $L_1: 2, 7, 16, 35$
 - $L_2: 5, 20, 25, 40, 50$
- Merge L_1 and L_2 into a third list, say L_3 .
- The merge process is as follows:
 repeatedly compare, using a loop, the elements of L_1 with the elements of L_2 and copy the smaller element into L_3 .

Example

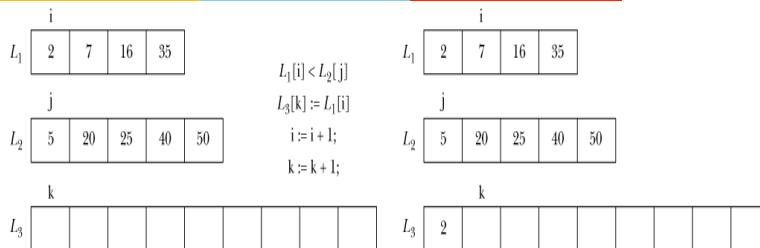


FIGURE 9.23 L_1 , L_2 , and L_3 before and after the first iteration

- First compare $L_1[1]$ with $L_2[1]$ and see that $L_1[1] < L_2[1]$, so copy $L_1[1]$ into $L_3[1]$
- Time :** If both the list has n elements each then the merging process takes $2n$ time in the worst case.



Merge Sort Complexity (Con't)



• Running time analysis (Con't):

$$\begin{aligned}
 T(1) &= 1 && \text{Initial condition} \\
 T(n) &= 2T\left(\frac{n}{2}\right) + n \\
 &= 2(2T\left(\frac{n}{4}\right) + \frac{n}{2}) + n \\
 &= 4T\left(\frac{n}{4}\right) + 2n \\
 &= 4(2T\left(\frac{n}{8}\right) + \frac{n}{4}) + 2n \\
 &= 8T\left(\frac{n}{8}\right) + 3n \\
 &\vdots \\
 &= 2^k T\left(\frac{n}{2^k}\right) + kn && \text{Since } n = 2^k, \text{ we have } k = \log_2 n \\
 &= nT(1) + n \log_2 n && \text{Since } T(1) = 1 \\
 &= n + n \log_2 n \\
 &= O(n \log n)
 \end{aligned}$$



Merge Sort (Complexity)



• Running time analysis:

- $T(n)$: worst-case running time of merge sort to sort n numbers (assume n is a power of 2)

Running time analysis can be modeled as an recurrence equation:

$$T(1) = 1, \text{ if } n=1$$

$$T(n) = 2T(n/2) + n, \text{ if } n>1$$



Design Strategy



Divide and Conquer

- is a general algorithm design paradigm:
 - Divide:** divide the input data S in two or more disjoint subsets S_1, S_2, \dots
 - Recur:** solve the subproblems recursively
 - Conquer:** combine the solutions for S_1, S_2, \dots , into a solution for S
- The base case for the recursion are subproblems of constant size
- Analysis can be done using **recurrence equations**



Master Method (Appendix)

- Many divide-and-conquer recurrence equations have the form:

$$T(n) = \begin{cases} c & \text{if } n < d \\ aT(n/b) + f(n) & \text{if } n \geq d \end{cases}$$

- The Master Theorem:**

- if $f(n)$ is $O(n^{\log_b a - \varepsilon})$, then $T(n)$ is $\Theta(n^{\log_b a})$
- if $f(n)$ is $\Theta(n^{\log_b a} \log^k n)$, then $T(n)$ is $\Theta(n^{\log_b a} \log^{k+1} n)$
- if $f(n)$ is $\Omega(n^{\log_b a + \varepsilon})$, then $T(n)$ is $\Theta(f(n))$, provided $af(n/b) \leq \delta f(n)$ for some $\delta < 1$.

10

BITS Pilani, Pilani Campus

Heap Data Structure

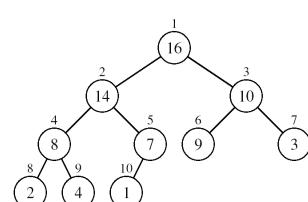
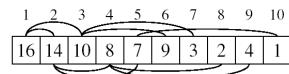
- A **heap** is a complete binary tree with the following two properties:
 - Structural property:** all levels are full, except possibly the last one, which is filled from left to right
 - Order (heap) property:** for any node x $\text{Parent}(x) \geq x$
- From the heap property, it follows that:
 - The root is the maximum element of the heap!**
 - All leaf nodes automatically have the heap property!**

11

BITS Pilani, Pilani Campus

Heap represented as Array

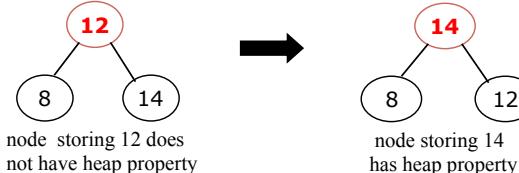
- A heap can be stored as an array A .
- Root of tree is $A[1]$
- Left child of $A[i] = A[2i]$
- Right child of $A[i] = A[2i + 1]$
- Parent of $A[i] = A[\lfloor i/2 \rfloor]$



- Height of an heap storing n elements is $O(\log n)$**

12

BITS Pilani, Pilani Campus



- Given a node that does not have the heap property, you can give it the heap property by exchanging its value with the value of the larger child

- Notice that the child may have lost the heap property
- Recursively fix the children until all of them satisfy the max-heap property.

13

BITS Pilani, Pilani Campus



Max-Heapify

MaxHeapify(A, i)

```

1.  $l \leftarrow \text{left}(i)$ 
2.  $r \leftarrow \text{right}(i)$ 
3. if  $l \leq \text{heap-size}[A]$  and  $A[l] > A[i]$ 
4.   then  $\text{largest} \leftarrow l$ 
5. else  $\text{largest} \leftarrow i$ 
6. if  $r \leq \text{heap-size}[A]$  and  $A[r] > A[\text{largest}]$ 
7.   then  $\text{largest} \leftarrow r$ 
8. if  $\text{largest} \neq i$ 
9.   then exchange  $A[i] \leftrightarrow A[\text{largest}]$ 
10.    MaxHeapify( $A, \text{largest}$ )
    
```

• Running time for MaxHeapify is $O(\log n)$

Assumption:

Left(i) and Right(i) are max-heaps.

14

BITS Pilani, Pilani Campus



Building a Heap

- Use **MaxHeapify** to convert an array A into a max-heap.
- How?**
- Call MaxHeapify on each element in a bottom-up manner.

BuildMaxHeap(A)

```

1.  $\text{heap-size}[A] \leftarrow \text{length}[A]$ 
2. for  $i \leftarrow \lfloor \text{length}[A]/2 \rfloor$  downto 1
3.   do MaxHeapify( $A, i$ )
    • Why from  $i = \lfloor \text{length}[A]/2 \rfloor$  ?
    • Running time  $O(n \log n)$ 
    
```

15

BITS Pilani, Pilani Campus

Heap Sort



- Start by building a max-heap on all elements in A .
 - Maximum element is in the root, $A[1]$.
- Move the maximum element to its correct final position.
 - Exchange $A[1]$ with $A[n]$.
- Discard $A[n]$ – it is now sorted.
 - Decrement $\text{heap-size}[A]$.
- Restore the max-heap property on $A[1..n-1]$.
 - Call **MaxHeapify($A, 1$)**.
- Repeat until $\text{heap-size}[A]$ is reduced to 2.

16
BITS Pilani, Pilani Campus



Heap Sort

Heap Sort(A)

- Build-Max-Heap(A)
- for $i \leftarrow \text{length}[A]$ downto 2
 - do exchange $A[1] \leftrightarrow A[i]$
 - $\text{heap-size}[A] \leftarrow \text{heap-size}[A] - 1$
 - MaxHeapify($A, 1$)
- Build-Max-Heap takes $O(n \log n)$ and each of the $n-1$ calls to Max-Heapify takes time $O(\log n)$.
- Therefore, $T(n) = O(n \log n)$

17

BITS Pilani, Pilani Campus



Course Name :
Data Structures & Algorithms

BITS Pilani
Pilani Campus

Bharat Deshpande
Computer Science & Information Systems



Lower Bound for Sorting

- **Merge sort**
 - worst-case running time is $O(N \log N)$
- **Insertion Sort**
 - Worst-case running time is $O(N^2)$
- **Heap Sort**
 - worst-case running time is $O(N \log N)$
- Are there better algorithms?
- **Goal:** Prove that any sorting algorithm based on only comparisons takes at least $O(N \log N)$ comparisons in the worst case (worse-case input) to sort N elements.

BITS Pilani, Pilani Campus

Lower Bound for Sorting (con't..)



- Suppose we want to sort N distinct elements
- How many possible orderings do we have for N elements?
- We can have $N!$ possible orderings (e.g., the sorted output for a,b,c can be $a\ b\ c$, $b\ a\ c$, $a\ c\ b$, $c\ a\ b$, $c\ b\ a$, $b\ c\ a$.)

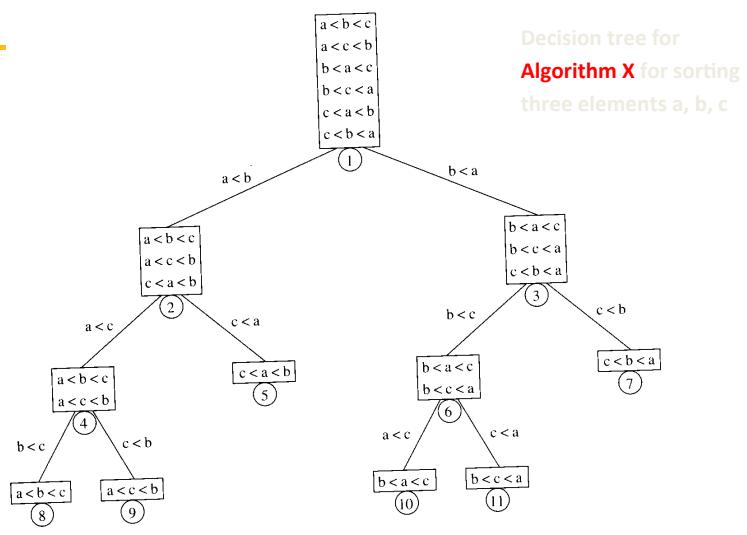
BITS Pilani, Pilani Campus

Lower Bound for Sorting (con't..)



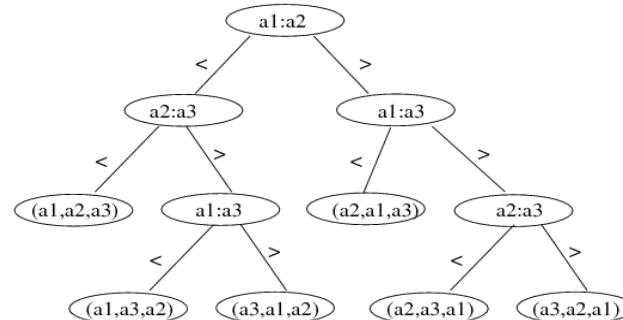
- Any comparison-based sorting process can be represented as a binary **decision tree**.
 - Each node represents a set of possible orderings, consistent with all the comparisons that have been made
 - The tree edges are results of the comparisons

BITS Pilani, Pilani Campus



Lower Bound for Sorting (con't)

- A different algorithm would have a different decision tree
- Decision tree for **Insertion Sort** on 3 elements:



BITS Pilani, Pilani Campus

Lower Bound for Sorting (con't)

- The worst-case number of comparisons used by the sorting algorithm is equal to the **depth of the deepest leaf**
 - The average number of comparisons used is equal to the average depth of the leaves
- A decision tree to sort N elements must have **N!** leaves
 - a binary tree of depth d has at most 2^d leaves
 - ⇒ the tree must have depth at least $\lceil \log_2(N!) \rceil$
- Therefore, any sorting algorithm based on only comparisons between elements requires at least $\lceil \log_2(N!) \rceil$ comparisons in the worst case.

BITS Pilani, Pilani Campus

Lower Bound for Sorting (con't..)

$$\begin{aligned}
 \log_2(N!) &= \log(N(N-1)(N-2)\cdots(2)(1)) \\
 &= \log N + \log(N-1) + \log(N-2) + \cdots + \log 2 + \log 1 \\
 &\geq \log N + \log(N-1) + \log(N-2) + \cdots + \log(N/2) \\
 &\geq \frac{N}{2} \log \frac{N}{2} \\
 &= \frac{N}{2} \log N - \frac{N}{2} \\
 &= \Omega(N \log N)
 \end{aligned}$$

- Any sorting algorithm based on comparisons between elements requires $\Omega(N \log N)$ comparisons.

BITS Pilani, Pilani Campus



Linear time sorting

- Can we do better? (linear time algorithm)

Yes, if the input has special structure

- Counting sort, radix sort, Bucket sort

BITS Pilani, Pilani Campus



Bucket Sort

- The bucket sort makes assumptions about the data being sorted
- The n elements to be sorted are integers in the range $[0, N-1]$
- Consequently, we can achieve better than $O(n \log n)$ run times

Idea:

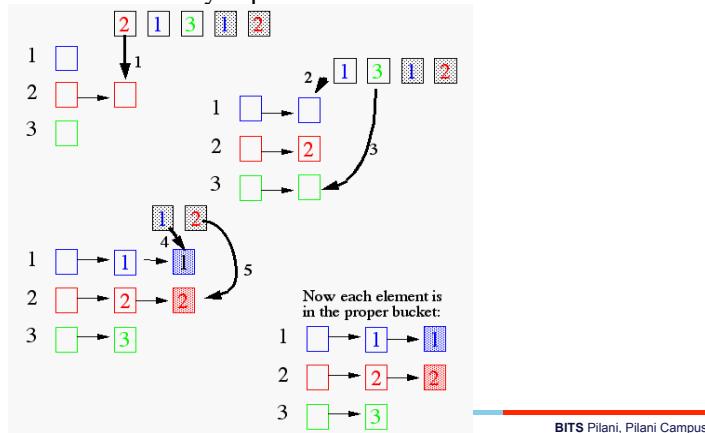
- Not based on Comparison
- But using keys as indices into a bucket array.

BITS Pilani, Pilani Campus



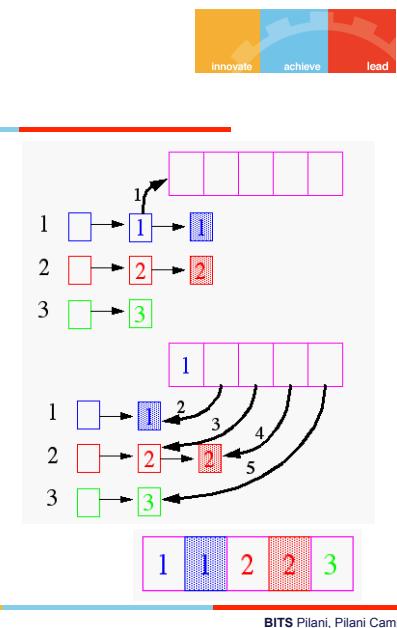
Bucket Sort (Example 1)

Each element of the array is put in one of the N “buckets”



Bucket Sort

Now, pull the elements from the buckets into the array



Bucket Sort Algorithm



```
bucketSort(s)
Let B be an array of n lists, each of which is initially empty
for each item x in S do
    let k be the key of x
    remove x from S and insert it at the end of bucket B[k]
for i = 0 to n-1 do
    for each item x in list B[i] do
        remove x from B[i] and insert it at the end of S.
```

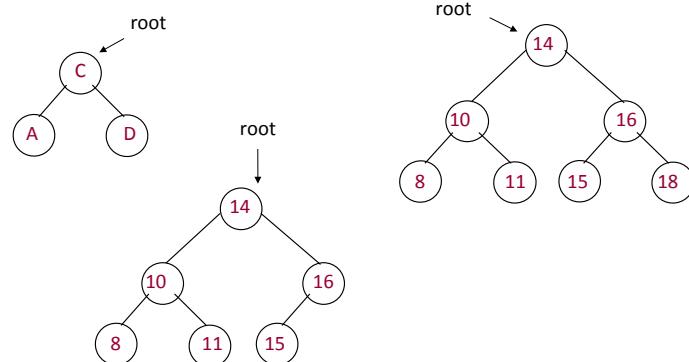
Complexity

$O(n + N)$

If N is $O(n)$ then we can sort in $O(n)$ time.

13
BITS Pilani, Pilani Campus

Binary Search Trees: Examples



BITS Pilani, Pilani Campus

Binary Search Trees



A **Binary Search Tree (BST)** is a binary tree with the following properties:

- The key of a node is *always greater* than the keys of the nodes in its left subtree
- The key of a node is *always smaller* than the keys of the nodes in its right subtree

BITS Pilani, Pilani Campus

Building a BST



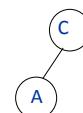
Build a BST from a sequence of nodes read one at a time

Example: Inserting **C A B L M** (in this order!)

1) Insert C



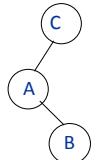
2) Insert A



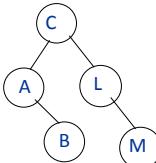
BITS Pilani, Pilani Campus

Building a BST

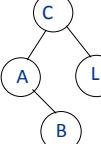
3) Insert B



5) Insert M



4) Insert L



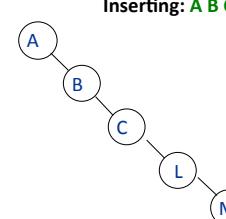
BITS Pilani, Pilani Campus

Building a BST

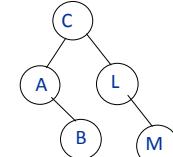
Is there a unique BST for letters A B C L M ?

NO! Different input sequences result in different trees

Inserting: A B C L M



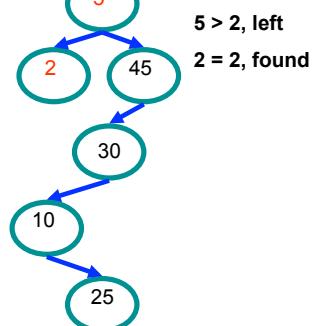
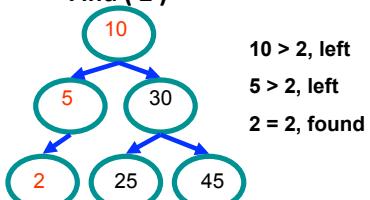
Inserting: C A B L M



BITS Pilani, Pilani Campus

Example Binary Searches

- Find (2)



BITS Pilani, Pilani Campus

Recursive Search of Binary Tree

```

Node Find( Node n, Value key) {
    if (n == null)                                // Not found
        return( n );
    else if (n.data == key)                         // Found it
        return( n );
    else if (n.data > key)                          // In left subtree
        return Find( n.left );
    else                                              // In right subtree
        return Find( n.right );
}

```

BITS Pilani, Pilani Campus

Complexity of Search



- Running time of searching in a BST is proportional to the height of the tree.

If n is the number of nodes in a BST, then

- **Best Case** – $O(\log n)$
- **Worst Case** – $O(n)$

Binary Search Tree - Insertion

Insert Algorithm

- If value we want to insert < key of current node, we have to go to the left subtree
- Otherwise we have to go to the right subtree
- If the current node is empty (not existing) create a node with the value we are inserting and place it here.

21

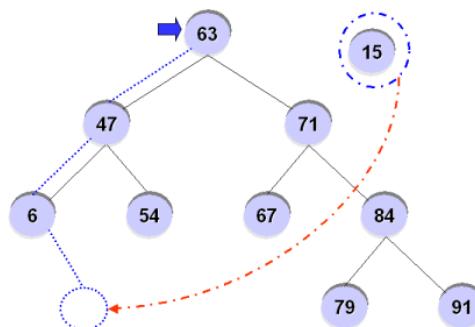
BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

Insertion - Example



For example, inserting '15' into the BST?



BITS Pilani, Pilani Campus



Binary Search Tree - Deletion

- How do we delete a node from BST?
Similar to the insert function, after deletion of a node, the property of the BST must be maintained.

BITS Pilani, Pilani Campus

Binary Search Tree - Deletion



There are 3 possible cases

Case1 : Node to be deleted has no children

→ We just delete the node.

Case2 : Node to be deleted has only one child

→ Replace the node with its child
and make the parent of the
deleted node to be a parent of the
child of the deleted node

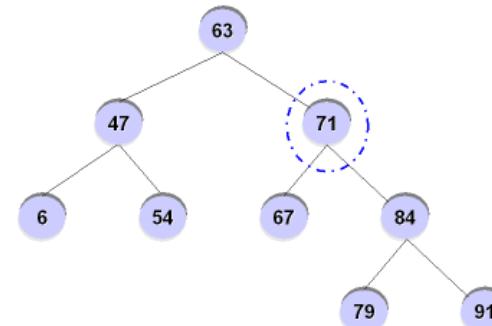
Case3 : Node to be deleted has two children

BITS Pilani, Pilani Campus

Binary Search Tree - Deletion



Node to be deleted has two children



BITS Pilani, Pilani Campus

Binary Search Tree - Deletion



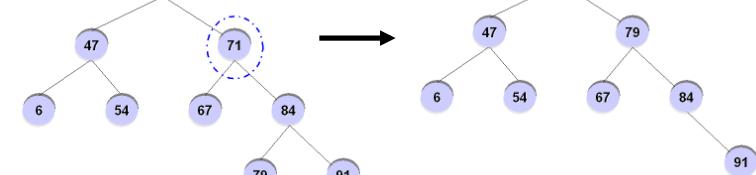
Node to be deleted has two children

Steps:

- Find minimum value of right subtree
- Delete minimum node of right subtree but keep its value
- Replace the value of the node to be deleted by the minimum value whose node was deleted earlier.

BITS Pilani, Pilani Campus

Binary Search Tree - Deletion



BITS Pilani, Pilani Campus



Course Name :
Data Structures & Algorithms

BITS Pilani
Pilani Campus

Bharat Deshpande
Computer Science & Information Systems



Binary Search Trees - Disadvantages

- Unbalanced Binary Search Trees are bad
- Worst case Performance is $O(n)$
- No better than sequential searching.
- We look at correcting this problem.

BITS Pilani, Pilani Campus

Height Balance Property



• Height Balance Property

For every internal node v of a binary search tree, the *heights of the children of v can differ by at most 1*.

- It is a property which characterizes the structure of a BST

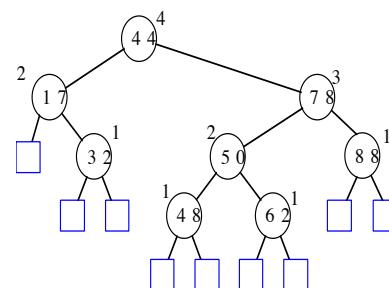
Definition: Any Binary Search Tree T that satisfies height balance property is said to be an **AVL tree**. (Adelson-Velskii & Landis)

BITS Pilani, Pilani Campus

AVL Tree



Example



An example of an AVL tree where the heights are shown next to the nodes:

BITS Pilani, Pilani Campus

AVL tree



- Simple Observation

Subtree of an AVL tree is itself an AVL tree.

- Important Consequence

Keeps the height small

BITS Pilani, Pilani Campus

Height of an AVL Tree



- Knowing $n(h-1) > n(h-2)$, we get $n(h) > 2n(h-2)$.

- So $n(h) > 2n(h-2)$, $n(h) > 4n(h-4)$, $n(h) > 8n(h-6)$, ...

by induction,

$$n(h) > 2^i n(h-2i)$$

Choosing i such that $h - 2i$ is either 1 or 2

We get: $n(h) > 2^{h/2-1}$

Taking logarithms: $h < 2\log n(h) + 2$

Thus the height of an AVL tree is $O(\log n)$

Height of an AVL Tree



- Fact:** The **height** of an AVL tree storing n keys is $O(\log n)$.

- Proof:** Let us bound $n(h)$: the minimum number of internal nodes of an AVL tree of height h .

- We easily see that $n(1) = 1$ and $n(2) = 2$

- For $h > 2$, AVL tree with minimum no. of nodes is such that both its subtrees are AVL trees with minimum no. of nodes

- Such an AVL tree contains the root node, one AVL subtree of height $n-1$ and another of height $n-2$.

- That is, $n(h) = 1 + n(h-1) + n(h-2)$



BITS Pilani, Pilani Campus

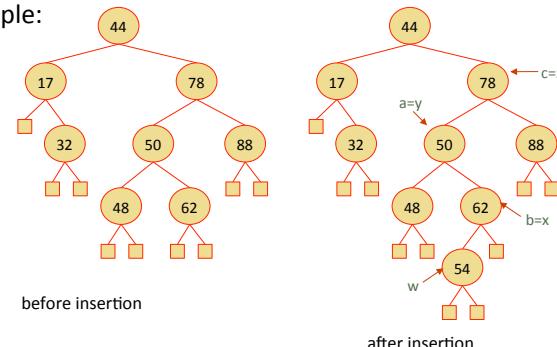
Insertion in an AVL Tree



- Insertion is as in a binary search tree

- Always done by expanding an external node.

- Example:



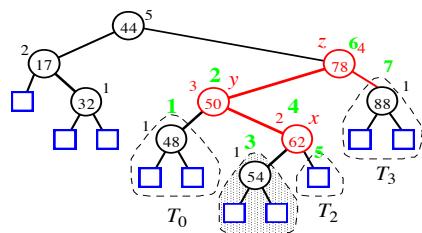
It is no longer balanced

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

Insertion

- Identifying the node at which height balance property is violated.



9
BITS Pilani, Pilani Campus

Insertion

- Relabel x, y, z as a, b, c such that a precedes b and b precedes c in an inorder traversal of T.
- There are four possible cases to do this mapping.
- Restructuring (Unified way)**

Replace z with the node called b, make the children of this node be a and c while maintaining the inorder relationships of all nodes in T.

Search & Repair strategy

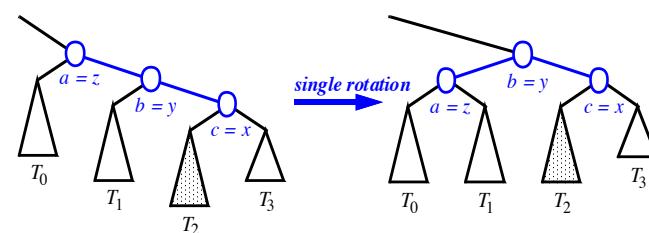
- Let z be the 1st node encountered in going up from the inserted node towards the root at which the height balance property is violated.
- Let y be the child of z with higher height
- Let x be the child of y with higher height
- Trinode Structuring**

Balance the subtree rooted at z.

10
BITS Pilani, Pilani Campus

Restructuring (Single Rotation)

- Single Rotations:



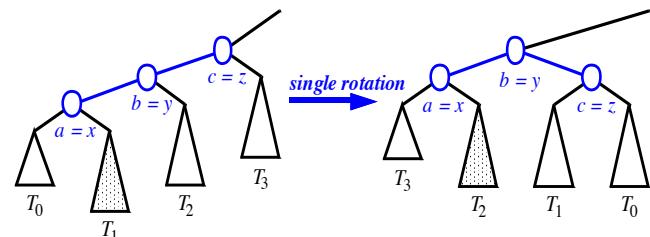
BITS Pilani, Pilani Campus

11
BITS Pilani, Pilani Campus

Restructuring (Single Rotation)



- Single Rotations:

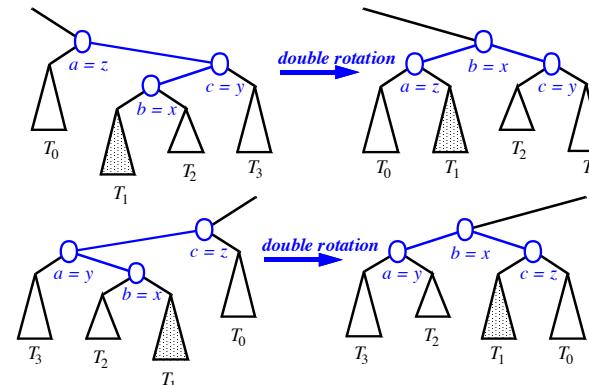


BITS Pilani, Pilani Campus

Restructuring (Double Rotation)

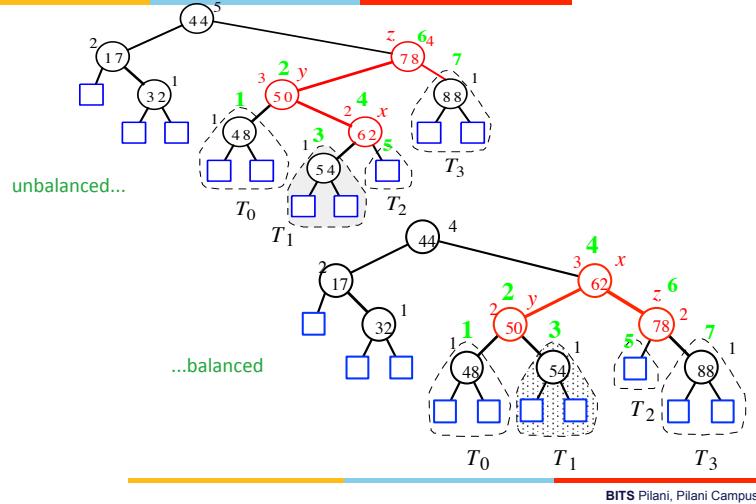


- double rotations:



BITS Pilani, Pilani Campus

Insertion Example, continued



BITS Pilani, Pilani Campus

Deletion in an AVL Tree

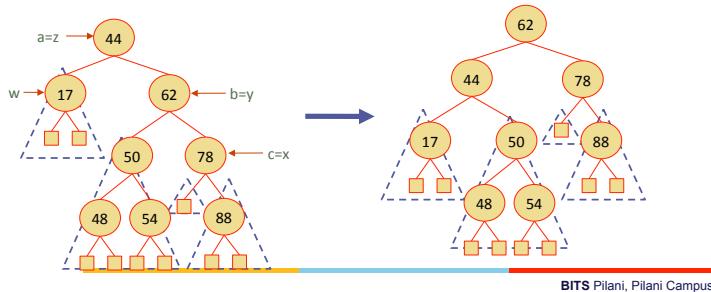


- Let ***z*** be the **first unbalanced** node encountered while travelling up the tree from ***w***. Also,
- let ***y*** be the child of ***z*** with the larger height,
- let ***x*** be the child of ***y*** defined as follows;
 - If one of the children of ***y*** is taller than the other, choose ***x*** as the taller child of ***y***.
 - If both children of ***y*** have the same height, select ***x*** be the child of ***y*** on the same side as ***y*** (i.e., if ***y*** is the left child of ***z***, then ***x*** is the left child of ***y***; and if ***y*** is the right child of ***z*** then ***x*** is the right child of ***y***.)

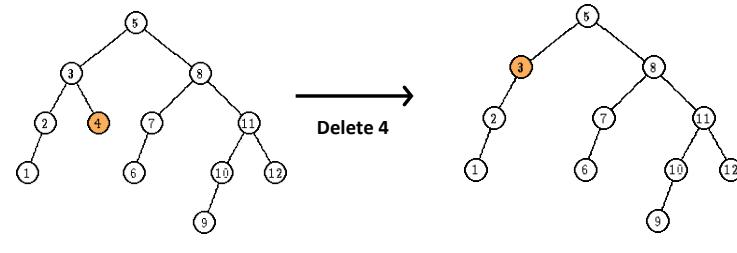
BITS Pilani, Pilani Campus

Trinode Restructuring

- We perform `restructure(x)` to restore balance at z.
- As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of T is reached

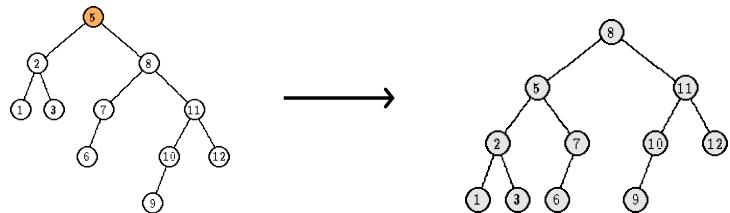


Deletion - Example



BITS Pilani, Pilani Campus

Example – Contd.



Imbalance at 5
Perform rotation with 8

BITS Pilani, Pilani Campus

Running Times for AVL Trees

- a single restructure is $O(1)$
 - using a linked-structure binary tree
- find is $O(\log n)$
 - height of tree is $O(\log n)$, no restructures needed
- insert is $O(\log n)$
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$
- remove is $O(\log n)$
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$

BITS Pilani, Pilani Campus

Multi-Way Search Trees



- Each internal node of a multi-way search tree T:
 - has at least two children, i.e., each internal node is a d-node, $d \geq 2$
 - stores a collection of items of the form (k, x) , where k is a key and x is an element
 - Each d – node with children v_1, \dots, v_d contains $d - 1$ items, stored in increasing order
 - “contains” 2 pseudo-items: $k_0 = -\infty$, $k_d = \infty$. For each item (k, x) stored at a node in the subtree of v_i rooted at v_i , $i = 1, \dots, D$, we have $k_{i-1} \leq k \leq k_i$
- Children of each internal node are “between” items
- all keys in the subtree rooted at the child fall between keys of those items

21

BITS Pilani, Pilani Campus

Performance



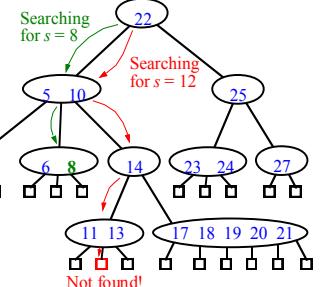
- **Worst Case** – $O(h \log d)$
where d is the maximum number of children of any node
- If d is constant, then running time is $O(h)$
- **Prime Goal**
Try to keep h as small as possible

BITS Pilani, Pilani Campus

Multi-way Searching



- Similar to binary searching
 - If search key $s < k_1$ search the leftmost child
 - If $s > k_{d-1}$, search the rightmost child
- That's it in a binary tree; what about if $d > 2$?
 - Find two keys k_{i-1} and k_i between which s falls, and search the child v_i .



BITS Pilani, Pilani Campus

(2, 4) or (2, 3, 4) Trees



(2, 4) tree is a multi-way search tree with the following two properties:

Size property:

Nodes may contain 1, 2 or 3 items & a node with k items has $k + 1$ children

Depth property:

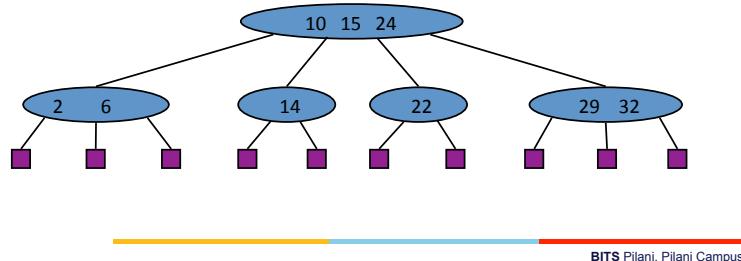
All the external nodes have same depth

BITS Pilani, Pilani Campus

(2,4) Trees



- Depending on the number of children, an internal node of a (2,4) tree is called a 2-node, 3-node or 4-node



Proof



Let h be the height of (2,4) tree T storing n items.
Then the number of external nodes in T is at most 4^h
& the number of external nodes in T is at least 2^h
Therefore,

$$2^h \leq n + 1 \leq 4^h$$

Taking logarithm in base 2, we get that
 $h \leq \log(n+1)$ and $\log(n+1) \leq 2h$, which proves the result.

Height of (2, 4) tree



Result:

A multi-way search tree storing n items has $n + 1$ external nodes.

- Can be proved by induction

Main Result :

Height of a (2, 4) tree storing n items is $O(\log n)$

Insertion in (2, 4) tree



- To insert an item with key k in a (2,4) tree T
- Assume that tree T has no element with key k .
- Perform search for k .
This search will terminate at an external node, say z .
- Let v be parent of z
- Insert new into node v and add a new child w (an external node) to v on the left of z .

Observe: Insertion method preserves the depth property.

But it may violate size property.

- If v was previously a 4-node, after insertion will become a 5-node, which is not allowed.
- This violation of size property is called **overflow** at v .
- Overflow needs to be resolved

Insertion in (2, 4) tree

Resolving Overflow

Perform **Split Operation**

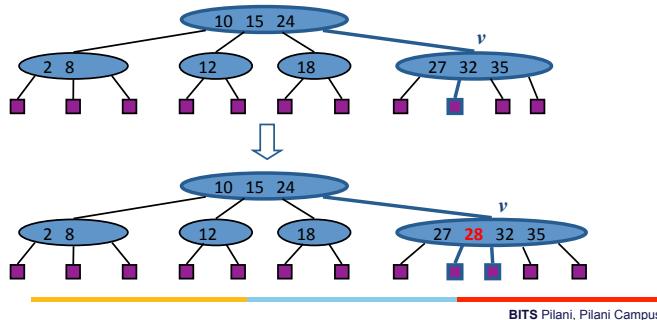
- Let v_1, v_2, v_3, v_4, v_5 be children of v .
- Let k_1, k_2, k_3, k_4 be the keys stored at v .
- Replace v with two nodes v' and v'' , where
- v' is a 3-node with children v_1, v_2, v_3 storing keys k_1, k_2
- v'' is a 2-node with children v_4, v_5 storing key k_4
- If v is a root of T , create a new root u or else u be parent of v .

29

BITS Pilani, Pilani Campus

Insertion

- We insert a new item (k, o) at the parent v of the leaf reached by searching for k
 - We preserve the depth property but
 - We may cause an **overflow** (i.e., node v may become a 5-node)
- Example: inserting key 28 causes an overflow



BITS Pilani, Pilani Campus

Insertion in (2, 4) tree

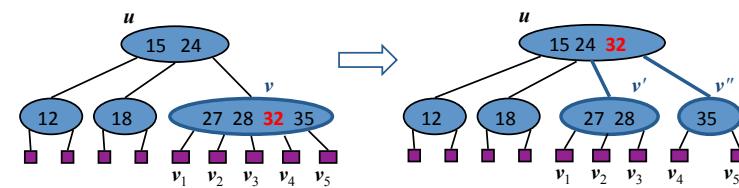
- Insert key k_3 into u and make v' & v'' children of u , so that if v was i th child of u , then v' & v'' become i th & $(i+1)$ th child of u .
- One split operation takes $O(1)$ time.
- As a consequence of split operation on a node v , a new overflow may occur at the parent u of v .
- In worst case this propagates all the way up to the root, where it is finally resolved.

30

BITS Pilani, Pilani Campus

Overflow and Split

- We handle an **overflow** at a 5-node v with a **split operation**:
- The overflow may propagate to the parent node u



BITS Pilani, Pilani Campus



Analysis of Insertion

Algorithm $\text{insert}(k, o)$

1. We search for key k to locate the insertion node v
2. We add the new entry (k, o) at node v
3. **while** $\text{overflow}(v)$
 if $\text{isRoot}(v)$
 create a new empty root above v
 $v \leftarrow \text{split}(v)$

- Let T be a $(2,4)$ tree with n items
 - Tree T has $O(\log n)$ height
 - Step 1 takes $O(\log n)$ time because we visit $O(\log n)$ nodes
 - Step 2 takes $O(1)$ time
 - Step 3 takes $O(\log n)$ time because each split takes $O(1)$ time and we perform $O(\log n)$ splits
- Thus, an insertion in a $(2,4)$ tree takes $O(\log n)$ time

BITs Pilani, Pilani Campus



Deletion in $(2, 4)$ tree

To remove an item with key k from a $(2, 4)$ tree T .

- Perform search in T for an item with key k .

Key Point:

Removing item can always be reduced to the case where the item to be removed is stored at a node v whose children are external nodes. (similar to BST)

- Suppose, item k is stored in the i th item at a node z that has only internal node children.
- Find the leftmost internal node v in the subtree rooted at $(i+1)$ th child of z , such that children of v are all external nodes.
- Swap the item to be deleted from z with the first item of v .

34

BITs Pilani, Pilani Campus

Deletion in $(2, 4)$ tree



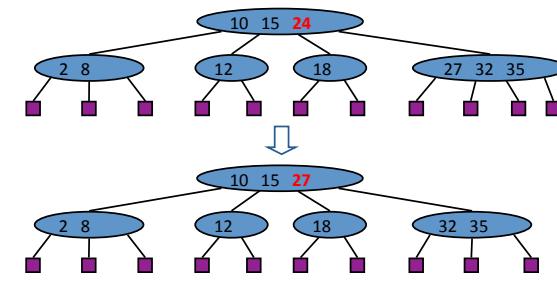
- Remove the item k from v and remove the i th external node of v .
- Removal preserves the depth property.
- But size property may get violated
 A 2-node may become a 1-node, which is not allowed.
- This is called **underflow**.
- Underflow needs to be resolved.

35
BITs Pilani, Pilani Campus

Deletion



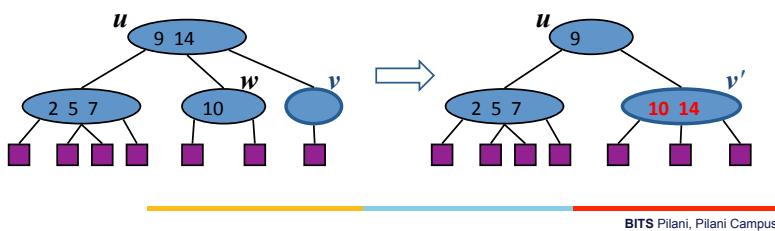
- **Example:** to delete key 24, we replace it with 27 (inorder successor)



BITs Pilani, Pilani Campus

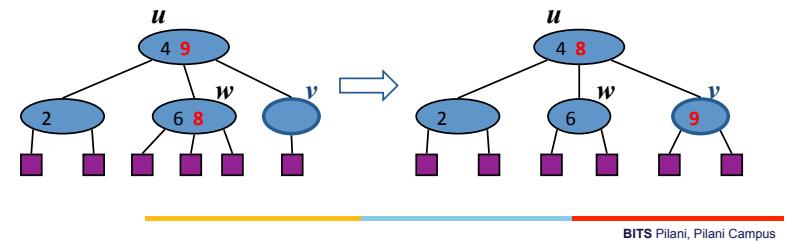
Underflow and Fusion

- Deleting an entry from a node v may cause an **underflow**, where node v becomes a 1-node with one child and no keys
- To handle an underflow at node v with parent u , we consider two cases
- Case 1:** the adjacent siblings of v are 2-nodes
 - Fusion operation:** we merge v with an adjacent sibling w and move an entry from u to the merged node v'
 - After a fusion, the underflow may propagate to the parent u



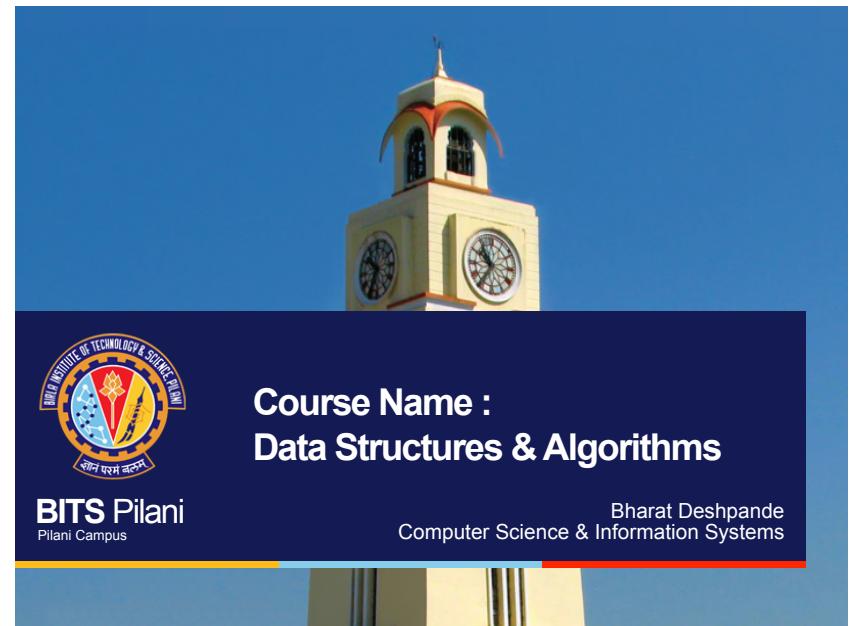
Underflow and Transfer

- Case 2:** an immediate sibling w of v is a 3-node or a 4-node
 - Transfer operation:**
 - we move a child of w to v
 - we move an item from u to v
 - we move an item from w to u
 - After a transfer, no underflow occurs



Analysis of Deletion

- Let T be a $(2,4)$ tree with n items
 - Tree T has $O(\log n)$ height
- In a deletion operation
 - We visit $O(\log n)$ nodes to locate the node from which to delete the entry
 - We handle an underflow with a series of $O(\log n)$ fusions, followed by at most one transfer
 - Each fusion and transfer takes $O(1)$ time
- Thus, deleting an item from a $(2,4)$ tree takes $O(\log n)$ time



Multi-Way Search Trees



- Each internal node of a multi-way search tree T:
 - has at least two children, i.e., each internal node is a d-node, $d \geq 2$
 - stores a collection of items of the form (k, x) , where k is a key and x is an element
 - Each d – node with children v_1, \dots, v_d contains $d - 1$ items, stored in increasing order
 - “contains” 2 pseudo-items: $k_0 = -\infty$, $k_d = \infty$. For each item (k, x) stored at a node in the subtree of v_i rooted at v_i , $i = 1, \dots, D$, we have $k_{i-1} \leq k \leq k_i$
- Children of each internal node are “between” items
- all keys in the subtree rooted at the child fall between keys of those items

2
BITS Pilani, Pilani Campus

Performance



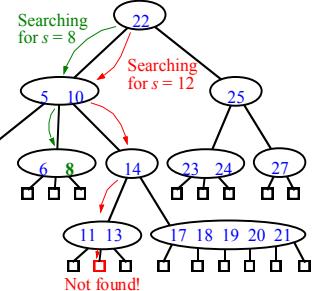
- **Worst Case** – $O(h \log d)$
where d is the maximum number of children of any node
- If d is constant, then running time is $O(h)$
- **Prime Goal**
Try to keep h as small as possible

BITS Pilani, Pilani Campus

Multi-way Searching



- Similar to binary searching
 - If search key $s < k_1$ search the leftmost child
 - If $s > k_{d-1}$, search the rightmost child
- That's it in a binary tree; what about if $d > 2$?
 - Find two keys k_{i-1} and k_i between which s falls, and search the child v_i .



BITS Pilani, Pilani Campus

(2, 4) or (2, 3, 4) Trees



(2, 4) tree is a multi-way search tree with the following two properties:

Size property:

Nodes may contain 1, 2 or 3 items & a node with k items has $k + 1$ children

Depth property:

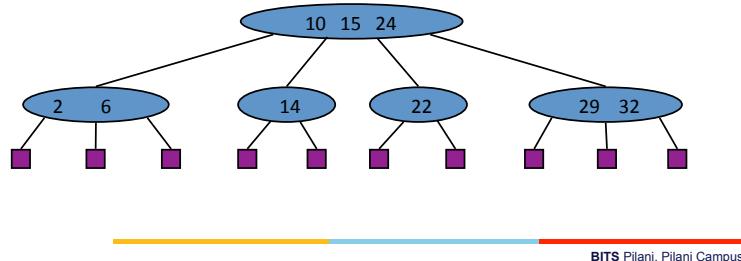
All the external nodes have same depth

BITS Pilani, Pilani Campus

(2,4) Trees



- Depending on the number of children, an internal node of a (2,4) tree is called a 2-node, 3-node or 4-node



Proof



Let h be the height of (2,4) tree T storing n items.
Then the number of external nodes in T is at most 4^h
& the number of external nodes in T is at least 2^h
Therefore,

$$2^h \leq n + 1 \leq 4^h$$

Taking logarithm in base 2, we get that
 $h \leq \log(n+1)$ and $\log(n+1) \leq 2h$, which proves the result.

Height of (2, 4) tree



Result:

A multi-way search tree storing n items has $n + 1$ external nodes.

- Can be proved by induction

Main Result :

Height of a (2, 4) tree storing n items is $O(\log n)$

Insertion in (2, 4) tree



- To insert an item with key k in a (2,4) tree T
- Assume that tree T has no element with key k .
- Perform search for k .
This search will terminate at an external node, say z .
- Let v be parent of z
- Insert new into node v and add a new child w (an external node) to v on the left of z .

Observe: Insertion method preserves the depth property.

But it may violate size property.

- If v was previously a 4-node, after insertion will become a 5-node, which is not allowed.
- This violation of size property is called **overflow** at v .
- Overflow needs to be resolved

Insertion in (2, 4) tree

Resolving Overflow

Perform **Split Operation**

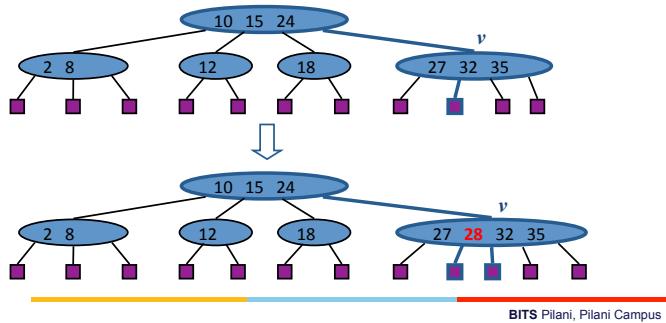
- Let v_1, v_2, v_3, v_4, v_5 be children of v .
- Let k_1, k_2, k_3, k_4 be the keys stored at v .
- Replace v with two nodes v' and v'' , where
- v' is a 3-node with children v_1, v_2, v_3 storing keys k_1, k_2
- v'' is a 2-node with children v_4, v_5 storing key k_4
- If v is a root of T , create a new root u or else u be parent of v .

10

BITS Pilani, Pilani Campus

Insertion

- We insert a new item (k, o) at the parent v of the leaf reached by searching for k
 - We preserve the depth property but
 - We may cause an **overflow** (i.e., node v may become a 5-node)
- Example: inserting key 28 causes an overflow



Insertion in (2, 4) tree

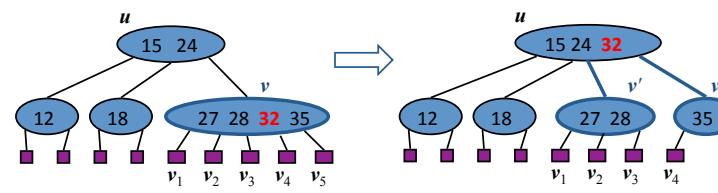
- Insert key k_3 into u and make v' & v'' children of u , so that if v was i th child of u , then v' & v'' become i th & $(i+1)$ th child of u .
- One split operation takes $O(1)$ time.
- As a consequence of split operation on a node v , a new overflow may occur at the parent u of v .
- In worst case this propagates all the way up to the root, where it is finally resolved.

11

BITS Pilani, Pilani Campus

Overflow and Split

- We handle an **overflow** at a 5-node v with a **split operation**:
- The overflow may propagate to the parent node u



BITS Pilani, Pilani Campus

Analysis of Insertion



Algorithm $\text{insert}(k, o)$

1. We search for key k to locate the insertion node v
2. We add the new entry (k, o) at node v
3. **while** $\text{overflow}(v)$
 if $\text{isRoot}(v)$
 create a new empty root above v
 $v \leftarrow \text{split}(v)$

- Let T be a $(2,4)$ tree with n items
 - Tree T has $O(\log n)$ height
 - Step 1 takes $O(\log n)$ time because we visit $O(\log n)$ nodes
 - Step 2 takes $O(1)$ time
 - Step 3 takes $O(\log n)$ time because each split takes $O(1)$ time and we perform $O(\log n)$ splits
- Thus, an insertion in a $(2,4)$ tree takes $O(\log n)$ time

BITs Pilani, Pilani Campus

Deletion in $(2, 4)$ tree



- Remove the item k from v and remove the i th external node of v .
- Removal preserves the depth property.
- But size property may get violated
 A 2-node may become a 1-node, which is not allowed.
- This is called **underflow**.
- Underflow needs to be resolved.

16
BITs Pilani, Pilani Campus

Deletion in $(2, 4)$ tree



To remove an item with key k from a $(2, 4)$ tree T .

- Perform search in T for an item with key k .

Key Point:

Removing item can always be reduced to the case where the item to be removed is stored at a node v whose children are external nodes. (similar to BST)

- Suppose, item k is stored in the i th item at a node z that has only internal node children.
- Find the leftmost internal node v in the subtree rooted at $(i+1)$ th child of z , such that children of v are all external nodes.
- Swap the item to be deleted from z with the first item of v .

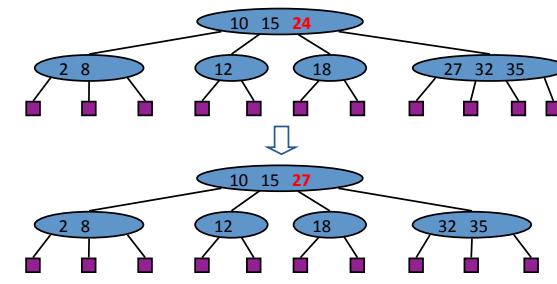
15

BITs Pilani, Pilani Campus

Deletion



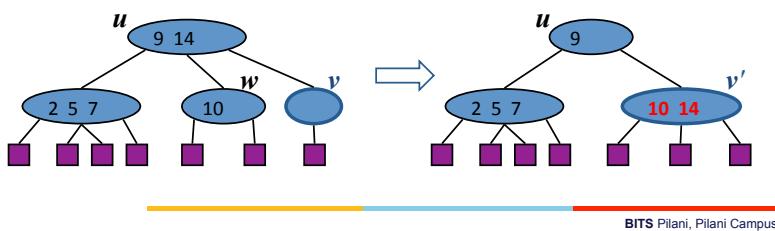
- **Example:** to delete key 24, we replace it with 27 (inorder successor)



BITs Pilani, Pilani Campus

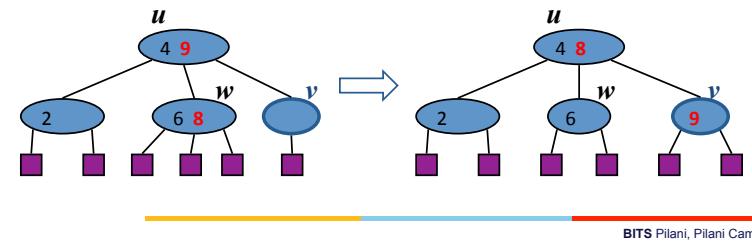
Underflow and Fusion

- Deleting an entry from a node v may cause an **underflow**, where node v becomes a 1-node with one child and no keys
- To handle an underflow at node v with parent u , we consider two cases
- Case 1:** the adjacent siblings of v are 2-nodes
 - Fusion operation:** we merge v with an adjacent sibling w and move an entry from u to the merged node v'
 - After a fusion, the underflow may propagate to the parent u



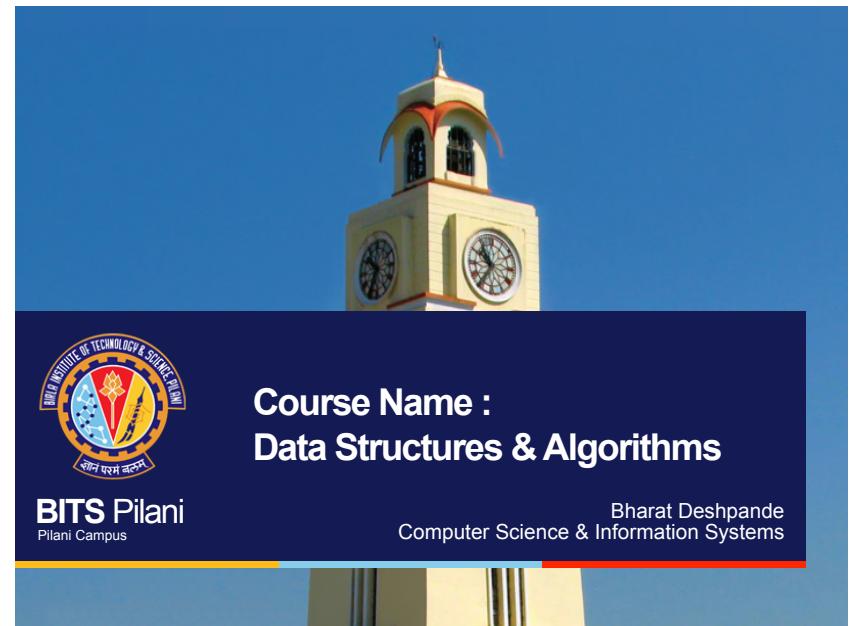
Underflow and Transfer

- Case 2:** an immediate sibling w of v is a 3-node or a 4-node
 - Transfer operation:**
 - we move a child of w to v
 - we move an item from u to v
 - we move an item from w to u
 - After a transfer, no underflow occurs



Analysis of Deletion

- Let T be a $(2,4)$ tree with n items
 - Tree T has $O(\log n)$ height
- In a deletion operation
 - We visit $O(\log n)$ nodes to locate the node from which to delete the entry
 - We handle an underflow with a series of $O(\log n)$ fusions, followed by at most one transfer
 - Each fusion and transfer takes $O(1)$ time
- Thus, deleting an item from a $(2,4)$ tree takes $O(\log n)$ time



(2, 4) or (2, 3, 4) Trees



(2, 4) tree is a multi-way search tree with the following two properties:

Size property:

Nodes may contain 1, 2 or 3 items & a node with k items has $k + 1$ children

Depth property:

All the external nodes have same depth

BITs Pilani, Pilani Campus

Insertion in (2, 4) tree



- To insert an item with key k in a (2, 4) tree T
- Assume that tree T has no element with key k .
- Perform search for k .
This search will terminate at an external node, say z .
- Let v be parent of z .
- Insert new into node v and add a new child w (an external node) to v on the left of z .

Observe: Insertion method preserves the depth property.

But it may violate size property.

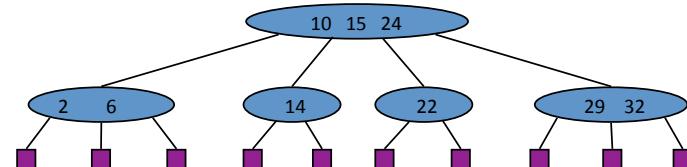
- If v was previously a 4-node, after insertion will become a 5-node, which is not allowed.
- This violation of size property is called **overflow** at v .
- Overflow needs to be resolved

BITs Pilani, Pilani Campus

(2,4) Trees



- Depending on the number of children, an internal node of a (2,4) tree is called a 2-node, 3-node or 4-node



BITs Pilani, Pilani Campus

Insertion in (2, 4) tree



Resolving Overflow

Perform **Split Operation**

- Let v_1, v_2, v_3, v_4, v_5 be children of v .
- Let k_1, k_2, k_3, k_4 be the keys stored at v .
- Replace v with two nodes v' and v'' , where
 - v' is a 3-node with children v_1, v_2, v_3 storing keys k_1, k_2
 - v'' is a 2-node with children v_4, v_5 storing key k_4
- If v is a root of T , create a new root u or else u be parent of v .

Insertion in (2, 4) tree

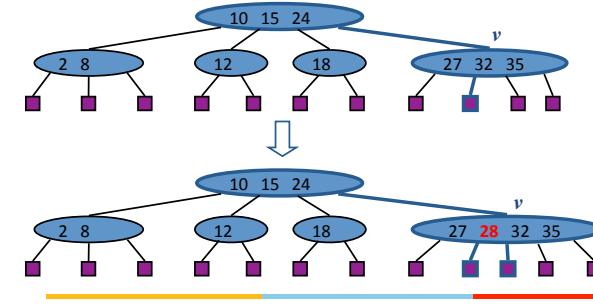
- Insert key k_3 into u and make v' & v'' children of u , so that if v was i th child of u , then v' & v'' become i th & $(i+1)$ th child of u .
- One split operation takes $O(1)$ time.
- As a consequence of split operation on a node v , a new overflow may occur at the parent u of v .
- In worst case this propagates all the way up to the root, where it is finally resolved.

6

BITS Pilani, Pilani Campus

Insertion

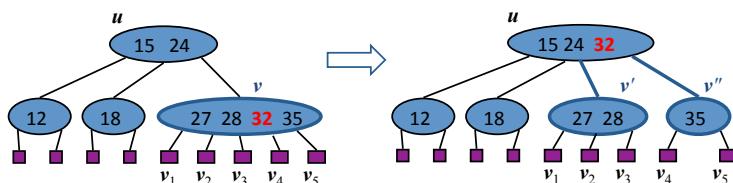
- We insert a new item (k, o) at the parent v of the leaf reached by searching for k
 - We preserve the depth property but
 - We may cause an **overflow** (i.e., node v may become a 5-node)
- Example: inserting key 28 causes an overflow



BITS Pilani, Pilani Campus

Overflow and Split

- We handle an **overflow** at a 5-node v with a **split operation**:
- The overflow may propagate to the parent node u



BITS Pilani, Pilani Campus

Analysis of Insertion

Algorithm $\text{insert}(k, o)$

- We search for key k to locate the insertion node v
- We add the new entry (k, o) at node v
- while** $\text{overflow}(v)$
 - if** $\text{isRoot}(v)$
 - create a new empty root above v
 - $v \leftarrow \text{split}(v)$

- Let T be a (2,4) tree with n items
 - Tree T has $O(\log n)$ height
 - Step 1 takes $O(\log n)$ time because we visit $O(\log n)$ nodes
 - Step 2 takes $O(1)$ time
 - Step 3 takes $O(\log n)$ time because each split takes $O(1)$ time and we perform $O(\log n)$ splits
- Thus, an insertion in a (2,4) tree takes $O(\log n)$ time

Deletion in (2, 4) tree

To remove an item with key k from a (2, 4) tree T.

- Perform search in T for an item with key k.

Key Point:

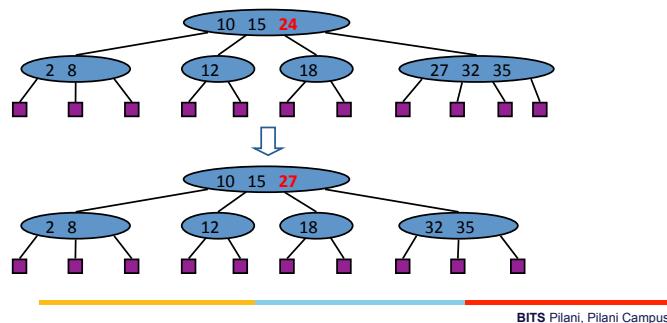
Removing item can always be reduced to the case where the item to be removed is stored at a node v whose children are external nodes. (similar to BST)

- Suppose, item k is stored in the ith item at a node z that has only internal node children.
- Find the leftmost internal node v in the subtree rooted at (i+1)th child of z, such that children of v are all external nodes.
- Swap the item to be deleted from z with the first item of v.



Deletion

- **Example:** to delete key 24, we replace it with 27 (inorder successor)



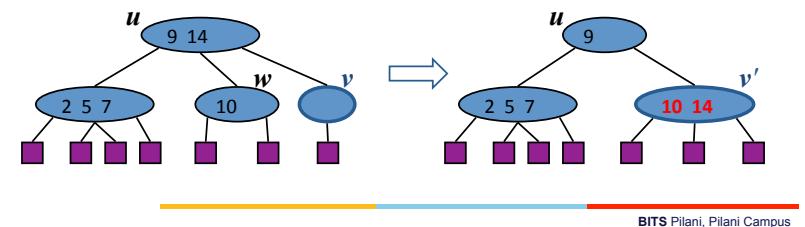
Deletion in (2, 4) tree

- Remove the item k from v and remove the ith external node of v.
- Removal preserves the depth property.
- But size property may get violated
- A 2-node may become a 1-node, which is not allowed.
- This is called **underflow**.
- Underflow needs to be resolved.



Underflow and Fusion

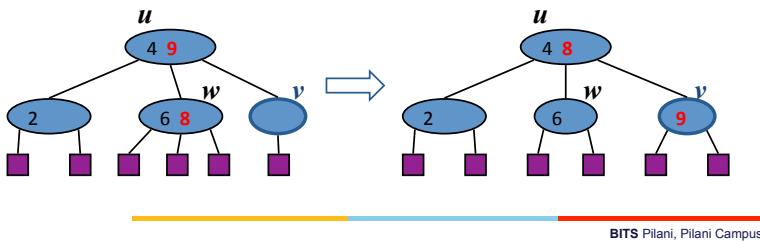
- Deleting an entry from a node **v** may cause an **underflow**, where node **v** becomes a 1-node with one child and no keys
- To handle an underflow at node **v** with parent **u**, we consider two cases
- **Case 1:** the adjacent siblings of **v** are 2-nodes
 - **Fusion operation:** we merge **v** with an adjacent sibling **w** and move an entry from **u** to the merged node **v'**
 - After a fusion, the underflow may propagate to the parent **u**



Underflow and Transfer



- **Case 2:** an immediate sibling w of v is a 3-node or a 4-node
 - **Transfer operation:**
 1. we move a child of w to v
 2. we move an item from u to v
 3. we move an item from w to u
 - After a transfer, no underflow occurs



BITS Pilani, Pilani Campus

Analysis of Deletion

- Let T be a $(2,4)$ tree with n items
 - Tree T has $O(\log n)$ height
- In a deletion operation
 - We visit $O(\log n)$ nodes to locate the node from which to delete the entry
 - We handle an underflow with a series of $O(\log n)$ fusions, followed by at most one transfer
 - Each fusion and transfer takes $O(1)$ time
- Thus, deleting an item from a $(2,4)$ tree takes $O(\log n)$ time

Hash Tables



Aim

To develop a search procedure with running time $O(1)$.

Array is an immediate answer.

Example: Consider a sorted sequence stored in an array.

11	22	28	33	44	49	55	58	66	77	79	88
----	----	----	----	----	----	----	----	----	----	----	----

Focus on finding 44.

Binary search does it in $O(\log n)$ time

Can we retrieve faster than binary search?

BITS Pilani, Pilani Campus



Hash Tables

If 44 was stored in $A[44]$, search would be done in one step.

- Problem is if keys are very large.

Lot of cells will be empty and so inefficient in terms of space requirement.

Answer : Hash Tables

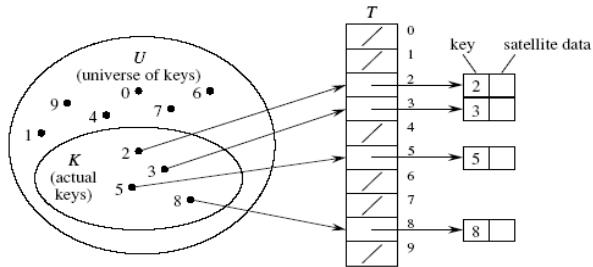
- Uses an array of size proportional to the number of keys used.
- Instead of using key as an array index, array index is computed from the key.

BITS Pilani, Pilani Campus



Array as table

- It is also called **Direct-address Hash Table**.
- Each **slot**, or position, corresponds to a key in U .
- If there's an element x with key k , then $T[k]$ contains a pointer to x .
- Otherwise, $T[k]$ is empty, represented by NIL.



BITS Pilani, Pilani Campus



Array as table

- Store the records in a huge array where the index corresponds to the key
 - add - **very fast** $O(1)$
 - delete - **very fast** $O(1)$
 - search - **very fast** $O(1)$

BITS Pilani, Pilani Campus

Direct Address table



- Disadvantage** of Direct Address table
 - If the universe U is very large, storing a table of size $|U|$ may be impractical.
 - The set K of keys actually stored may be small relative to U , so that most of the space allocated for table is wasted.

BITS Pilani, Pilani Campus



Hash function

- In direct addressing, an element with key k is stored in slot k .
- With hashing, the **hash function** is used to compute the slot $h(k)$

$$h : U \rightarrow \{0, 1, \dots, m-1\}$$

We say that $h(k)$ is the hash value of k .

- Aim**

To reduce the range of array of indices that needs to be handled

BITS Pilani, Pilani Campus



Compression Maps

- **Division Method**

$$h(k) = k \bmod m$$

- Certain values of m may not be good:
- Good values for m are prime numbers which are not close to exact powers of 2. For example, if you want to store 2000 elements then $m=701$ ($m = \text{hash table length}$) yields a hash function: $h(k) = k \bmod 701$

22

BITS Pilani, Pilani Campus

Hash function

- **Problem:**

Two keys may hash to the same slot – **collision**

Ideal situation – avoid collision altogether

- But it is generally **difficult** to design perfect hash. (e.g. when the potential key space is large)

- **Solution :**

Find effective techniques for resolving collision

24

BITS Pilani, Pilani Campus



Compression Maps

- **MAD Method (Multiply, Add and Divide)**

Define $h(k) = (ak + b) \bmod n$,

where n is a prime number and a and b are chosen randomly so that $a \bmod n \neq 0$.

- This function is chosen in order to eliminate repeated patterns.

23

BITS Pilani, Pilani Campus



Collision Resolving

- Collision resolution by chaining

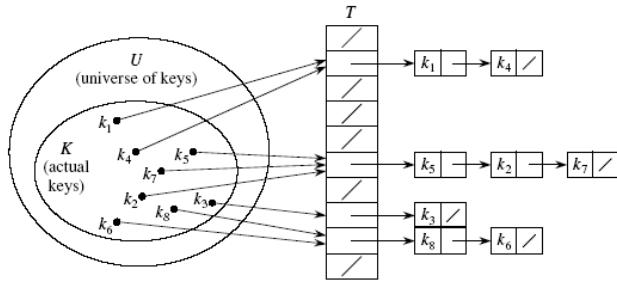
In chaining, we put all elements that hash to same slot in a linked list.

- Slot j contains a pointer to the head of the list of all stored elements that hash to j
- If there are no such elements, slot j contains NIL.

25

BITS Pilani, Pilani Campus

Chained Hash Table



26

BITS Pilani, Pilani Campus

Performance Analysis

- **Assume**
 - *Simple uniform hashing.*
 - Any key is equally likely to hash into any of the m slots, independent of where any other key hashes to.
 - $O(1)$ time to compute $h(k)$.
- Time to search for an element with key k depends linearly on the length of the list $T[h(k)]$.
- Expected length of a linked list = load factor = $\alpha = n/m$.
- We consider two cases:
 1. Unsuccessful search
 2. Successful search

28
BITS Pilani, Pilani Campus

Performance Analysis

• Worst Case

All keys are hashed to the same slot, worst cast is $O(n)$ time

• Average Performance

Given a hash table T with m slots that stores n keys,

Define **load factor α** for T as n/m

- average keys per slot.

27
BITS Pilani, Pilani Campus

27
BITS Pilani, Pilani Campus

Unsuccessful Search

• Theorem:

- An unsuccessful search takes expected time $\Theta(1+\alpha)$.

Proof:

- Any key not already in the table is equally likely to hash to any of the m slots.
- To search unsuccessfully for any key k , need to search to the end of the list $T[h(k)]$, whose expected length is α .
- Adding the time to compute the hash function, the total time required is $\Theta(1+\alpha)$.

BITS Pilani, Pilani Campus

Successful Search



- **Theorem:**

- A successful search takes expected time $\Theta(1+\alpha)$.

Proof:

- The probability that a list is searched is proportional to the number of elements it contains.
- Assume that the element being searched for is equally likely to be any of the n elements in the table.
- The number of elements examined during a successful search for an element x is 1 more than the number of elements examined when the sought for element was inserted.
- The expected length of list to which the i th element is added is $(i-1)/m$

BITS Pilani, Pilani Campus



Expected Cost of a Successful Search

Proof contd.

Therefore, the expected no. of elements examined in a successful search is $\left[\frac{1}{n} \sum_{i=1}^n \left(1 + \frac{i-1}{m} \right) \right]$

It can be shown that this is of order $\Theta(1+\alpha)$.

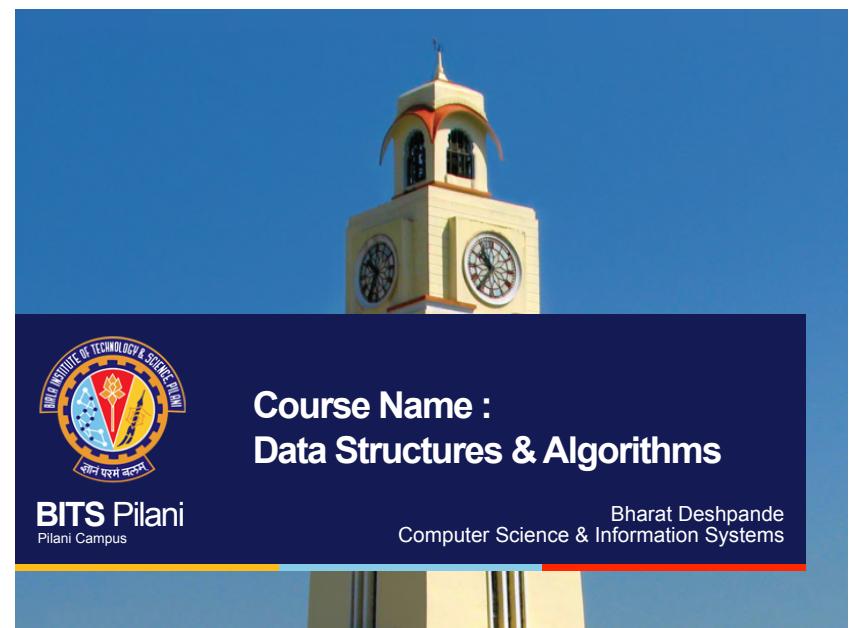
BITS Pilani, Pilani Campus

Expected Cost – Interpretation



- If $n = O(m)$, then $\alpha = n/m = O(m)/m = O(1)$.
⇒ Searching takes constant time on average.
- Insertion is $O(1)$ in the worst case.
- Deletion takes $O(1)$ worst-case time when lists are doubly linked.
- Hence, all dictionary operations take $O(1)$ time on average with hash tables with chaining.

BITS Pilani, Pilani Campus



The image shows the iconic yellow clock tower of BITS Pilani against a clear blue sky. Below the tower, the university's logo is displayed, featuring a circular emblem with a torch and the text "BITS PILANI". To the right of the logo, the course name is listed. At the bottom right, the name of the professor is mentioned.

Course Name :
Data Structures & Algorithms

Bharat Deshpande
Computer Science & Information Systems



Hash Tables

- Aim**

To develop a search procedure with running time $O(1)$.

Array is an immediate answer.

Example: Consider a sorted sequence stored in an array.

11	22	28	33	44	49	55	58	66	77	79	88
----	----	----	----	----	----	----	----	----	----	----	----

Focus on finding 44.

Binary search does it in $O(\log n)$ time

Can we retrieve faster than binary search?



Hash Tables

If 44 was stored in $A[44]$, search would be done in one step.

- Problem is if keys are very large.

Lot of cells will be empty and so inefficient in terms of space requirement.

- Answer : Hash Tables**

- Uses an array of size proportional to the number of keys used.

- Instead of using key as an array index, array index is computed from the key.

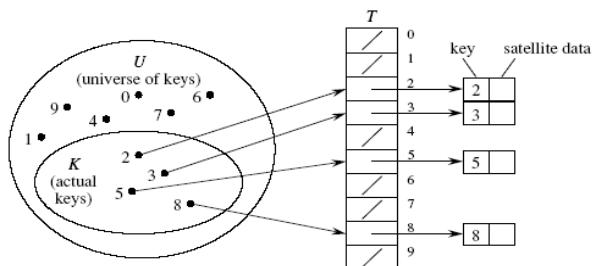
BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

Array as table



- It is also called **Direct-address Hash Table**.
- Each *slot*, or position, corresponds to a key in U .
- If there's an element x with key k , then $T[k]$ contains a pointer to x .
- Otherwise, $T[k]$ is empty, represented by NIL.



BITS Pilani, Pilani Campus



Array as table

- Store the records in a huge array where the index corresponds to the key
 - add - **very fast** $O(1)$
 - delete - **very fast** $O(1)$
 - search - **very fast** $O(1)$

BITS Pilani, Pilani Campus



Direct Address table

- **Disadvantage** of Direct Address table
 - If the universe U is very large, storing a table of size $|U|$ may be impractical.
 - The set K of keys actually stored may be small relative to U , so that most of the space allocated for table is wasted.



BITS Pilani, Pilani Campus

Compression Maps



- **Division Method**

$$h(k) = k \bmod m$$

- Certain values of m may not be good:
- Good values for m are prime numbers which are not close to exact powers of 2. For example, if you want to store 2000 elements then $m=701$ ($m = \text{hash table length}$) yields a hash function: $h(k) = k \bmod 701$

8
BITS Pilani, Pilani Campus

Hash function

- In direct addressing, an element with key k is stored in slot k .
- With hashing, the **hash function** is used to compute the slot $h(k)$
$$h : U \rightarrow \{0, 1, \dots, m-1\}$$

We say that $h(k)$ is the hash value of k .

- **Aim**

To reduce the range of array of indices that needs to be handled

BITS Pilani, Pilani Campus

Compression Maps



- **MAD Method** (Multiply, Add and Divide)

Define $h(k) = (ak + b) \bmod n$,
where n is a prime number and a and b are chosen randomly so that $a \bmod n \neq 0$.

- This function is chosen in order to eliminate repeated patterns.

9
BITS Pilani, Pilani Campus

Hash function



- **Problem:**

Two keys may hash to the same slot – **collision**

Ideal situation – avoid collision altogether

- But it is generally **difficult** to design perfect hash. (e.g. when the potential key space is large)

- **Solution :**

Find effective techniques for resolving collision

Collision Resolving



- Collision resolution by chaining

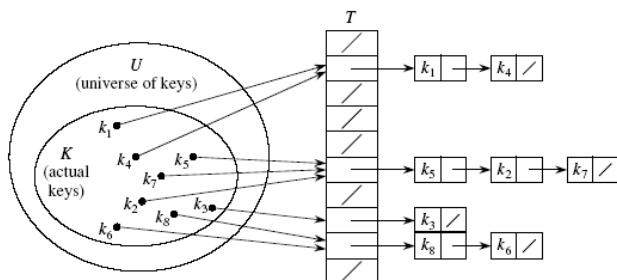
In chaining, we put all elements that hash to same slot in a linked list.

- Slot j contains a pointer to the head of the list of all stored elements that hash to j
- If there are no such elements, slot j contains NIL.

10

11
BITS Pilani, Pilani Campus

Chained Hash Table



12
BITS Pilani, Pilani Campus

Performance Analysis



- **Worst Case**

All keys are hashed to the same slot, worst cast is $O(n)$ time

- **Average Performance**

Given a hash table T with m slots that stores n keys,

Define **load factor α** for T as n/m

- average keys per slot.

13
BITS Pilani, Pilani Campus



- **Assume**
 - *Simple uniform hashing.*
 - Any key is equally likely to hash into any of the m slots, independent of where any other key hashes to.
 - *$O(1)$ time to compute $h(k)$.*
- Time to search for an element with key k depends linearly on the length of the list $T[h(k)]$.
- Expected length of a linked list = load factor = $\alpha = n/m$.
- We consider two cases:
 1. Unsuccessful search
 2. Successful search

14

BITS Pilani, Pilani Campus

Successful Search



- **Theorem:**
 - A successful search takes expected time $\Theta(1+\alpha)$.
- **Proof:**
 - The probability that a list is searched is proportional to the number of elements it contains.
 - Assume that the element being searched for is equally likely to be any of the n elements in the table.
 - **The number of elements examined during a successful search for an element x is 1 more than the number of elements examined when the sought for element was inserted.**
 - The expected length of list to which the i th element is added is $(i-1)/m$

BITS Pilani, Pilani Campus

Unsuccessful Search



- **Theorem:**
 - An unsuccessful search takes expected time $\Theta(1+\alpha)$.
- **Proof:**
 - Any key not already in the table is equally likely to hash to any of the m slots.
 - To search unsuccessfully for any key k , need to search to the end of the list $T[h(k)]$, whose expected length is α .
 - Adding the time to compute the hash function, the total time required is $\Theta(1+\alpha)$.

BITS Pilani, Pilani Campus

Expected Cost of a Successful Search



Proof contd.

Therefore, the expected no. of elements examined in a successful search is $\left[\frac{1}{n} \sum_{i=1}^n \left(1 + \frac{i-1}{m} \right) \right]$

It can be shown that this is of order $\Theta(1+\alpha)$.

BITS Pilani, Pilani Campus



Expected Cost – Interpretation



- If $n = O(m)$, then $\alpha = n/m = O(m)/m = O(1)$.
⇒ Searching takes constant time on average.
- Insertion is $O(1)$ in the worst case.
- Deletion takes $O(1)$ worst-case time when lists are doubly linked.
- Hence, all dictionary operations take $O(1)$ time on average with hash tables with chaining.

BITS Pilani, Pilani Campus

Probe Sequence



- Sequence of slots examined during a key search constitutes a **probe sequence**.
- Probe sequence must be a permutation of the slot numbers.
 - We examine every slot in the table, if we have to.
 - We don't examine any slot more than once.
- The hash function is extended to:

$$h : U \times \{0, 1, \dots, m-1\} \rightarrow \{0, 1, \dots, m-1\}$$

probe number slot number
- $\langle h(k,0), h(k,1), \dots, h(k,m-1) \rangle$ should be a permutation of $\langle 0, 1, \dots, m-1 \rangle$.

BITS Pilani, Pilani Campus

Open Addressing

An alternative to chaining for handling collisions.

- **Idea:**
 - Store all keys in the hash table itself.
 - Each slot contains either a key or NIL.
 - To **search** for key k :
 - Examine slot $h(k)$. Examining a slot is known as a **probe**.
 - If slot $h(k)$ contains key k , the search is successful. If the slot contains NIL, the search is unsuccessful.
 - There's a third possibility: **slot $h(k)$ contains a key that is not k** .
 - Compute the index of some other slot, based on k and which probe we are on.
 - Keep probing until we either find key k or we find a slot holding NIL.
 - **Advantages:** Avoids pointers; so can use a larger table.

BITS Pilani, Pilani Campus

Insert Operation



Act as though we were searching, and insert at the first NIL slot found

Hash-Insert(T, k)

1. $i \leftarrow 0$
2. **repeat** $j \leftarrow h(k, i)$
3. **if** $T[j] = \text{NIL}$
4. **then** $T[j] \leftarrow k$
5. **return** j
6. **else** $i \leftarrow i + 1$
7. **until** $i = m$
8. **error** "hash table overflow"

BITS Pilani, Pilani Campus



Search Operation

The search algorithm for key k probes the same sequence of slots that the insertion algorithm examined when k was inserted.

Hash-Search (T, k)

1. $i \leftarrow 0$
2. **repeat** $j \leftarrow h(k, i)$
3. **if** $T[j] = k$
4. **then return** j
5. $i \leftarrow i + 1$
6. **until** $T[j] = \text{NIL}$ **or** $i = m$
7. **return** NIL

BITS Pilani, Pilani Campus

Probe Sequences



- The ideal situation is ***uniform hashing***:
 - Generalization of simple uniform hashing.
 - Each key is equally likely to have any of the $m!$ permutations of $\langle 0, 1, \dots, m-1 \rangle$ as its probe sequence.
- It is **hard to implement** true uniform hashing.
 - **Approximate** with techniques that at least guarantee that the probe sequence is a permutation of $\langle 0, 1, \dots, m-1 \rangle$.
- **Three commonly used techniques:**
 - Linear Probing.
 - Quadratic Probing.
 - Double Hashing.
- All guarantee that $\langle h(k,0), h(k,1), \dots, h(k,m-1) \rangle$ is a permutation of $\langle 0, 1, \dots, m-1 \rangle$.
- None of these fulfills assumptions of uniform hashing.
 - Can't produce all $m!$ probe sequences.

BITS Pilani, Pilani Campus



Deletion Operation

- Cannot just turn the slot containing the key we want to delete to contain NIL.
- Doing so might make it impossible to retrieve any key k during whose insertion we had probed this slot & found it occupied.
- Use a special value **DELETED** instead of NIL when marking a slot as empty during deletion.
 - **Search** should treat DELETED as though the slot holds a key that does not match the one being searched for.
 - **Insert** should treat DELETED as though the slot were empty, so that it can be reused.
- **Disadvantage: Search time is no longer dependent on α .**
 - Hence, chaining is more common when keys have to be deleted.

BITS Pilani, Pilani Campus

Linear Probing



- $$h(k, i) = (h'(k) + i) \bmod m.$$

key
Probe number
Auxiliary hash function
- The initial probe determines the entire probe sequence.
 - $T[h'(k)], T[h'(k)+1], \dots, T[m-1], T[0], T[1], \dots, T[h'(k)-1]$
 - Hence, **only m distinct probe sequences** are possible.
 - Easy to Implement
- Suffers from **primary clustering**:
 - Long runs of occupied sequences build up.
 - Long runs tend to get longer, since an empty slot preceded by i full slots gets filled next with probability $(i+1)/m$.
 - Hence, average search and insertion times increase.

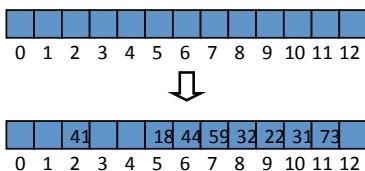
BITS Pilani, Pilani Campus

Linear Probing : Example



Example:

- $h'(x) = x \bmod 13$
- $h(x) = (h'(x) + i) \bmod 13$
- Insert keys 18, 41, 22, 44, 59, 32, 31, 73, in this order



BITS Pilani, Pilani Campus

Double Hashing



- $h(k, i) = (h_1(k) + i h_2(k)) \bmod m$
- Two auxiliary hash functions.
 - h_1 gives the initial probe. h_2 gives the remaining probes.
- Must have $h_2(k)$ relatively prime to m , so that the probe sequence is a full permutation of $\{0, 1, \dots, m-1\}$.
 - Choose m to be a power of 2 and have $h_2(k)$ always return an odd number. Or,
 - Let m be prime, and have $1 < h_2(k) < m$.
- $\Theta(m^2)$ different probe sequences.
 - One for each possible combination of $h_1(k)$ and $h_2(k)$.
 - Close to the ideal uniform hashing.
 - Best method available for open addressing

BITS Pilani, Pilani Campus

Quadratic Probing



$$h(k, i) = (h'(k) + c_1 i + c_2 i^2) \bmod m \quad c_1 \neq c_2$$

key Probe number Auxiliary hash function

- The initial probe position is $T[h'(k)]$, later probe positions are offset by amounts that depend on a quadratic function of the probe number i .
- Must constrain c_1 , c_2 , and m to ensure that we get a full permutation of $\{0, 1, \dots, m-1\}$.
- Can suffer from **secondary clustering**:
 - If two keys have the same initial probe position, then their probe sequences are the same.

27

BITS Pilani, Pilani Campus

Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.

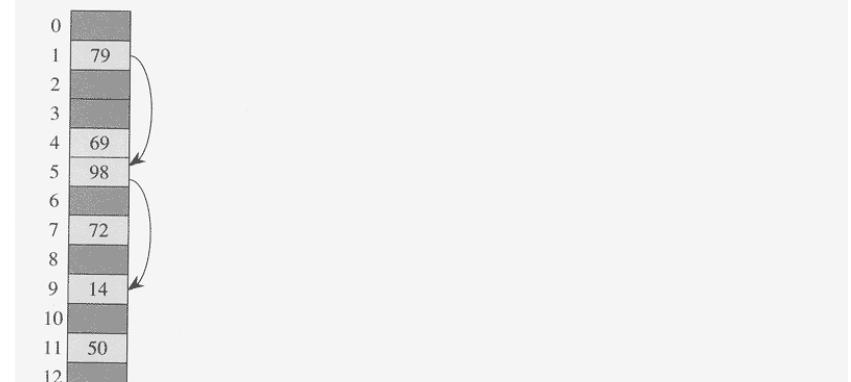


Figure 11.5 Insertion by double hashing. Here we have a hash table of size 13 with $h_1(k) = k \bmod 13$ and $h_2(k) = 1 + (k \bmod 11)$. Since $14 \equiv 1 \pmod{13}$ and $14 \equiv 3 \pmod{11}$, the key 14 is inserted into empty slot 9, after slots 1 and 5 are examined and found to be occupied.

BITS Pilani, Pilani Campus

Performance of Open Addressing



Theorem:

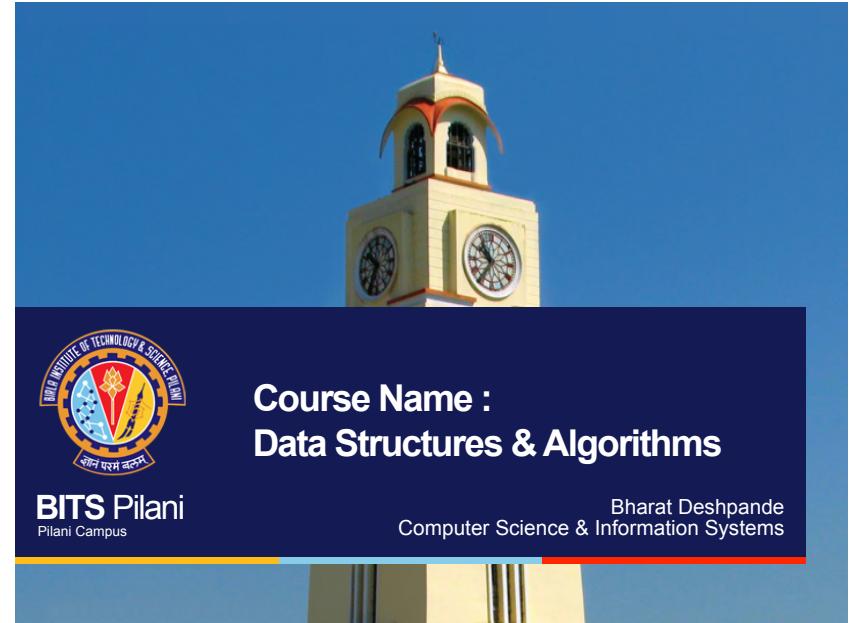
The expected number of probes in an unsuccessful search in an open-address hash table is at most $1/(1-\alpha)$.

Corollary:

Inserting an element into an open-address table takes $\leq 1/(1-\alpha)$ probes on average.

- If α is a constant, search insertion takes $O(1)$ time.

BITS Pilani, Pilani Campus



What is a Graph?



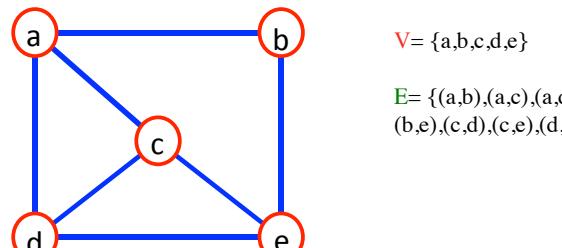
- A graph $G = (V, E)$ is composed of:

V : set of **vertices**

E : set of **edges** connecting the **vertices** in V

- An **edge** $e = (u, v)$ is a pair of **vertices**

- Example:**



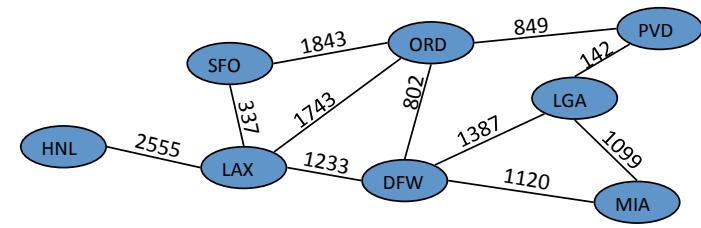
BITS Pilani, Pilani Campus

Graph-Example



- Example:**

- A vertex represents an airport and stores the three-letter airport code
- An edge represents a flight route between two airports and stores the mileage of the route

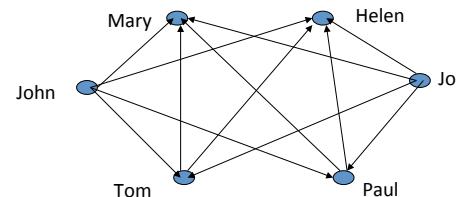


BITS Pilani, Pilani Campus



A “Real-life” Example of a Graph

- $V = \text{set of 6 people: John, Mary, Joe, Helen, Tom, and Paul, of ages 12, 15, 12, 15, 13, and 13, respectively.}$
- $E = \{(x,y) \mid \text{if } x \text{ is younger than } y\}$

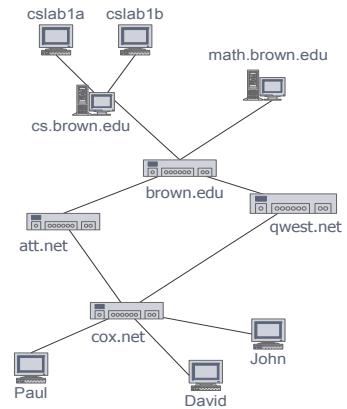


BITS Pilani, Pilani Campus



Applications

- Electronic circuits
 - Printed circuit board
 - Integrated circuit
- Transportation networks
 - Highway network
 - Flight network
- Computer networks
 - Local area network
 - Internet
 - Web
- Databases
 - Entity-relationship diagram

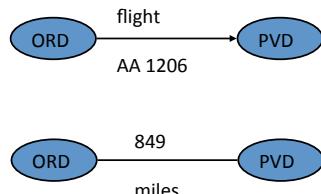


BITS Pilani, Pilani Campus



Edge Types

- **Directed edge**
 - ordered pair of vertices (u,v)
 - first vertex u is the origin
 - second vertex v is the destination
 - e.g., a flight
- **Undirected edge**
 - unordered pair of vertices (u,v)
 - e.g., a flight route
- **Directed graph (Digraph)**
 - all the edges are directed
 - e.g., flight network
- **Undirected graph**
 - all the edges are undirected
 - e.g., route network



BITS Pilani, Pilani Campus



Terminology

- If (v_0, v_1) is an edge in an undirected graph,
 - v_0 and v_1 are **adjacent**
 - The edge (v_0, v_1) is incident on vertices v_0 and v_1
- If $\langle v_0, v_1 \rangle$ is an edge in a directed graph
 - v_0 is **adjacent to** v_1 , and v_1 is **adjacent from** v_0
 - The edge $\langle v_0, v_1 \rangle$ is incident on v_0 and v_1

Terminology:



- **Degree of a Vertex**

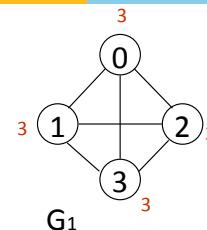
The **degree** of a vertex is the number of edges incident to that vertex

- For directed graph,

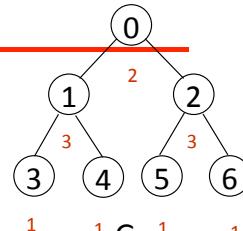
- the **in-degree** of a vertex v is the number of edges that have v as the head
- the **out-degree** of a vertex v is the number of edges that have v as the tail
- if d_i is the degree of a vertex i in a graph G with n vertices and e edges, the number of edges is $e = (\sum_{i=1}^{n-1} d_i) / 2$ **Why?** Since adjacent vertices each count the adjoining edge, it will be counted twice

BITS Pilani, Pilani Campus

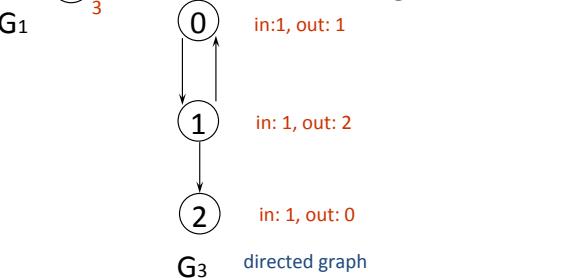
Examples



G_1



G_2



G_3 directed graph

in: 1, out: 1

in: 1, out: 2

in: 1, out: 0

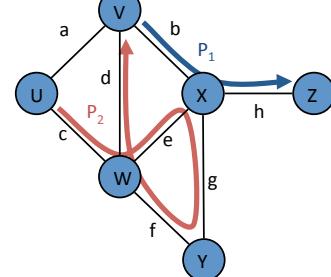
BITS Pilani, Pilani Campus

Terminology (cont.)



- **Path**

- sequence of alternating vertices and edges
- begins with a vertex
- ends with a vertex
- each edge is preceded and followed by its endpoints



BITS Pilani, Pilani Campus

- **Simple path**

- path such that all its vertices and edges are distinct

- **Examples**

- $P_1=(V,b,X,h,Z)$ is a simple path
- $P_2=(U,c,W,e,X,g,Y,f,W,d,V,a)$ is a path that is not simple

Terminology (cont.)



- **Cycle**

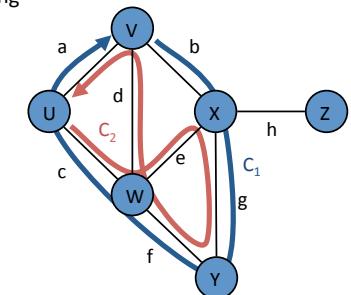
- circular sequence of alternating vertices and edges
- each edge is preceded and followed by its endpoints

- **Simple cycle**

- cycle such that all its vertices and edges are distinct

- **Examples**

- $C_1=(V,b,X,g,Y,f,W,c,U,a)$ is a simple cycle
- $C_2=(U,c,W,e,X,g,Y,f,W,d,V,a)$ is a cycle that is not simple



BITS Pilani, Pilani Campus

Properties



Property 1

$$\sum_v \deg(v) = 2m$$

Proof:
each edge is counted twice

Property 2

In an undirected graph with no self-loops and no multiple edges
 $m \leq n(n - 1)/2$

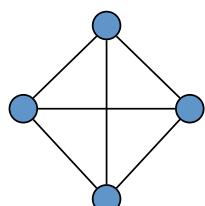
Proof:
each vertex has degree at most $(n - 1)$

Notation

n	number of vertices
m	number of edges
$\deg(v)$	degree of vertex v

Example

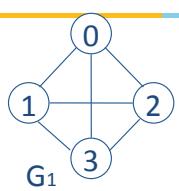
- $n = 4$
- $m = 6$
- $\deg(v) = 3$



Graphs

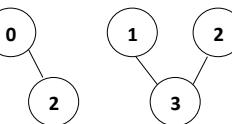
12 BITS Pilani, Pilani Campus

Subgraph Examples

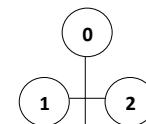


(i)

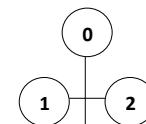
Some of the subgraph of G_1



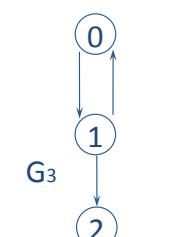
(ii)



(iii)



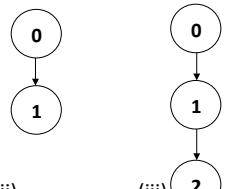
(iv)



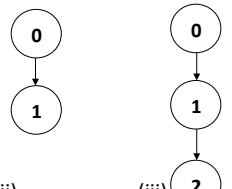
(i)

Some of the subgraph of G_3

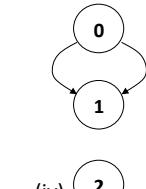
(ii)



(iii)



(iv)



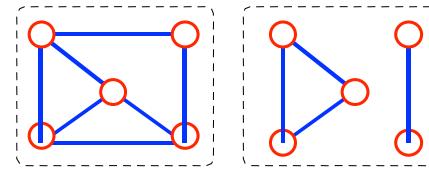
(v)

BITS Pilani, Pilani Campus

Even More Terminology



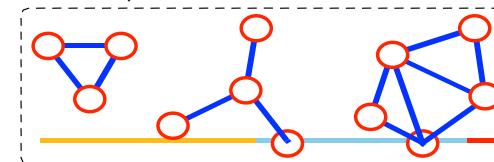
- **connected graph:** any two vertices are connected by some path



connected

not connected

- **subgraph:** subset of vertices and edges forming a graph
- **connected component:** maximal connected subgraph. E.g., the graph below has 3 connected components.



BITS Pilani, Pilani Campus

Trees



- Tree is a special case of a graph. Each node has zero or more child nodes, which are below it in the tree.
- A tree is a connected acyclic graph
- Already seen.
- **Forest** - collection of trees

BITS Pilani, Pilani Campus

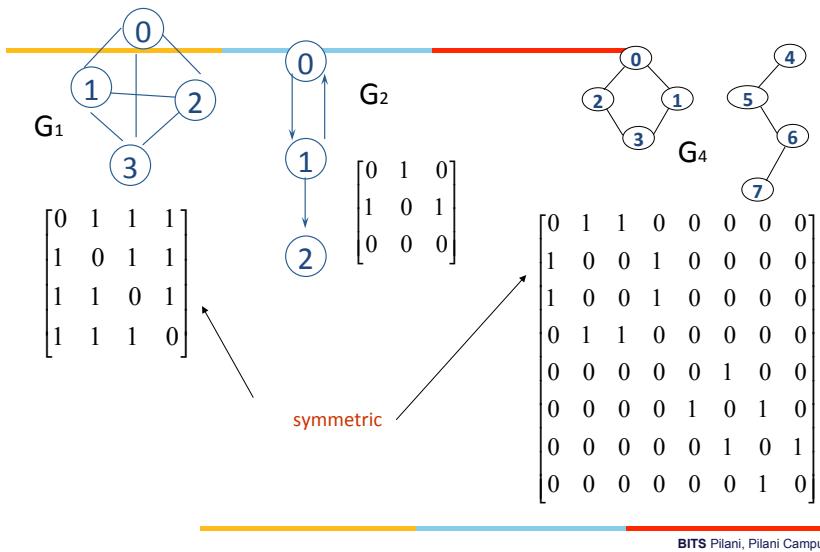
Graph Representations



- Adjacency Matrix
- Adjacency Lists

BITS Pilani, Pilani Campus

Examples for Adjacency Matrix



BITS Pilani, Pilani Campus

Adjacency Matrix



- Let $G=(V,E)$ be a graph with n vertices.
- The **adjacency matrix** of G is a two-dimensional n by n array, say adj_mat
- If the edge (v_i, v_j) is in $E(G)$, $\text{adj_mat}[i][j]=1$
- If there is no such edge in $E(G)$, $\text{adj_mat}[i][j]=0$
- The adjacency matrix for an undirected graph is **symmetric**; the adjacency matrix for a digraph need not be symmetric

BITS Pilani, Pilani Campus

Merits of Adjacency Matrix



- From the adjacency matrix, to determine the connection of vertices is easy
- The degree of a vertex is $\sum_{j=0}^{n-1} \text{adj_mat}[i][j]$
- For a digraph the row sum is the *out_degree*, while the column sum is the *in_degree*

$$\text{ind}(vi) = \sum_{j=0}^{n-1} A[j,i] \quad \text{outd}(vi) = \sum_{j=0}^{n-1} A[i,j]$$

Cons : No matter how few edges the graph has, the matrix takes $O(n^2)$ in memory

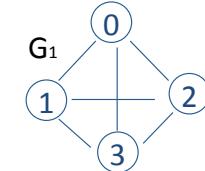
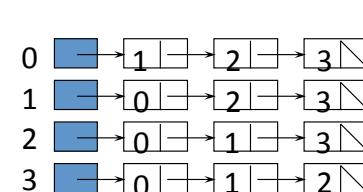
BITS Pilani, Pilani Campus

Adjacency Lists Representation

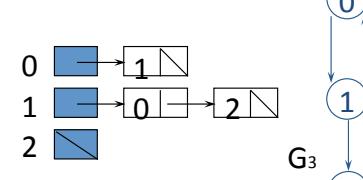
- A graph of n nodes is represented by a one-dimensional array L of linked lists, where
 - $L[i]$ is the linked list containing all the nodes adjacent from node i .
 - The nodes in the list $L[i]$ are in no particular order



BITS Pilani, Pilani Campus



An undirected graph with n vertices and e edges => n head nodes and $2e$ list nodes



BITS Pilani, Pilani Campus

Pros and Cons of Adjacency Lists



- Pros**
 - Saves on space (memory): the representation takes as many memory words as there are nodes and edge.
- Cons**
 - It can take up to $O(n)$ time to determine if a pair of nodes (i,j) is an edge: one would have to search the linked list $L[i]$, which takes time proportional to the length of $L[i]$.

BITS Pilani, Pilani Campus



Graph Traversal

- Problem:** Search for a certain node or traverse all nodes in the graph.
- Depth First Search**
 - Once a possible path is found, continue the search until the end of the path
- Breadth First Search**
 - Start several paths at a time, and advance in each one step at a time

BITS Pilani, Pilani Campus

Exploring a Labyrinth Without Getting Lost

- A **depth-first search (DFS)** in an undirected graph G is like wandering in a labyrinth with a string and a can of red paint without getting lost.
- We start at vertex s , tying the end of our string to the point and painting s "visited". Next we label s as our current vertex called u .
- Now we travel along an arbitrary edge (u, v) .
- If edge (u, v) leads us to an already visited vertex v we return to u .
- If vertex v is unvisited, we unroll our string and move to v , paint v "visited", set v as our current vertex, and repeat the previous steps.
- The process terminates when our backtracking leads us back to the start index s , and there are no more unexplored edges incident on s .

BITS Pilani, Pilani Campus

Depth-First Search

Algorithm DFS(v): Input: A vertex v in a graph

Output: A labeling of the edges as "discovery" edges and "backedges"
for each edge e incident on v **do**
 if edge e is unexplored **then** let w be the other endpoint of e
 if vertex w is unexplored **then** label e as a **discovery edge**
 recursively call **DFS(w)**
 else label e as a **backedge**

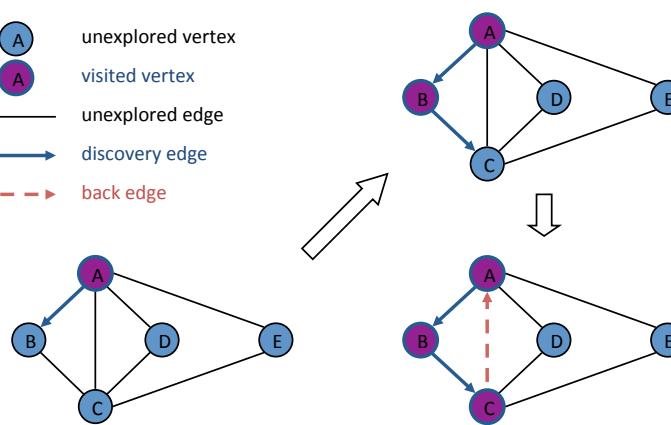
Property 1

$\text{DFS}(G, v)$ visits all the vertices and edges in the connected component of v

BITS Pilani, Pilani Campus

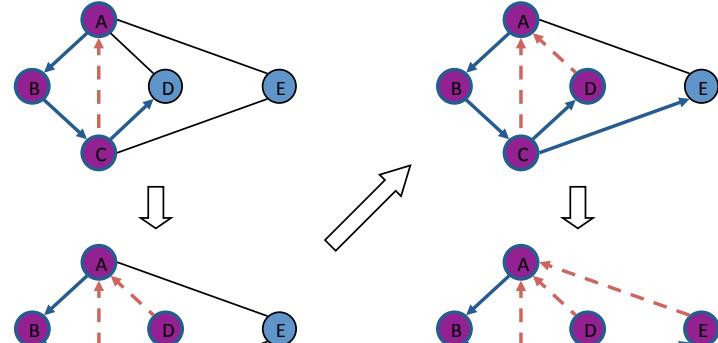
Example

- unexplored vertex
- visited vertex
- unexplored edge
- discovery edge
- back edge



BITS Pilani, Pilani Campus

Example (cont.)



BITS Pilani, Pilani Campus



Analysis of DFS

- Setting/getting a vertex/edge label takes $O(1)$ time
- Each vertex is labeled twice
 - once as **UNEXPLORED**
 - once as **VISITED**
- Each edge is labeled twice
 - once as **UNEXPLORED**
 - once as **DISCOVERY** or **BACK**
- Method `incidentEdges` is called once for each vertex
- DFS runs in $O(n + m)$ time provided the graph is represented by the adjacency list structure

BITS Pilani, Pilani Campus



Path Finding

- We can specialize the DFS algorithm to **find a path between two given vertices u and z** using the template method pattern
- We call **$DFS(G, u)$** with u as the start vertex
- We use a stack S to keep track of the path between the start vertex and the current vertex
- As soon as destination vertex z is encountered, we return the path as the contents of the stack

BITS Pilani, Pilani Campus



Cycle Finding

- We can specialize the DFS algorithm to find a simple cycle using the template method pattern
- We use a stack S to keep track of the path between the start vertex and the current vertex
- As soon as a back edge (v, w) is encountered, we return the cycle as the portion of the stack from the top to vertex w

BITS Pilani, Pilani Campus



Breadth-First Search

- Instead of going as far as possible, BFS tries to search all paths.
- BFS makes use of a queue to store visited (but not dead) vertices, expanding the path from the earliest visited vertices
- The starting vertex s has level 0, and, as in **DFS**, defines that point as an “anchor.”
- In the first round, the string is unrolled the length of one edge, and all of the nodes that are only one edge away from the anchor are visited.
- These edges are placed into level 1
- In the second round, all the new edges that can be reached by unrolling the string 2 edges are visited and placed in level 2.
- This continues until every vertex has been assigned a level.
- The label of any vertex v corresponds to the length of the shortest path from s to v .

BITS Pilani, Pilani Campus

BFS Pseudo-Code



Algorithm BFS(s): Input: A vertex s in a graph

Output: A labeling of the edges as “discovery” edges and “cross edges”

initialize container L_0 to contain vertex s

$i \leftarrow 0$

while L_i is not empty do

 create container L_{i+1} to initially be empty

 for each vertex v in L_i do

 if edge e incident on v do

 let w be the other endpoint of e

 if vertex w is unexplored then

 label e as a **discovery edge**

 insert w into L_{i+1}

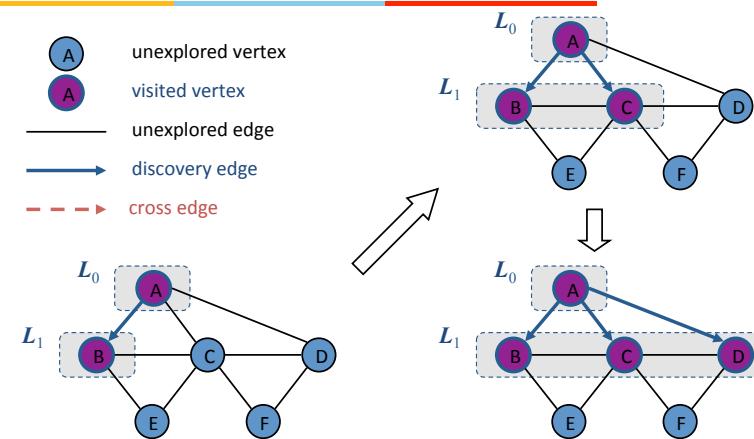
 else label e as a **cross edge**

$i \leftarrow i + 1$

BITs Pilani, Pilani Campus

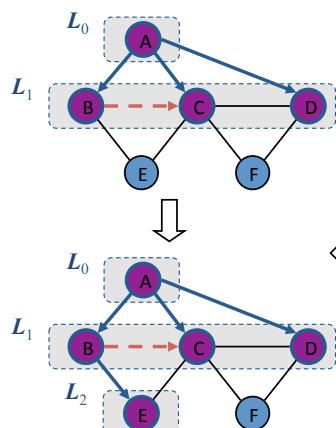
Example

- unexplored vertex
- visited vertex
- unexplored edge
- discovery edge
- cross edge



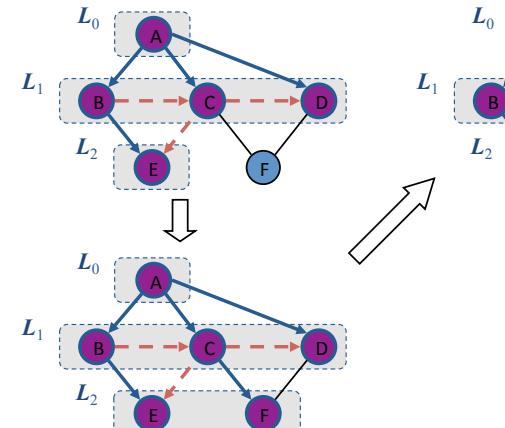
BITs Pilani, Pilani Campus

Example (cont.)



BITs Pilani, Pilani Campus

Example (cont.)



BITs Pilani, Pilani Campus



Analysis

- Setting/getting a vertex/edge label takes $O(1)$ time
- Each vertex is labeled twice
 - once as UNEXPLORED
 - once as VISITED
- Each edge is labeled twice
 - once as UNEXPLORED
 - once as DISCOVERY or CROSS
- Each vertex is inserted once into a sequence L_i
- Method incidentEdges is called once for each vertex
- BFS runs in $O(n + m)$ time provided the graph is represented by the adjacency list structure



Applications

We can specialize the BFS traversal of a graph G to solve the following problems in $O(n + m)$ time

- Compute the connected components of G
- Find a simple cycle in G , or report that G is a forest
- Given two vertices of G , find a path in G between them with the minimum number of edges, or report that no such path exists

BITS Pilani, Pilani Campus

BITS Pilani, Pilani Campus

BITS Pilani
Pilani Campus

Course Name :
Data Structures & Algorithms

Bharat Deshpande
Computer Science & Information Systems



Graph Traversal

- **Problem:** Search for a certain node or traverse all nodes in the graph.
- **Depth First Search**
 - Once a possible path is found, continue the search until the end of the path
- **Breadth First Search**
 - Start several paths at a time, and advance in each one step at a time

BITS Pilani, Pilani Campus

Exploring a Labyrinth Without Getting Lost

- A **depth-first search (DFS)** in an undirected graph G is like wandering in a labyrinth with a string and a can of red paint without getting lost.
- We start at vertex s , tying the end of our string to the point and painting s "visited". Next we label s as our current vertex called u .
- Now we travel along an arbitrary edge (u, v) .
- If edge (u, v) leads us to an already visited vertex v we return to u .
- If vertex v is unvisited, we unroll our string and move to v , paint v "visited", set v as our current vertex, and repeat the previous steps.
- The process terminates when our backtracking leads us back to the start index s , and there are no more unexplored edges incident on s .

BITS Pilani, Pilani Campus

Depth-First Search

Algorithm DFS(v): Input: A vertex v in a graph

Output: A labeling of the edges as "discovery" edges and "backedges"
for each edge e incident on v **do**
 if edge e is unexplored **then** let w be the other endpoint of e
 if vertex w is unexplored **then** label e as a **discovery edge**
 recursively call **DFS(w)**
 else label e as a **backedge**

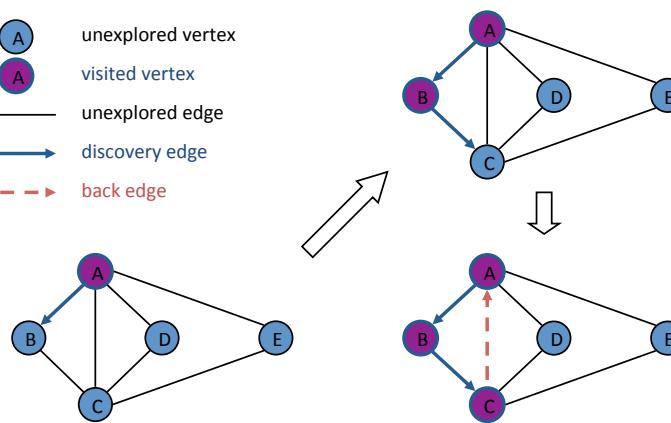
Property 1

$\text{DFS}(G, v)$ visits all the vertices and edges in the connected component of v

BITS Pilani, Pilani Campus

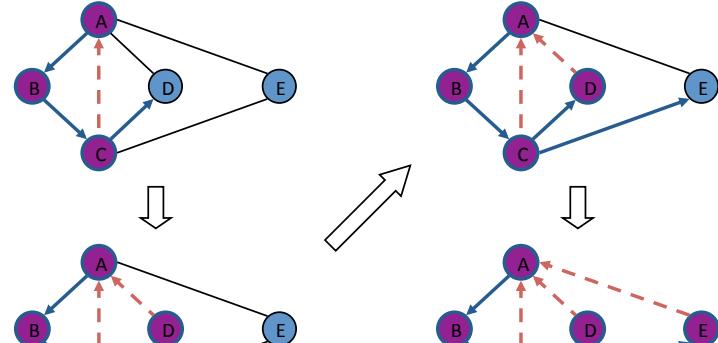
Example

- unexplored vertex
- visited vertex
- unexplored edge
- discovery edge
- back edge



BITS Pilani, Pilani Campus

Example (cont.)



BITS Pilani, Pilani Campus



Analysis of DFS

- Setting/getting a vertex/edge label takes $O(1)$ time
- Each vertex is labeled twice
 - once as **UNEXPLORED**
 - once as **VISITED**
- Each edge is labeled twice
 - once as **UNEXPLORED**
 - once as **DISCOVERY** or **BACK**
- Method `incidentEdges` is called once for each vertex
- DFS runs in $O(n + m)$ time provided the graph is represented by the adjacency list structure

BITS Pilani, Pilani Campus



Path Finding

- We can specialize the DFS algorithm to **find a path between two given vertices u and z** using the template method pattern
- We call **$DFS(G, u)$** with u as the start vertex
- We use a stack S to keep track of the path between the start vertex and the current vertex
- As soon as destination vertex z is encountered, we return the path as the contents of the stack

BITS Pilani, Pilani Campus



Cycle Finding

- We can specialize the DFS algorithm to find a simple cycle using the template method pattern
- We use a stack S to keep track of the path between the start vertex and the current vertex
- As soon as a back edge (v, w) is encountered, we return the cycle as the portion of the stack from the top to vertex w

BITS Pilani, Pilani Campus



Breadth-First Search

- Instead of going as far as possible, BFS tries to search all paths.
- BFS makes use of a queue to store visited (but not dead) vertices, expanding the path from the earliest visited vertices
- The starting vertex s has level 0, and, as in **DFS**, defines that point as an “anchor.”
- In the first round, the string is unrolled the length of one edge, and all of the nodes that are only one edge away from the anchor are visited.
- These edges are placed into level 1
- In the second round, all the new edges that can be reached by unrolling the string 2 edges are visited and placed in level 2.
- This continues until every vertex has been assigned a level.
- The label of any vertex v corresponds to the length of the shortest path from s to v .

BITS Pilani, Pilani Campus

BFS Pseudo-Code



Algorithm BFS(s): Input: A vertex s in a graph

Output: A labeling of the edges as “discovery” edges and “cross edges”

initialize container L_0 to contain vertex s

$i \leftarrow 0$

while L_i is not empty do

 create container L_{i+1} to initially be empty

 for each vertex v in L_i do

 if edge e incident on v do

 let w be the other endpoint of e

 if vertex w is unexplored then

 label e as a **discovery edge**

 insert w into L_{i+1}

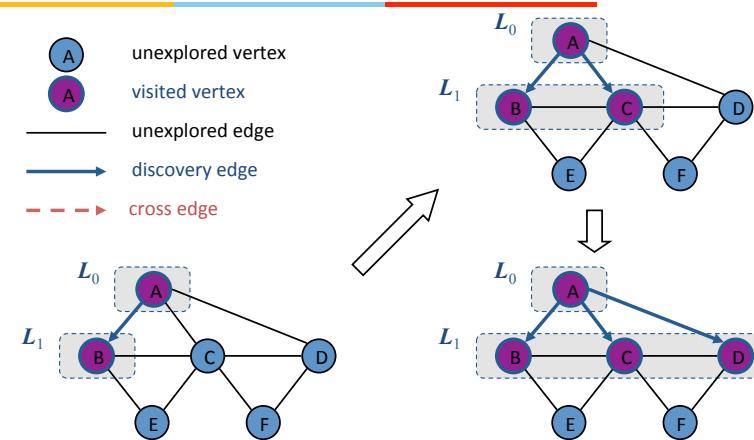
 else label e as a **cross edge**

$i \leftarrow i + 1$

BITs Pilani, Pilani Campus

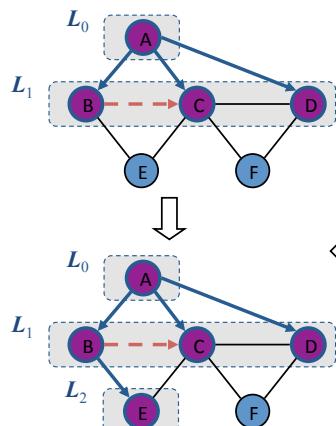
Example

- A unexplored vertex
- A visited vertex
- unexplored edge
- discovery edge
- - - cross edge



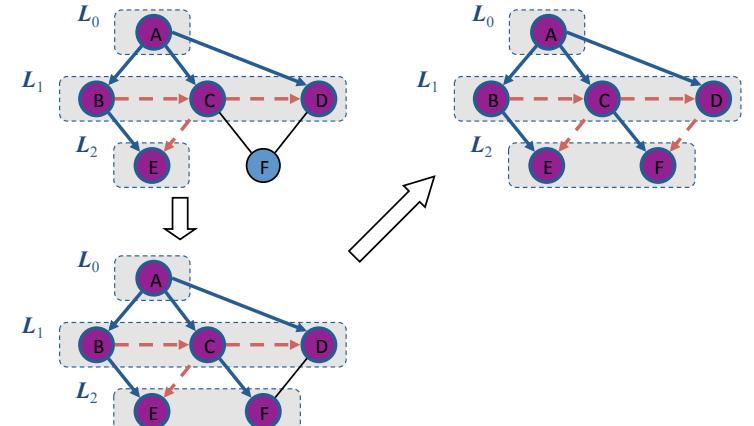
BITs Pilani, Pilani Campus

Example (cont.)



BITs Pilani, Pilani Campus

Example (cont.)



BITs Pilani, Pilani Campus

Analysis

- Setting/getting a vertex/edge label takes $O(1)$ time
- Each vertex is labeled twice
 - once as UNEXPLORED
 - once as VISITED
- Each edge is labeled twice
 - once as UNEXPLORED
 - once as DISCOVERY or CROSS
- Each vertex is inserted once into a sequence L_i
- Method incidentEdges is called once for each vertex
- BFS runs in $O(n + m)$ time provided the graph is represented by the adjacency list structure

BITS Pilani, Pilani Campus

Applications

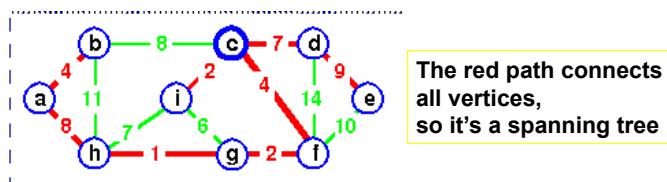
We can specialize the BFS traversal of a graph G to solve the following problems in $O(n + m)$ time

- Compute the connected components of G
- Find a simple cycle in G , or report that G is a forest
- Given two vertices of G , find a path in G between them with the minimum number of edges, or report that no such path exists

BITS Pilani, Pilani Campus

Graphs - Definitions

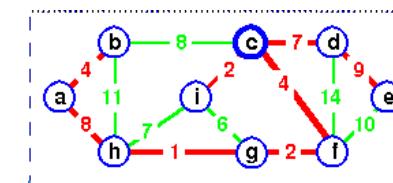
- Spanning Tree**
 - A **spanning tree** is a set of $|V|-1$ edges that connect all the vertices of a graph



BITS Pilani, Pilani Campus

Graphs - Definitions

- Minimum Spanning Tree**
 - Generally there is more than one spanning tree
 - If a cost c_{ij} is associated with edge $e_{ij} = (v_i, v_j)$ then the **minimum spanning tree** is the set of edges E_{span} such that $C = \sum (c_{ij} | \forall e_{ij} \in E_{\text{span}})$ is a minimum



BITS Pilani, Pilani Campus

Growing a MST



- The procedure manages a set A that is always a subset of some MST.
 - At each step, an edge (u, v) is added to A such that $A \cup \{(u, v)\}$ is also a subset of a MST.
 - Edge (u, v) is called a **safe edge**.
 - GENERIC-MST(G)**
- ```

 $A \leftarrow \emptyset$
while A does not form a spanning tree
 do find an edge (u, v) that is safe for A
 $A \leftarrow A \cup \{(u, v)\}$
return A

```
- We design two MST algorithms, each of which uses a specific rule to determine the safe edge to be added.

BITS Pilani, Pilani Campus

## Partition Property

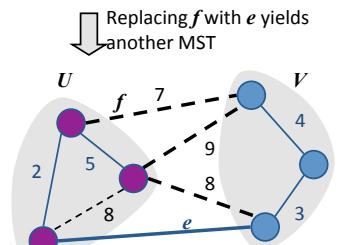
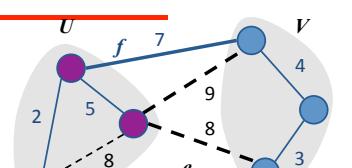


### Partition Property:

- Consider a partition of the vertices of  $G$  into subsets  $U$  and  $V$
- Let  $e$  be an edge of minimum weight across the partition
- There is a minimum spanning tree of  $G$  containing edge  $e$

### Proof:

- Let  $T$  be an MST of  $G$
- If  $T$  does not contain  $e$ , consider the cycle  $C$  formed by  $e$  with  $T$  and let  $f$  be an edge of  $C$  across the partition
- By the cycle property,  $\text{weight}(f) \leq \text{weight}(e)$
- Thus,  $\text{weight}(f) = \text{weight}(e)$
- We obtain another MST by replacing  $f$  with  $e$



BITS Pilani, Pilani Campus

## Cycle Property

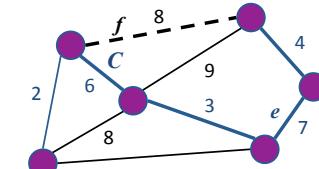
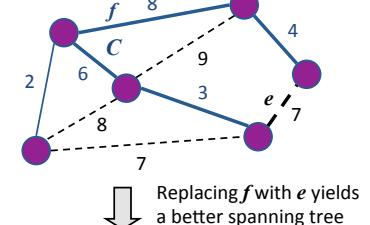


### Cycle Property:

- Let  $T$  be a minimum spanning tree of a weighted graph  $G$
- Let  $e$  be an edge of  $G$  that is not in  $T$  and  $C$  let be the cycle formed by  $e$  with  $T$
- For every edge  $f$  of  $C$ ,  $\text{weight}(f) \leq \text{weight}(e)$

### Proof:

- By contradiction
- If  $\text{weight}(f) > \text{weight}(e)$  we can get a spanning tree of smaller weight by replacing  $e$  with  $f$



BITS Pilani, Pilani Campus

## Prim-Jarnik's Algorithm

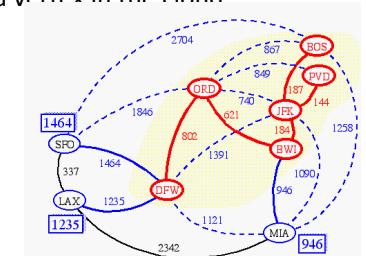


- Initially discovered in 1930 by Vojtěch Jarník, then rediscovered in 1957 by Robert C. Prim

- We pick an arbitrary vertex  $s$  and we grow the MST as a cloud of vertices, starting from  $s$
- We store with each vertex  $v$  a label  $d(v) =$  the smallest weight of an edge connecting  $v$  to a vertex in the cloud

### At each step:

- We add to the cloud the vertex  $u$  outside the cloud with the smallest distance label
- We update the labels of the vertices adjacent to  $u$



BITS Pilani, Pilani Campus

## Prim's Algorithm – Pseudo Code



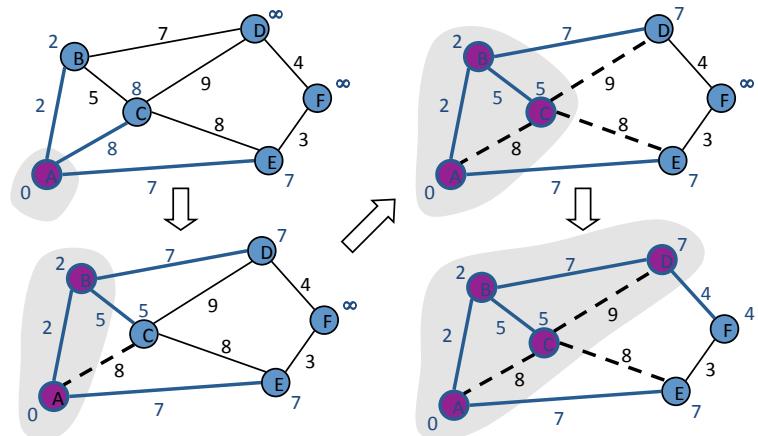
```

MST-PRIM(G, r)
Q ← V[G] // Q – priority queue initialize
for each u ∈ Q
 do d[u] ← ∞
D[r] = 0
Π[r] ← NIL
while Q ≠ φ
 do u ← EXTRACT-MIN(Q) The running time is $O(m \log n)$
 for each v ∈ Adj[u]
 do if v ∈ Q and w(u, v) < d[v]
 then Π[v] ← u
 d[v] ← w(u,v)

```

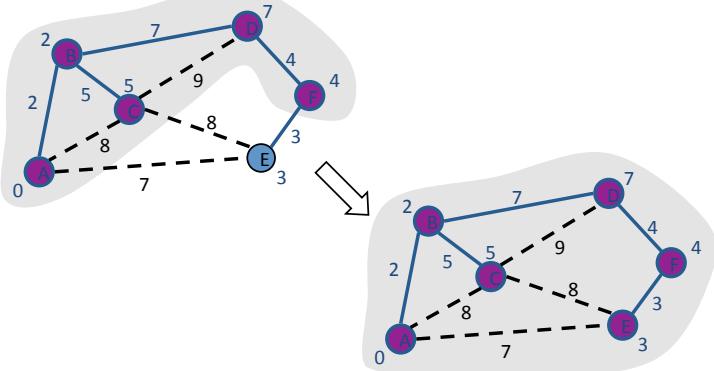
BITS Pilani, Pilani Campus

## Example



BITS Pilani, Pilani Campus

## Example (contd.)



BITS Pilani, Pilani Campus

## Kruskal's Algorithm



- Created in 1957 by Joseph Kruskal
- Algorithm**
- Starts by choosing an edge in the graph with minimum weight.
- Successively add edges with minimum weight that do not form a cycle with those edges already chosen.
- Terminates after  $n-1$  edges have been selected ( $n$  is the number of vertices)

### Note

In Kruskal's algorithm the set A is a forest.

The safe edge added to A is always a least-weight edge in the graph that connects two distinct components.

BITS Pilani, Pilani Campus

## Kruskal's algorithm – Pseudo code



### Algorithm Kruskal( $G$ )

**Input:** A weighted graph  $G$ .

**Output:** An MST  $T$  for  $G$ .

Let  $P$  be a partition of the vertices of  $G$ , where each vertex forms a separate set.

Let  $Q$  be a priority queue storing the edges of  $G$ , sorted by their weights

Let  $T$  be an initially-empty tree

**while**  $Q$  is not empty **do**

$(u,v) \leftarrow Q.\text{removeMinElement}()$

**if**  $P.\text{find}(u) \neq P.\text{find}(v)$  **then**

        Add  $(u,v)$  to  $T$

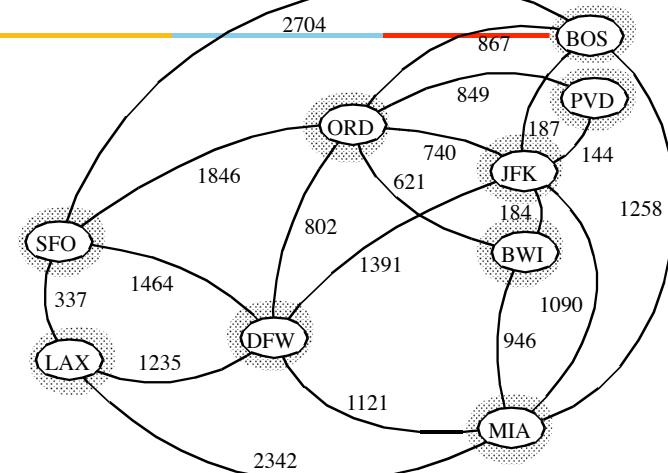
$P.\text{union}(u,v)$

**return**  $T$

**Running time:**  $O(m \log m)$

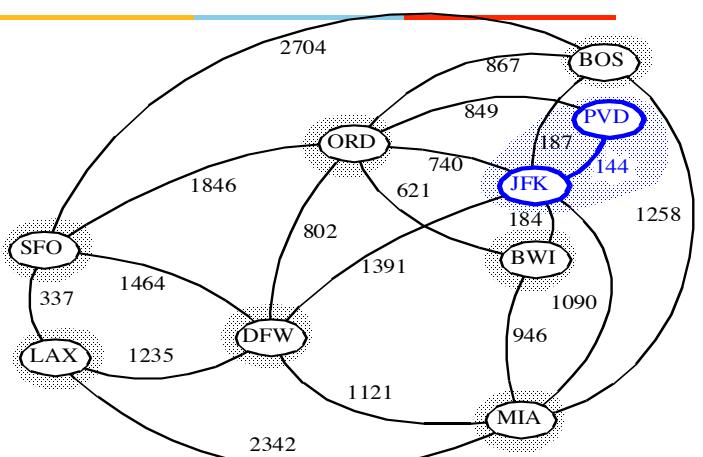
BITS Pilani, Pilani Campus

## Kruskal Example



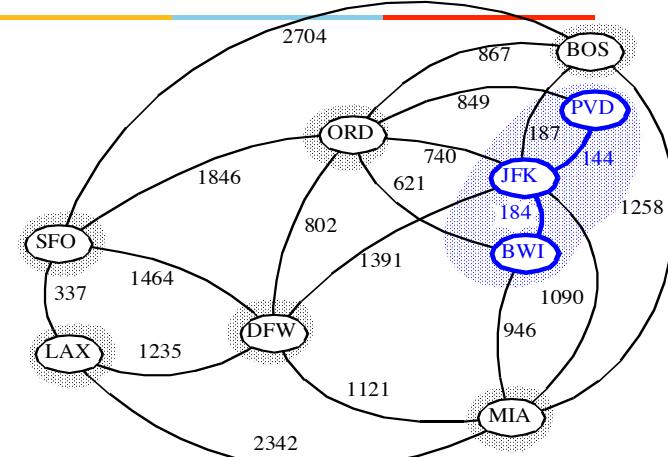
BITS Pilani, Pilani Campus

## Example



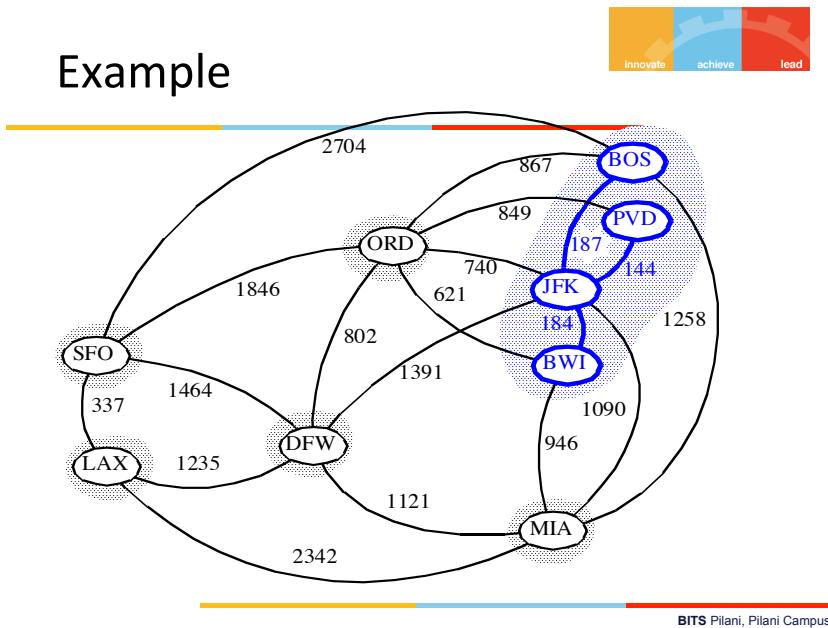
BITS Pilani, Pilani Campus

## Example

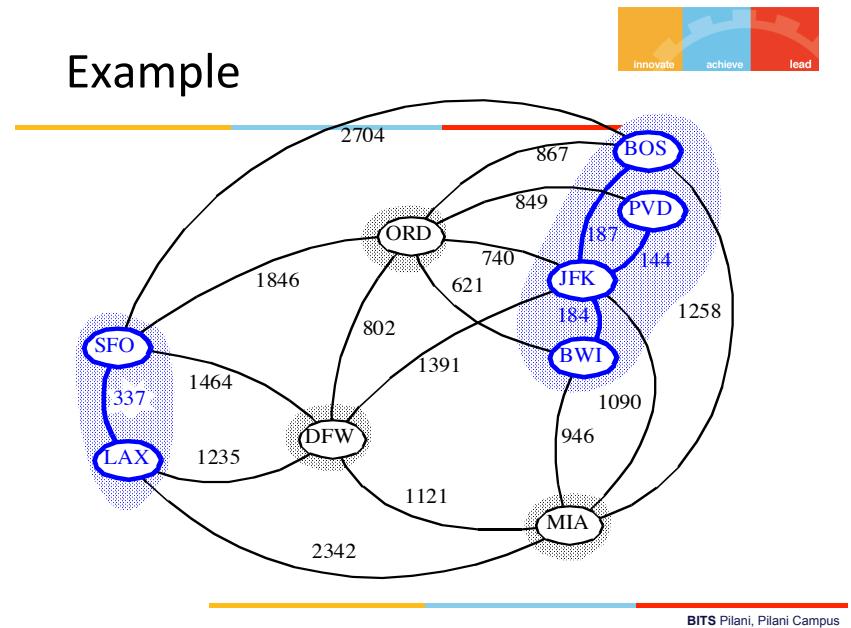


BITS Pilani, Pilani Campus

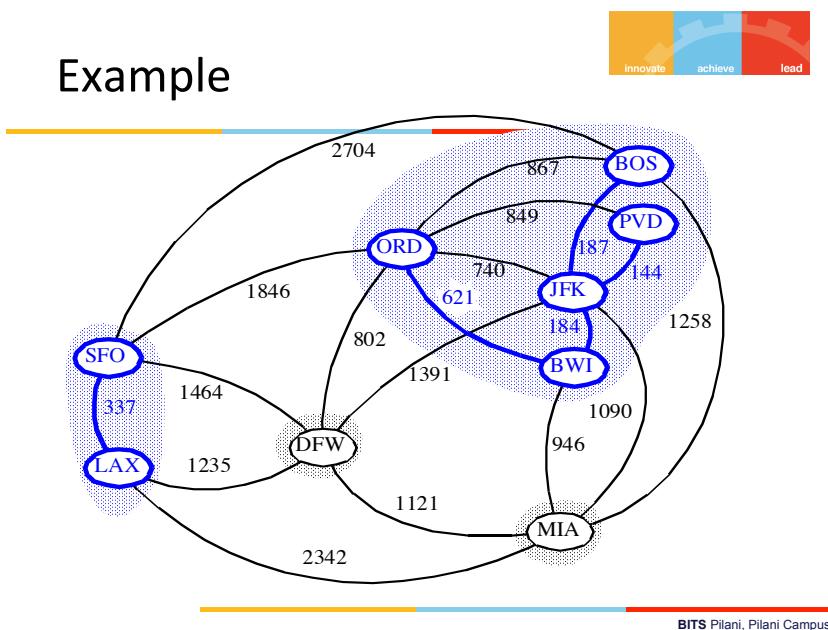
## Example



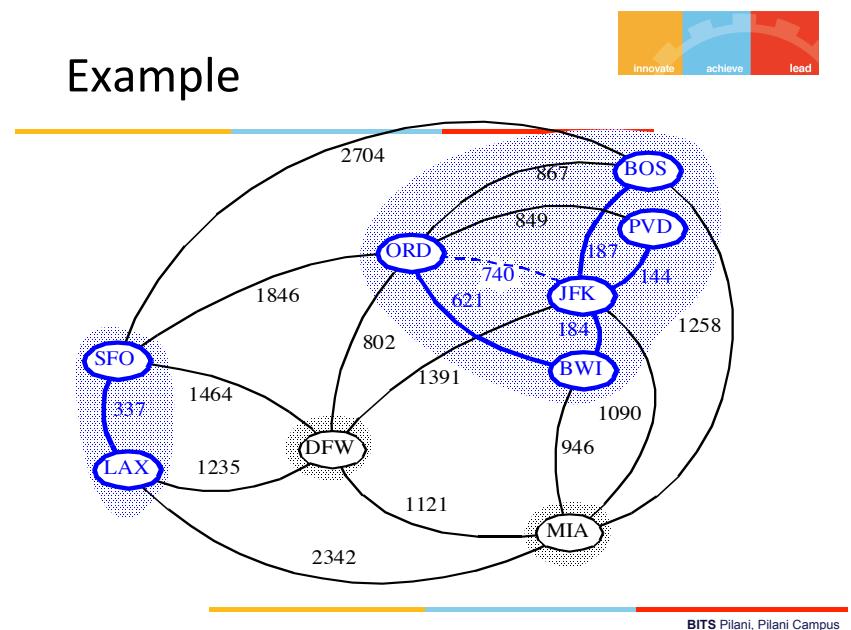
## Example



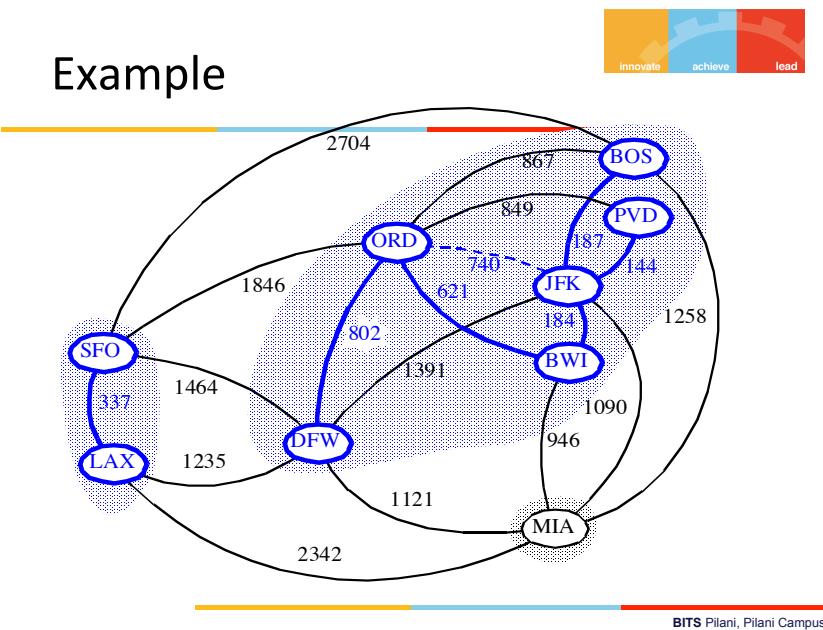
## Example



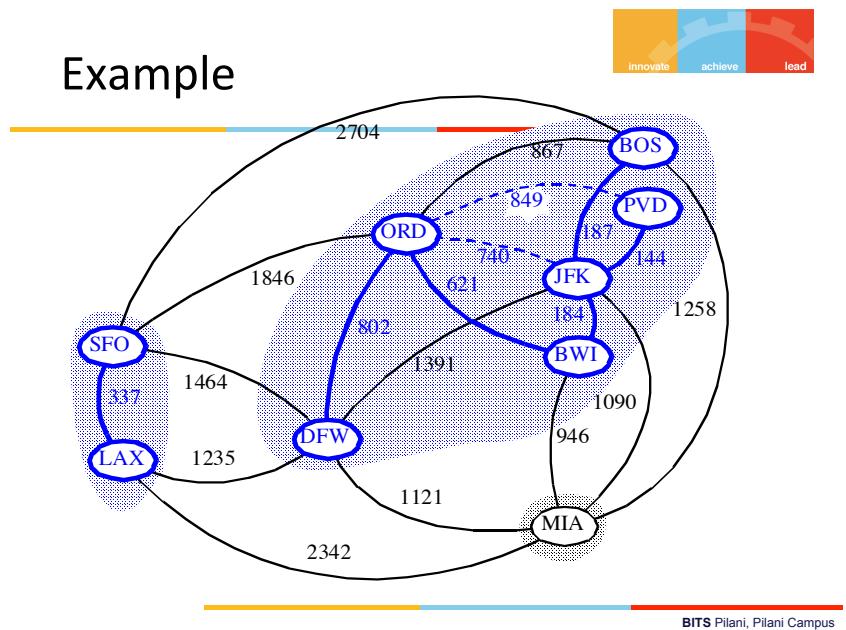
## Example



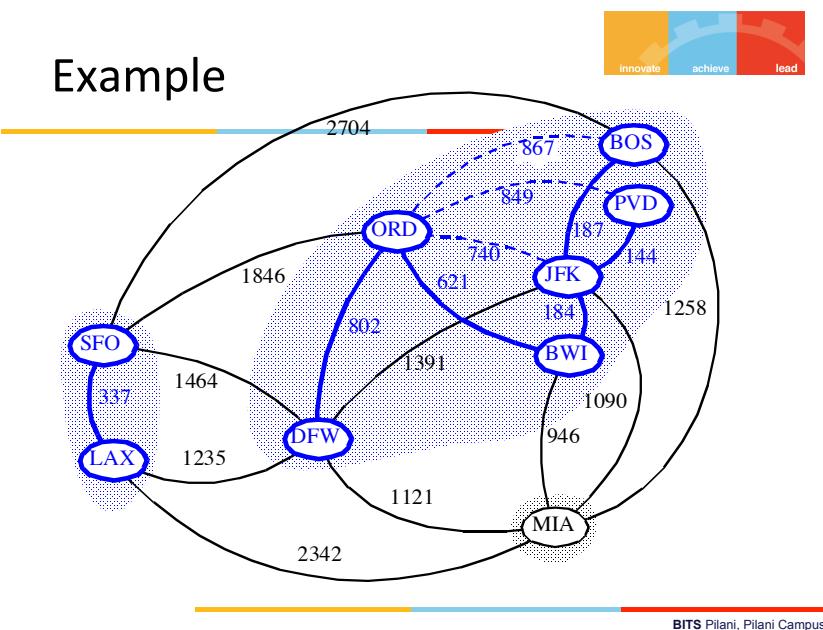
## Example



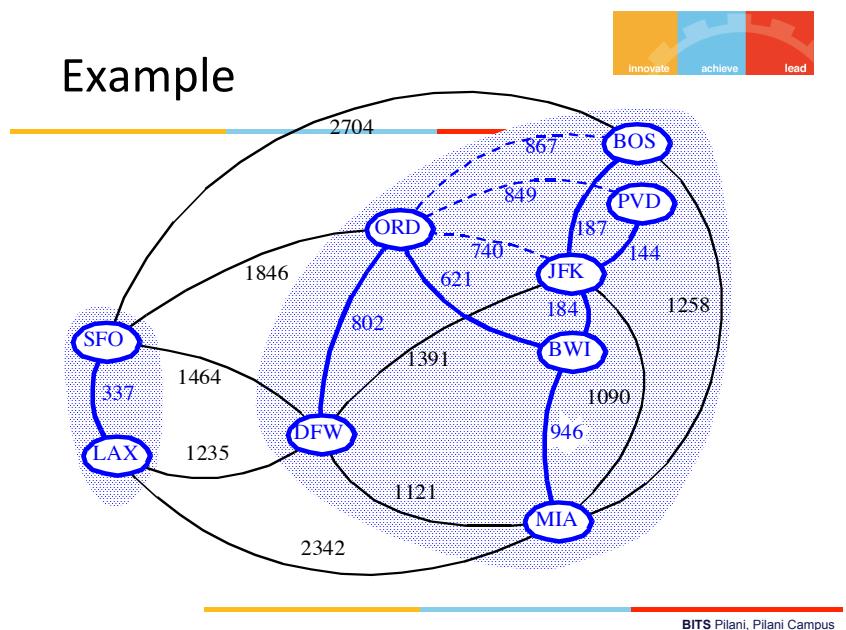
## Example



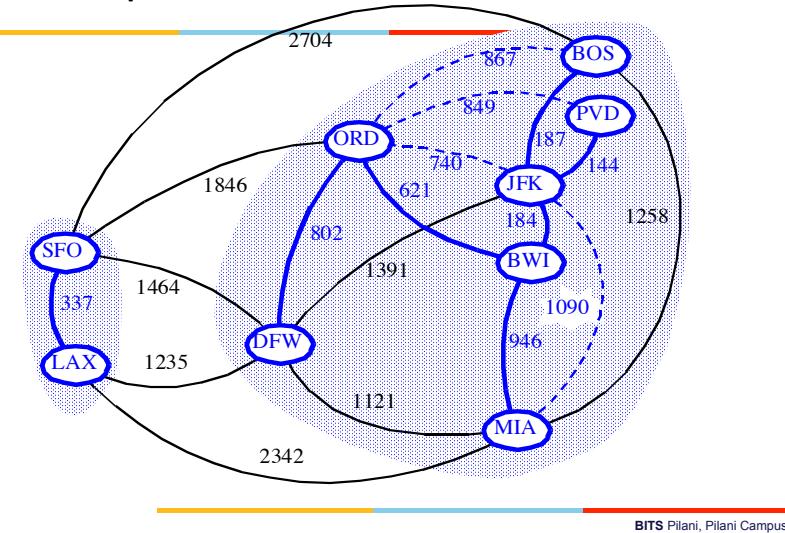
## Example



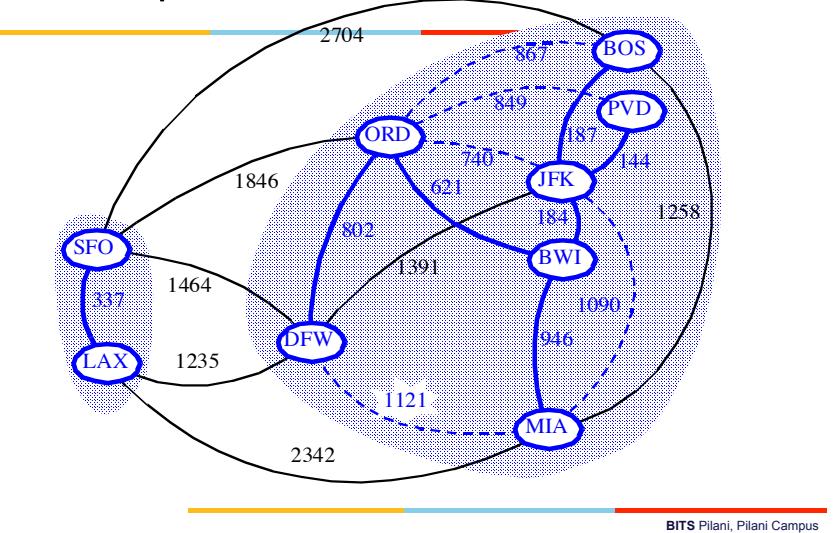
## Example



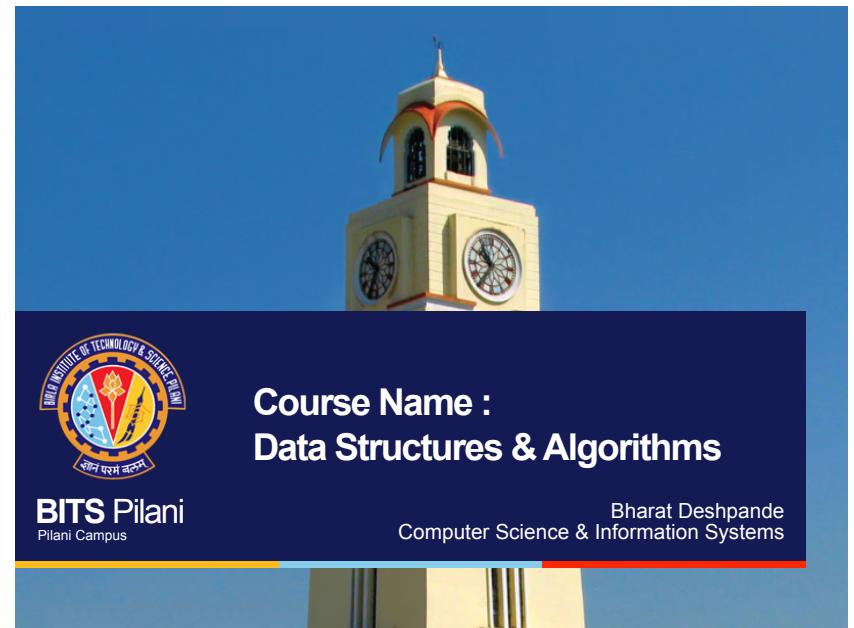
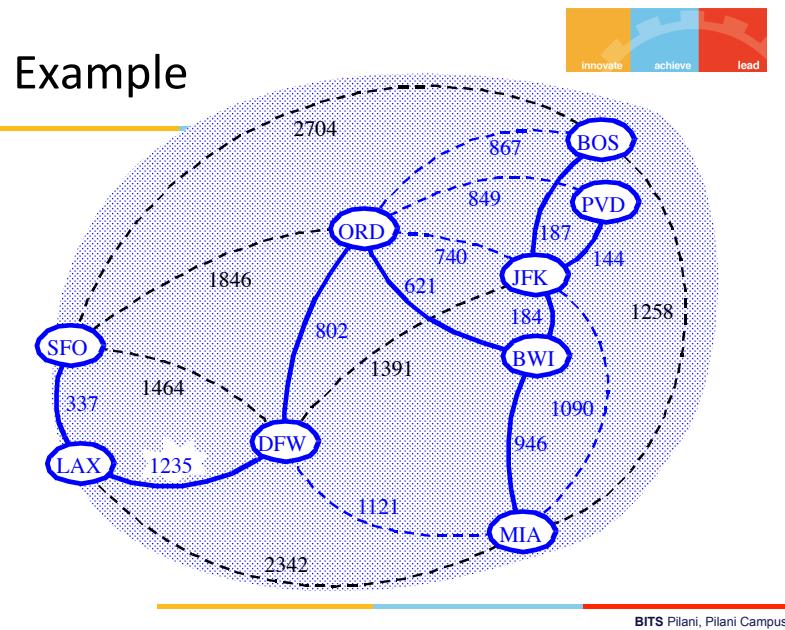
## Example



## Example



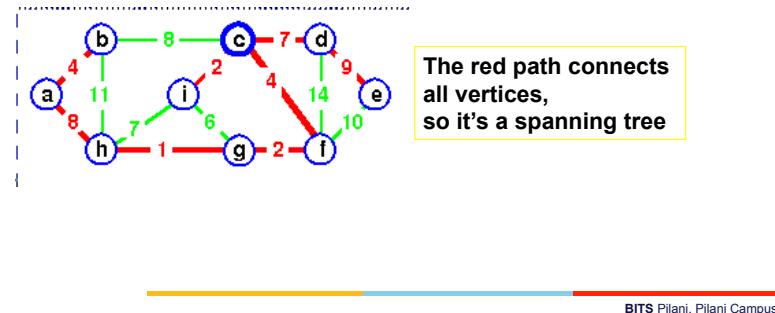
## Example



## Graphs - Definitions

- Spanning Tree

- A **spanning tree** is a set of  $|V|-1$  edges that connect all the vertices of a graph



## Growing a MST

- The procedure manages a set  $A$  that is always a subset of some MST.
  - At each step, an edge  $(u, v)$  is added to  $A$  such that  $A \cup \{(u, v)\}$  is also a subset of a MST.
  - Edge  $(u, v)$  is called a **safe edge**.
  - GENERIC-MST( $G$ )**
- ```

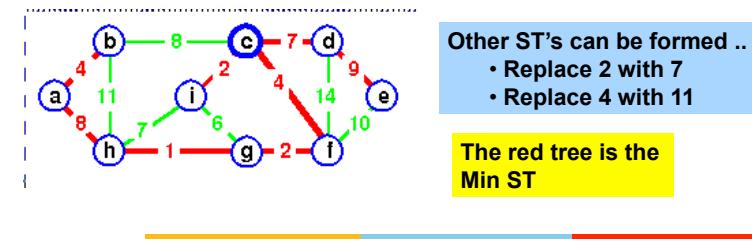
 $A \leftarrow \emptyset$ 
while  $A$  does not form a spanning tree
    do find an edge  $(u, v)$  that is safe for  $A$ 
         $A \leftarrow A \cup \{(u, v)\}$ 
return  $A$ 

```
- We design two MST algorithms, each of which uses a specific rule to determine the safe edge to be added.

Graphs - Definitions

- Minimum Spanning Tree**

- Generally there is more than one spanning tree
- If a cost c_{ij} is associated with edge $e_{ij} = (v_i, v_j)$ then the **minimum spanning tree** is the set of edges E_{span} such that $C = \sum (c_{ij} \mid \forall e_{ij} \in E_{\text{span}})$ is a minimum



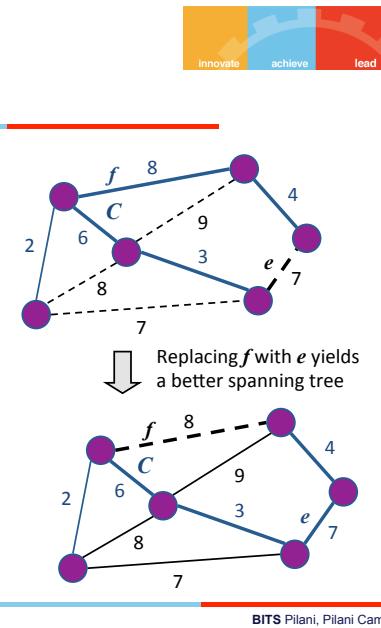
Cycle Property

Cycle Property:

- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T and C let be the cycle formed by e with T
- For every edge f of C , $\text{weight}(f) \leq \text{weight}(e)$

Proof:

- By contradiction
- If $\text{weight}(f) > \text{weight}(e)$ we can get a spanning tree of smaller weight by replacing e with f



Partition Property

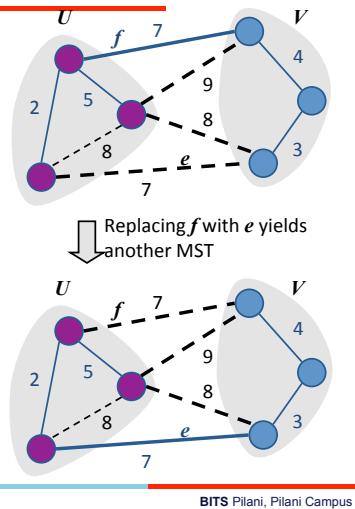


Partition Property:

- Consider a partition of the vertices of G into subsets U and V
- Let e be an edge of minimum weight across the partition
- There is a minimum spanning tree of G containing edge e

Proof:

- Let T be an MST of G
- If T does not contain e , consider the cycle C formed by e with T and let f be an edge of C across the partition
- By the cycle property,
 $\text{weight}(f) \leq \text{weight}(e)$
- Thus, $\text{weight}(f) = \text{weight}(e)$
- We obtain another MST by replacing f with e



BITS Pilani, Pilani Campus

Prim's Algorithm – Pseudo Code



MST-PRIM(G, r)

```

 $Q \leftarrow V[G] // Q - priority queue initialize$ 
 $\text{for each } u \in Q$ 
     $\text{do } d[u] \leftarrow \infty$ 
 $D[r] = 0$ 
 $\Pi[r] \leftarrow \text{NIL}$ 
 $\text{while } Q \neq \emptyset$ 
     $\text{do } u \leftarrow \text{EXTRACT-MIN}(Q)$ 
         $\text{for each } v \in \text{Adj}[u]$ 
             $\text{do if } v \in Q \text{ and } w(u, v) < d[v]$ 
                 $\text{then } \Pi[v] \leftarrow u$ 
                 $d[v] \leftarrow w(u, v)$ 

```

The running time is $O(m \log n)$

BITS Pilani, Pilani Campus

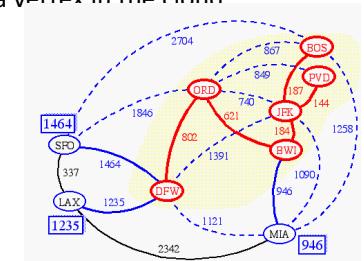
Prim-Jarnik's Algorithm



- Initially discovered in 1930 by Vojtěch Jarník, then rediscovered in 1957 by Robert C. Prim
- We pick an arbitrary vertex s and we grow the MST as a cloud of vertices, starting from s
- We store with each vertex v a label $d(v)$ = the smallest weight of an edge connecting v to a vertex in the cloud

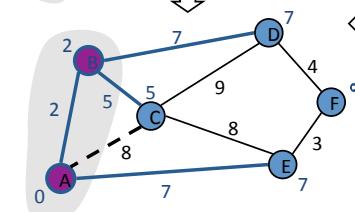
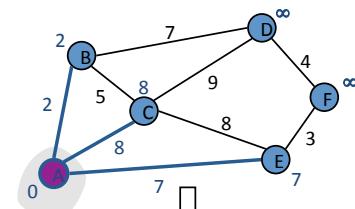
At each step:

- We add to the cloud the vertex u outside the cloud with the smallest distance label
- We update the labels of the vertices adjacent to u

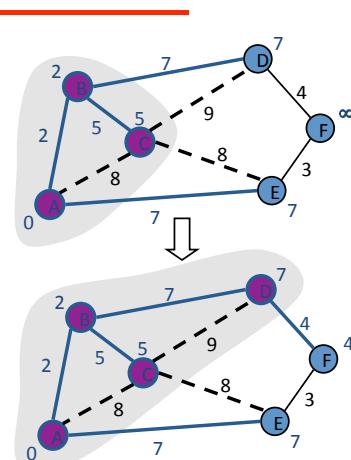


BITS Pilani, Pilani Campus

Example

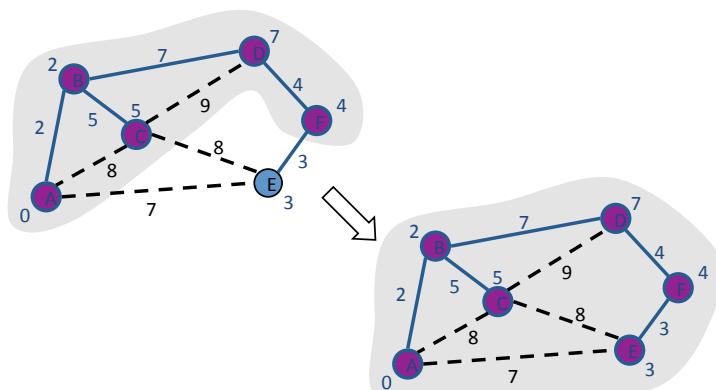


BITS Pilani, Pilani Campus



BITS Pilani, Pilani Campus

Example (contd.)



BITS Pilani, Pilani Campus

Kruskal's Algorithm

- Created in 1957 by Joseph Kruskal

Algorithm

- Starts by choosing an edge in the graph with minimum weight.
- Successively add edges with minimum weight that do not form a cycle with those edges already chosen.
- Terminates after $n-1$ edges have been selected (n is the number of vertices)

Note

In Kruskal's algorithm the set A is a forest.

The safe edge added to A is always a least-weight edge in the graph that connects two distinct components.



Kruskal's algorithm – Pseudo code



Algorithm Kruskal(G)

Input: A weighted graph G .

Output: An MST T for G .

Let P be a partition of the vertices of G , where each vertex forms a separate set.

Let Q be a priority queue storing the edges of G , sorted by their weights

Let T be an initially-empty tree

while Q is not empty **do**

$(u,v) \leftarrow Q.\text{removeMinElement}()$

if $P.\text{find}(u) \neq P.\text{find}(v)$ **then**

 Add (u,v) to T

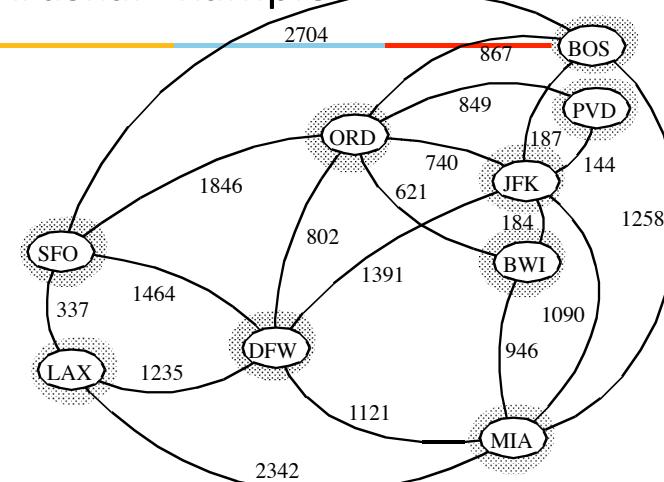
$P.\text{union}(u,v)$

Running time: $O(m \log m)$

return T

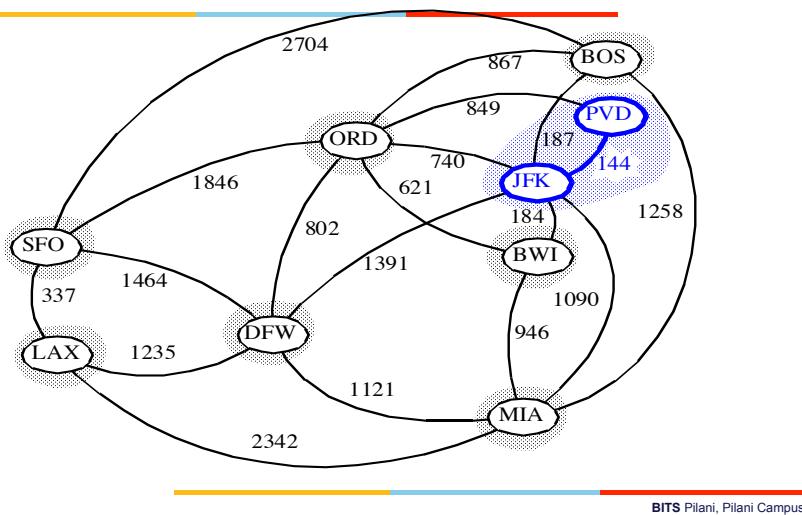
BITS Pilani, Pilani Campus

Kruskal Example

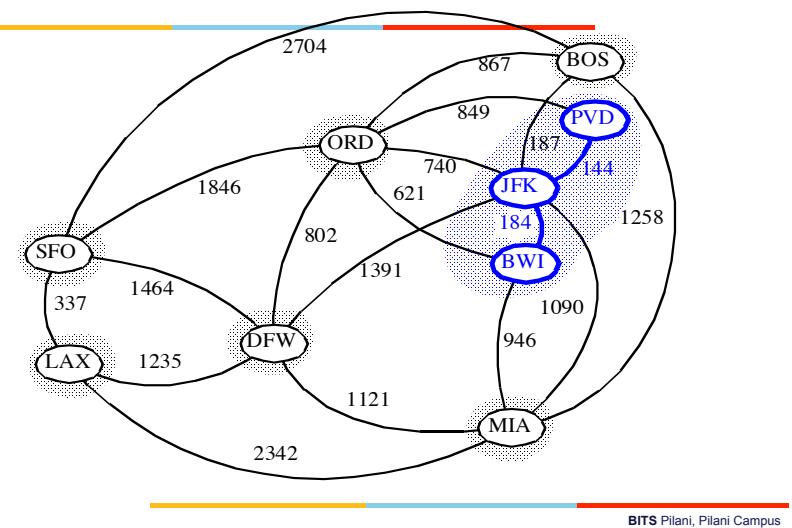


BITS Pilani, Pilani Campus

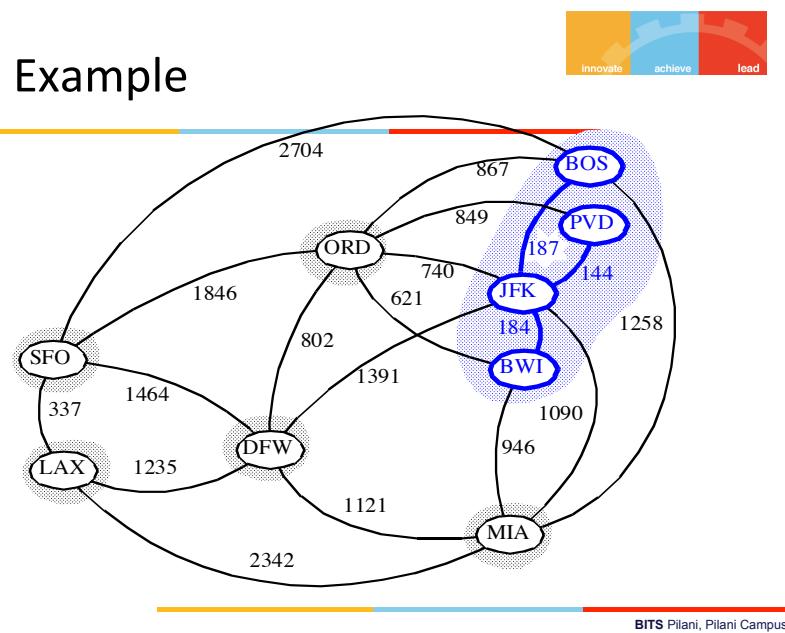
Example



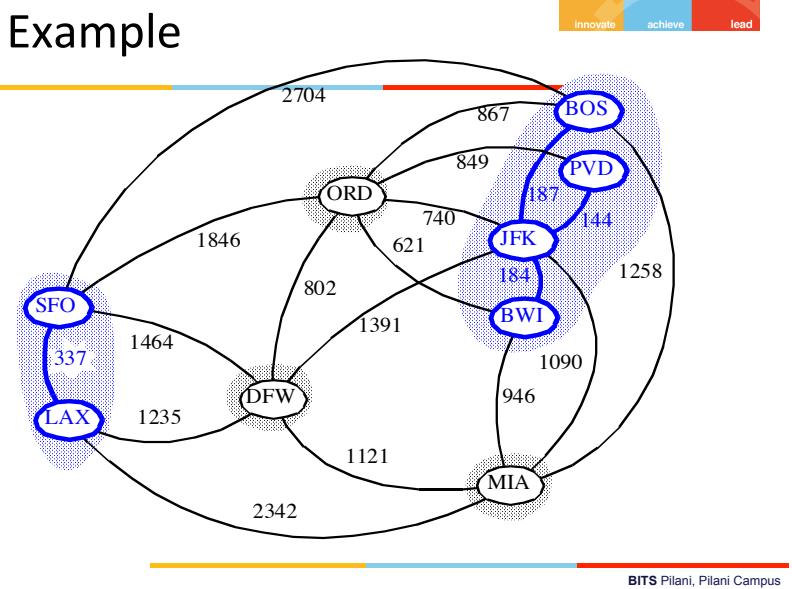
Example



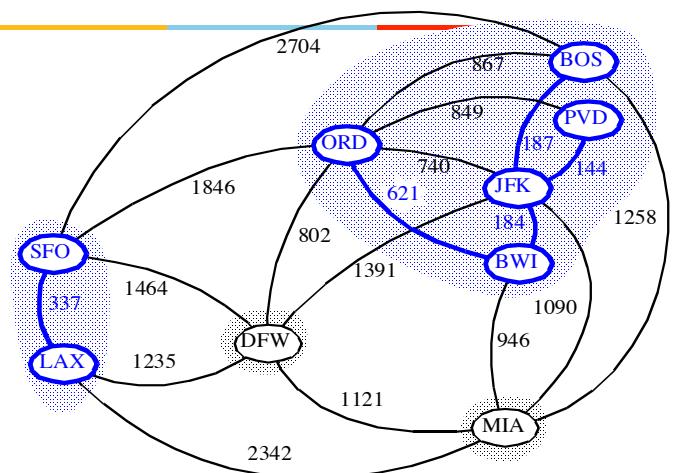
Example



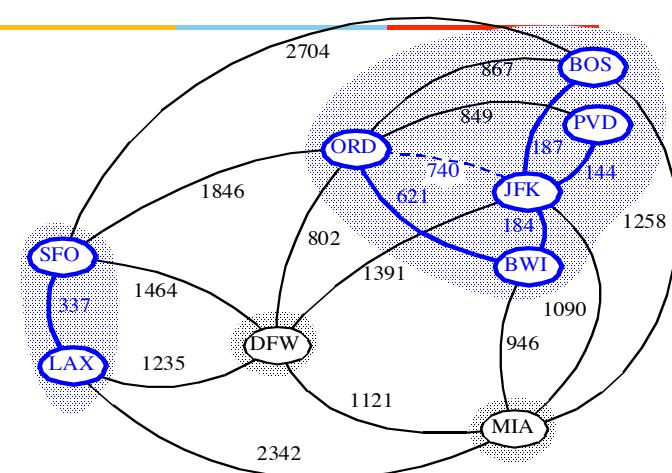
Example



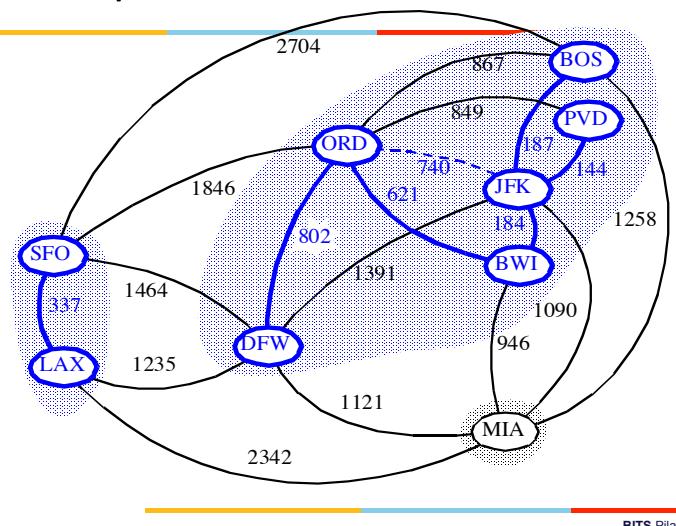
Example



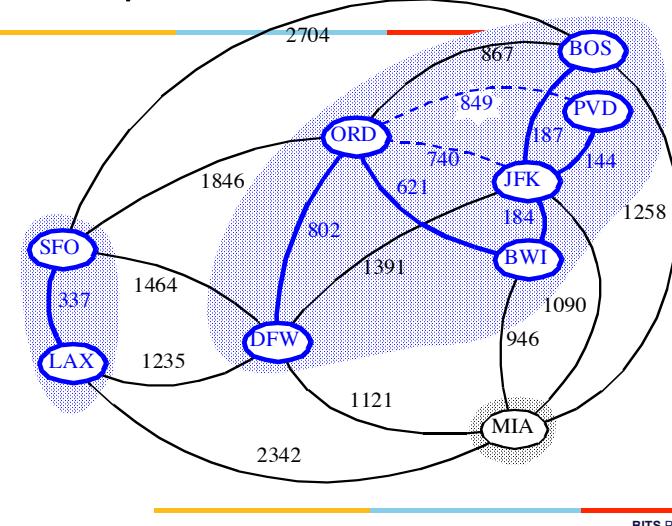
Example



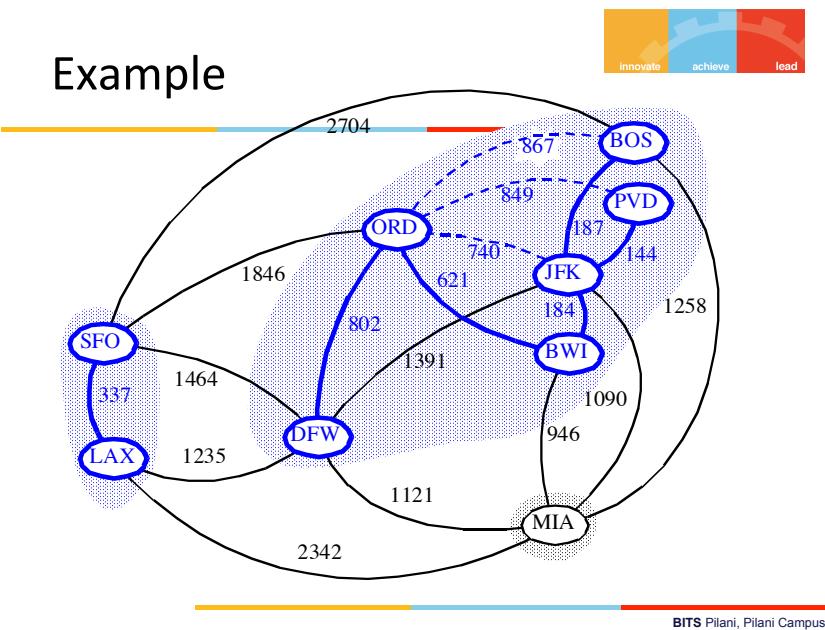
Example



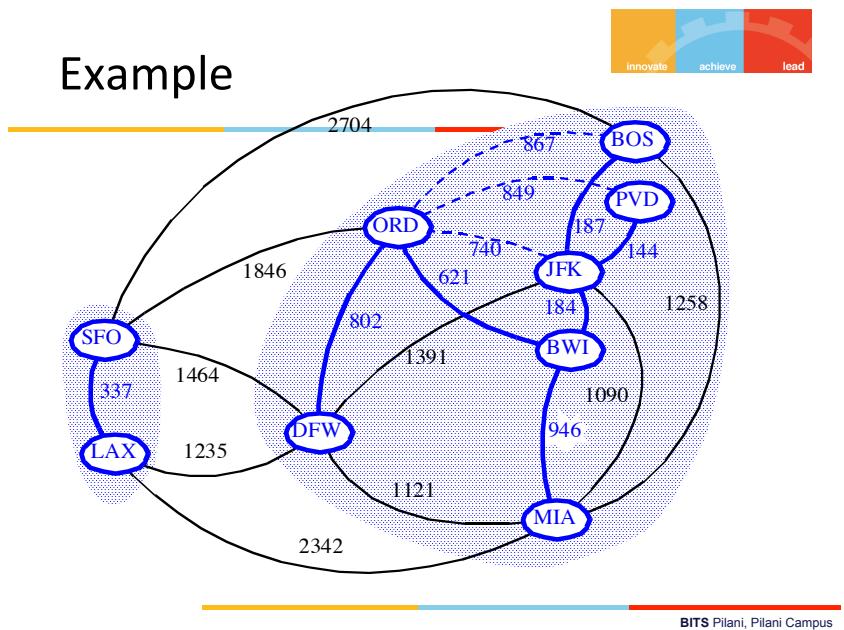
Example



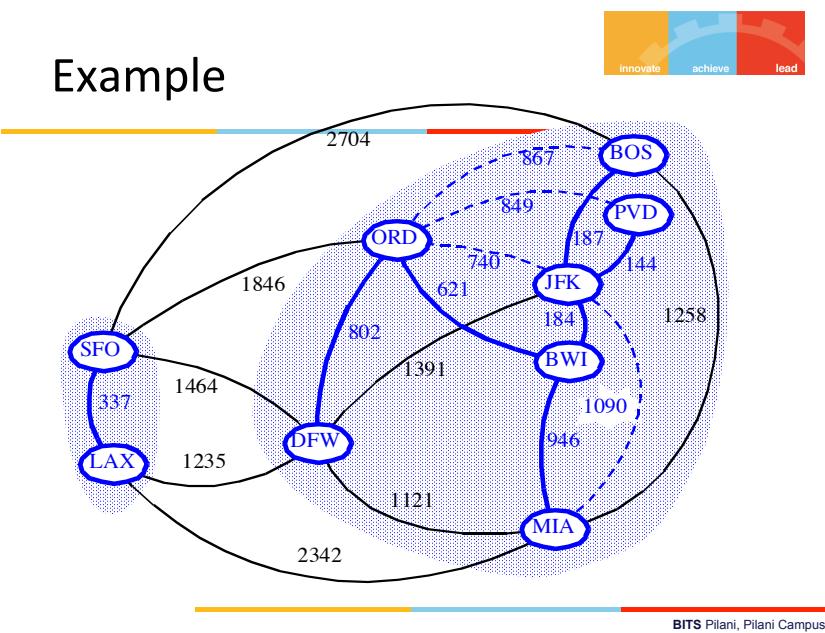
Example



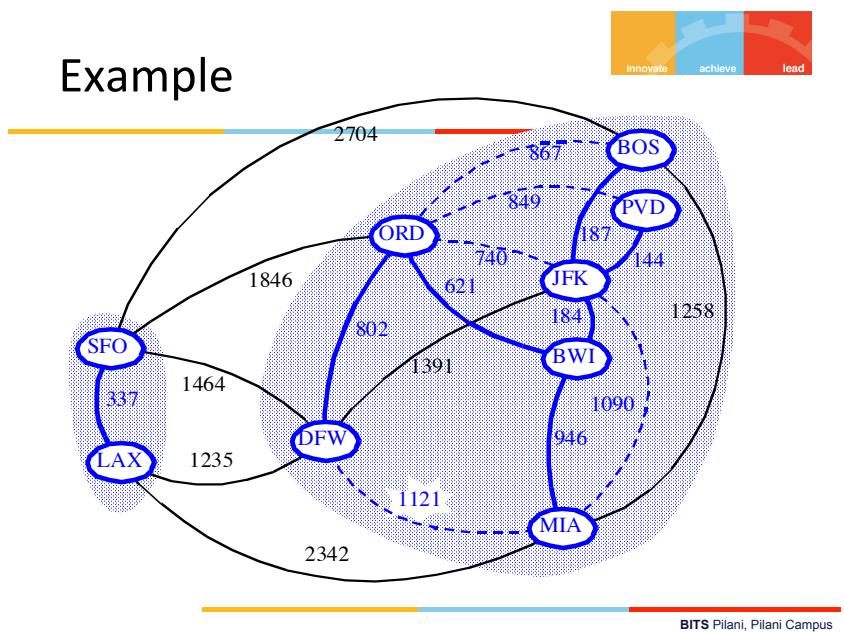
Example



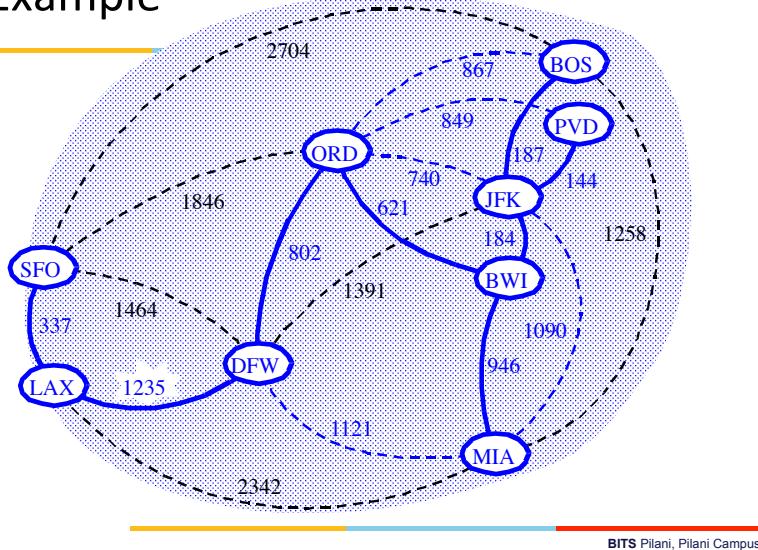
Example



Example



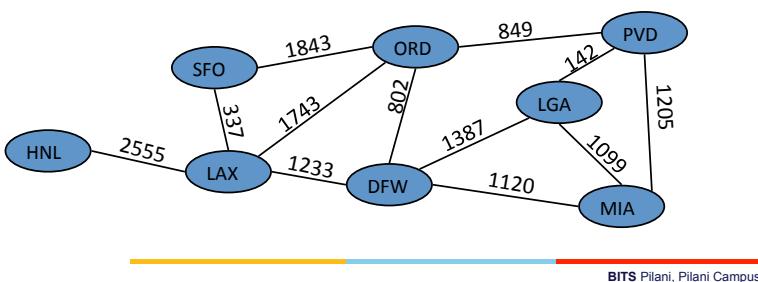
Example



Weighted Graphs



- In a weighted graph, each edge has an associated numerical value, called the **weight** of the edge
- Edge weights may represent, distances, costs, etc.
- Example:**
 - In a flight route graph, the weight of an edge represents the distance in miles between the endpoint airports

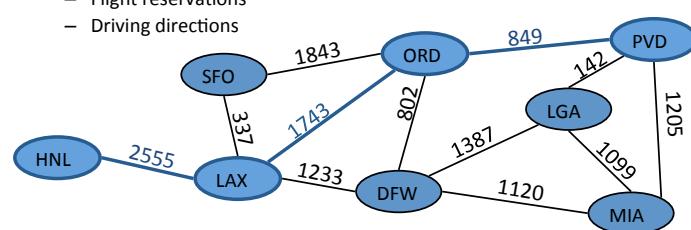


BITS Pilani, Pilani Campus

Shortest Paths



- Given a weighted graph and two vertices u and v , we want to find a path of minimum total weight between u and v .
 - Length of a path is the sum of the weights of its edges.
- Example:**
 - Shortest path between Providence (PVD) and Honolulu (HNL)
- Applications**
 - Internet packet routing
 - Flight reservations
 - Driving directions



BITS Pilani, Pilani Campus



Weighted shortest path

- We want to minimize the total mileage.
- Breadth-first search does not work!
 - Minimum number of hops does not mean minimum distance.

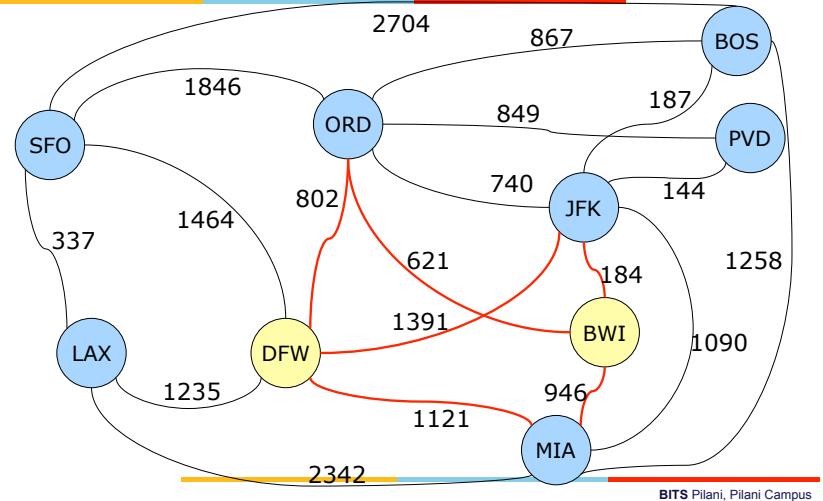
BITS Pilani, Pilani Campus

A greedy algorithm

- Assume that every city is infinitely far away.
 - I.e., every city is ∞ miles away from BWI (except BWI, which is 0 miles away).
 - Now perform something similar to breadth-first search, and *optimistically guess that we have found the best path to each city as we encounter it.*
 - If we later discover we are wrong and find a better path to a particular city, then update the distance to that city.

BITS Pilani, Pilani Campus

Three 2-hop routes to DFW



Intuition behind Dijkstra's algorithm

- For our airline-mileage problem, we can start by guessing that every city is ∞ miles away.
 - Mark each city with this guess.
- Find all cities one hop away from BWI, and check whether the mileage is less than what is currently marked for that city.
 - If so, then revise the guess.
- Continue for 2 hops, 3 hops, etc.

BITS Pilani, Pilani Campus

Dijkstra's Algorithm



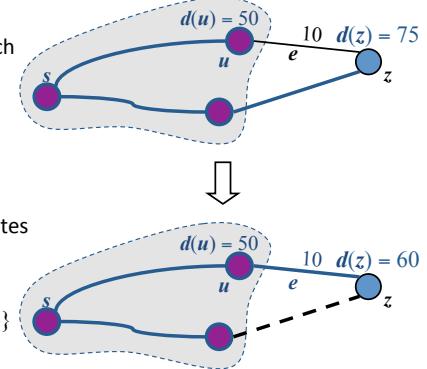
- **Dijkstra's algorithm** computes the distances of all the vertices from a given start vertex s
- **Assumptions:**
 - the graph is connected
 - the graph is undirected and simple
 - the edge weights are nonnegative
- Algorithm grows a **cloud** of vertices, beginning with s and eventually covering all the vertices
- We store with each vertex v a label $d(v)$ representing the distance of v from s in the subgraph consisting of the cloud and its adjacent vertices
- **At each step**
 - We add to the cloud the vertex u outside the cloud with the smallest distance label, $d(u)$
 - We update the labels of the vertices adjacent to u

BITS Pilani, Pilani Campus

Edge Relaxation



- Update procedure is known as **relaxation procedure**.
 - Consider an edge $e = (u,z)$ such that
 - u is the vertex most recently added to the cloud
 - z is not in the cloud
 - The relaxation of edge e updates distance $d(z)$ as follows:
- $$d(z) \leftarrow \min\{d(z), d(u) + \text{weight}(e)\}$$



BITS Pilani, Pilani Campus

Dijkstra's algorithm



- Algorithm initialization:
 - Label each node with the distance ∞ , except start node, which is labeled with distance 0.
 - $D[v]$ is the distance label for v .
- Put all nodes into a priority queue Q , using the distances as labels.

BITS Pilani, Pilani Campus

Edge Relaxation



- While Q is not empty do:
 - $u = Q.\text{removeMin}$
 - for each node z one hop away from u do:
 - if $D[u] + \text{miles}(u,z) < D[z]$ then
 - $D[z] = D[u] + \text{miles}(u,z)$
 - change key of z in Q to $D[z]$
- **Note** use of priority queue allows “finished” nodes to be found quickly (in $O(\log N)$ time).
- Running time is $O(m \log n)$

BITS Pilani, Pilani Campus

DIJKSTRA'S ALGORITHM - WHY IT WORKS



Property 1: Triangle inequality

If $\delta(u,v)$ is the shortest path length between u and v ,
 $\delta(u,v) \leq \delta(u,x) + \delta(x,v)$

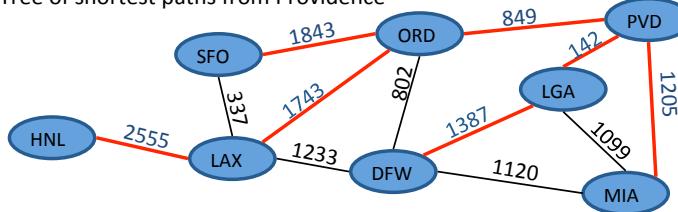
Property 2: A subpath of a shortest path is itself a shortest path

Property 3:

There is a tree of shortest paths from a start vertex to all the other vertices

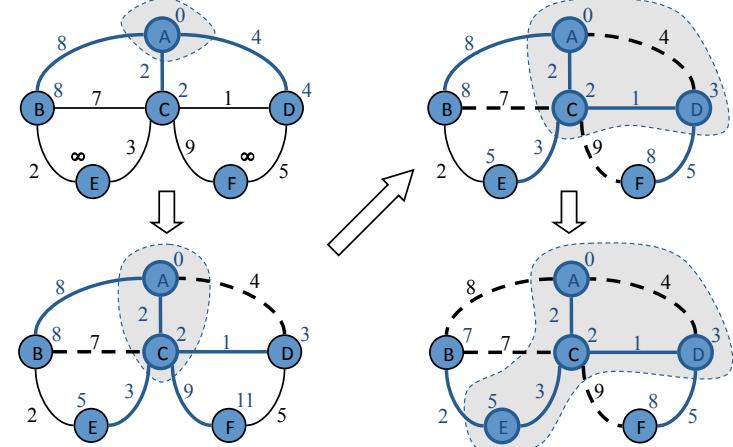
Example:

Tree of shortest paths from Providence



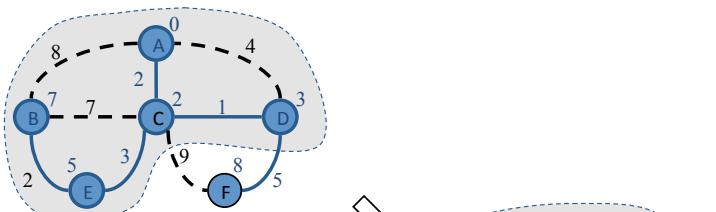
BITs Pilani, Pilani Campus

Example



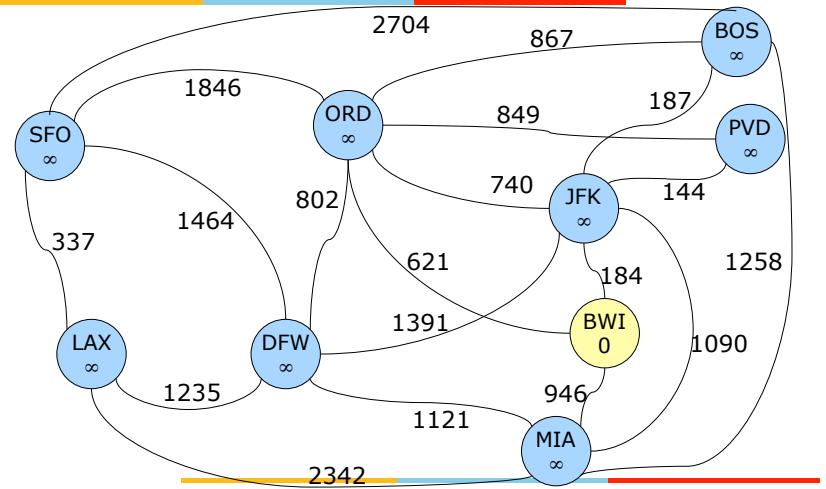
BITs Pilani, Pilani Campus

Example (cont.)



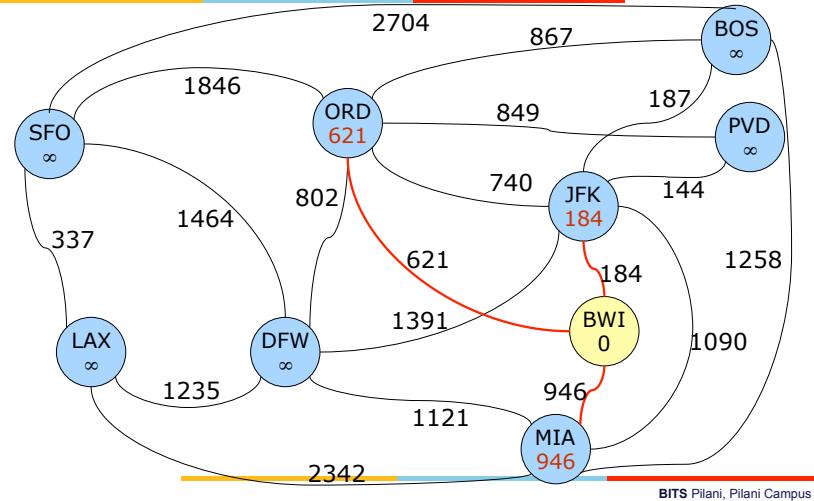
BITs Pilani, Pilani Campus

Shortest mileage from BWI

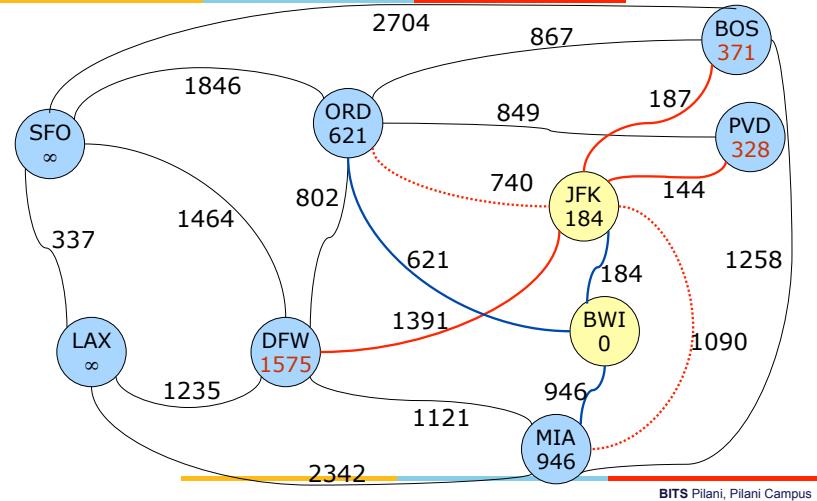


BITs Pilani, Pilani Campus

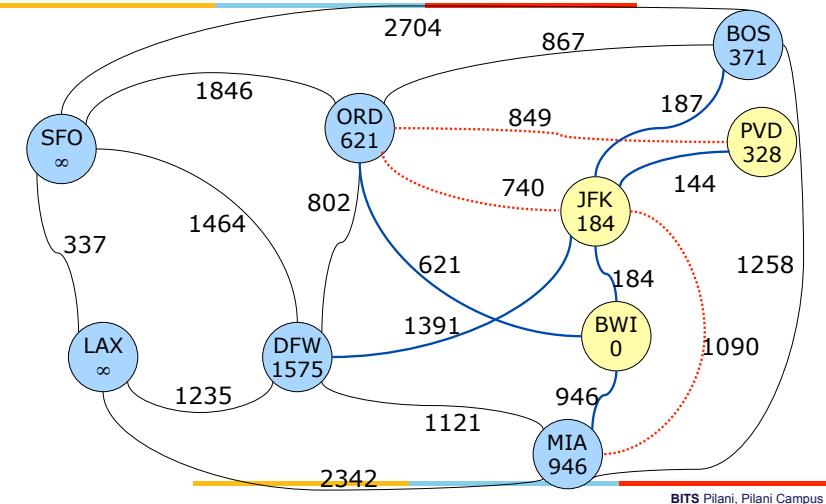
Shortest mileage from BWI



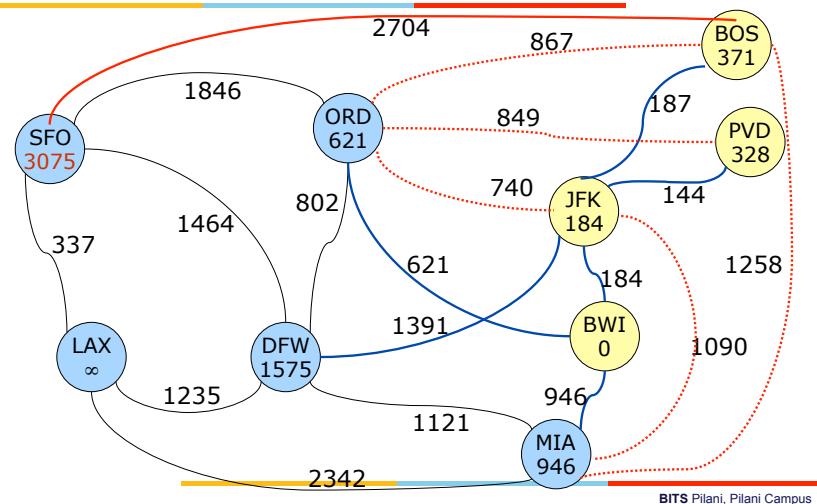
Shortest mileage from BWI



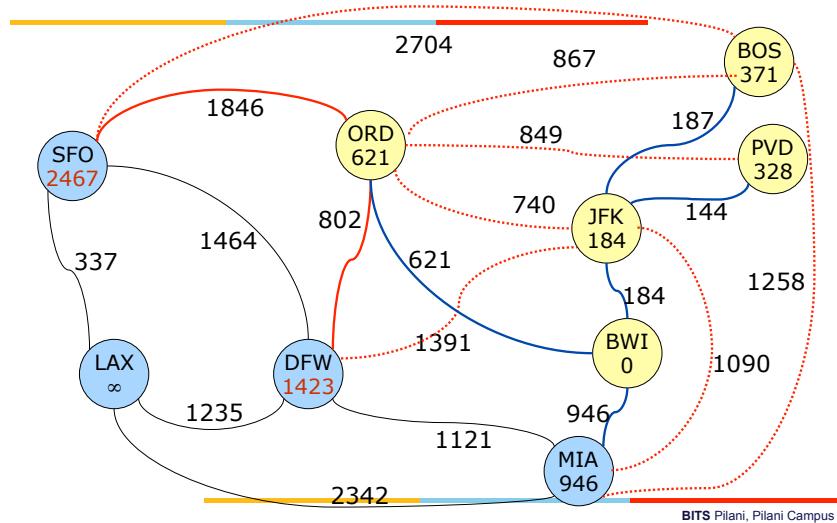
Shortest mileage from BWI



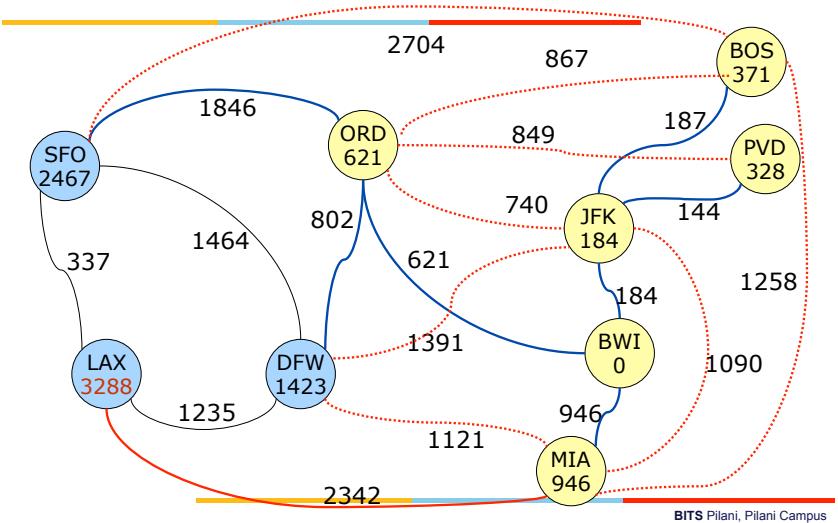
Shortest mileage from BWI



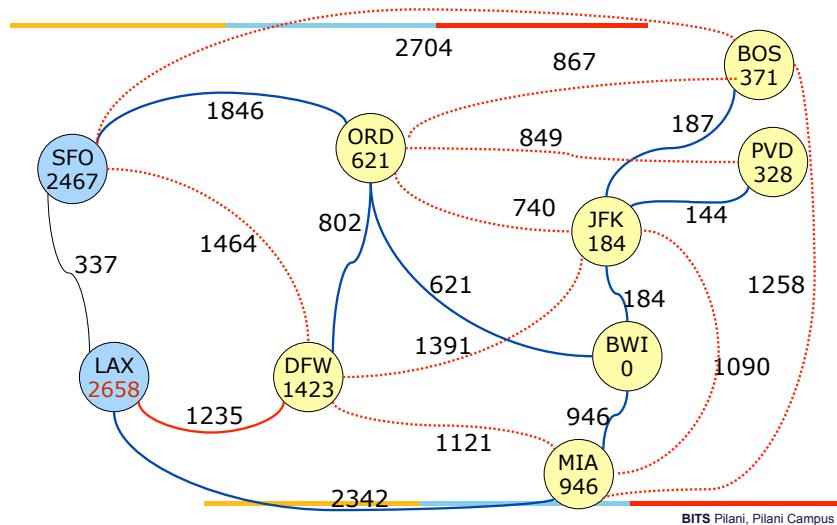
Shortest mileage from BWI



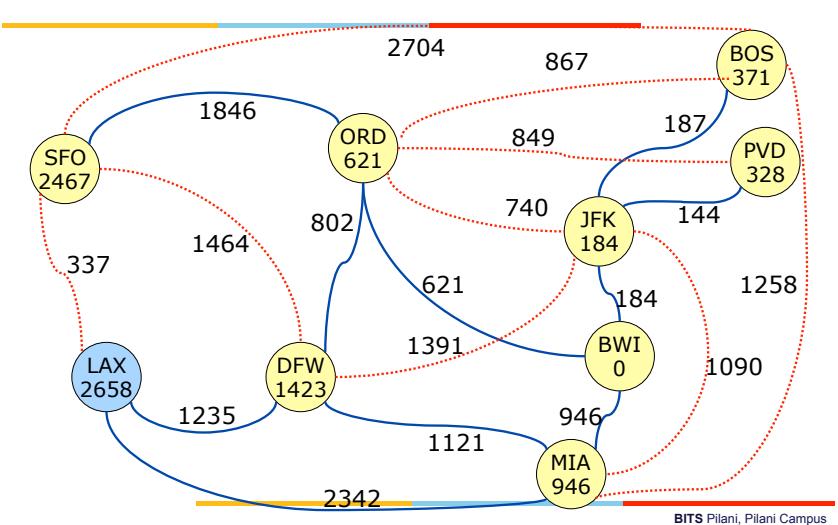
Shortest mileage from BWI



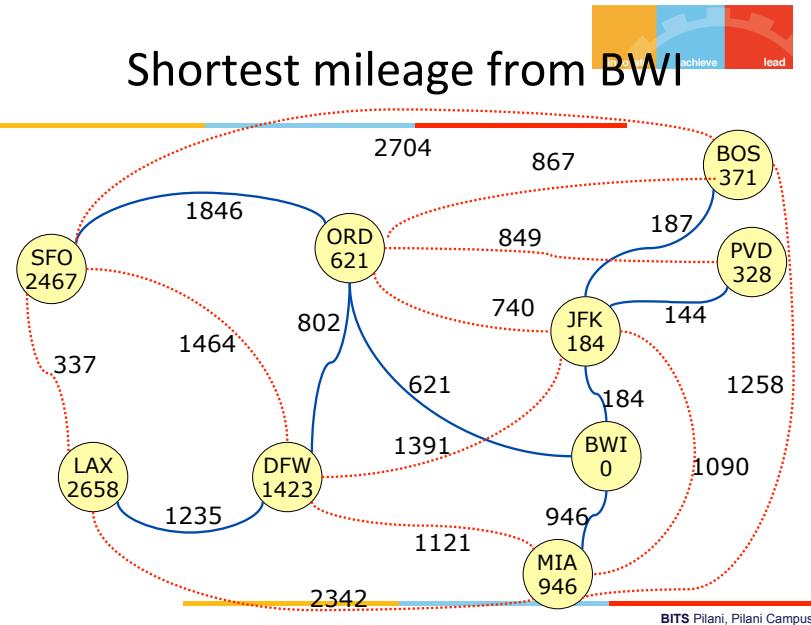
Shortest mileage from BWI



Shortest mileage from BWI



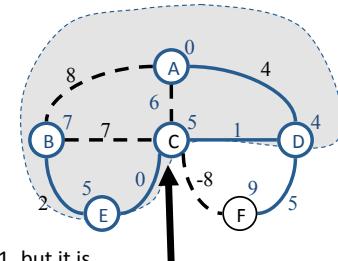
Shortest mileage from BWI



Why It Doesn't Work for Negative-Weight Edges

- Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.

- If a node with a negative incident edge were to be added late to the cloud, it could mess up distances for vertices already in the cloud.



BITS Pilani, Pilani Campus

Bellman-Ford Algorithm



- Works even with negative-weight edges
- Must assume directed edges (for otherwise we would have negative-weight cycles)
- Iteration i finds all shortest paths that use i edges.
- Can be extended to detect a negative-weight cycle if it exists
- Uses edge relaxation as in Dijkstra's but not in conjunction with the greedy method.

BITS Pilani, Pilani Campus

Bellman-Ford Algorithm

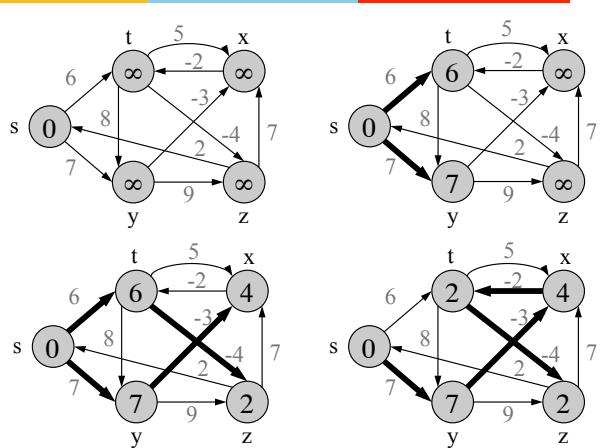


```

Bellman-Ford( $G, s$ )
  for each vertex  $u \in V$ 
     $d[u] \leftarrow \infty$ 
     $\text{parent}[u] \leftarrow \text{NIL}$ 
   $d[s] \leftarrow 0$ 
  for  $i \leftarrow 1$  to  $|V|-1$  do
    for each edge  $(u, v)$  outgoing from  $u$  do
      {Perform the relaxation operation on  $(u, v)$ }
      if  $d[u] + w(u, v) < d[v]$  then
         $d[v] \leftarrow d[u] + w(u, v)$ 
  if there are no more edges left with potential relaxation operation then
    return the label  $d[u]$  of each vertex  $u$ 
  else
    return  $G$  contains a negative weight cycle.
  
```

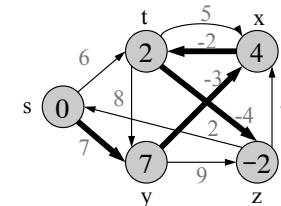
BITS Pilani, Pilani Campus

Bellman-Ford Example



BITS Pilani, Pilani Campus

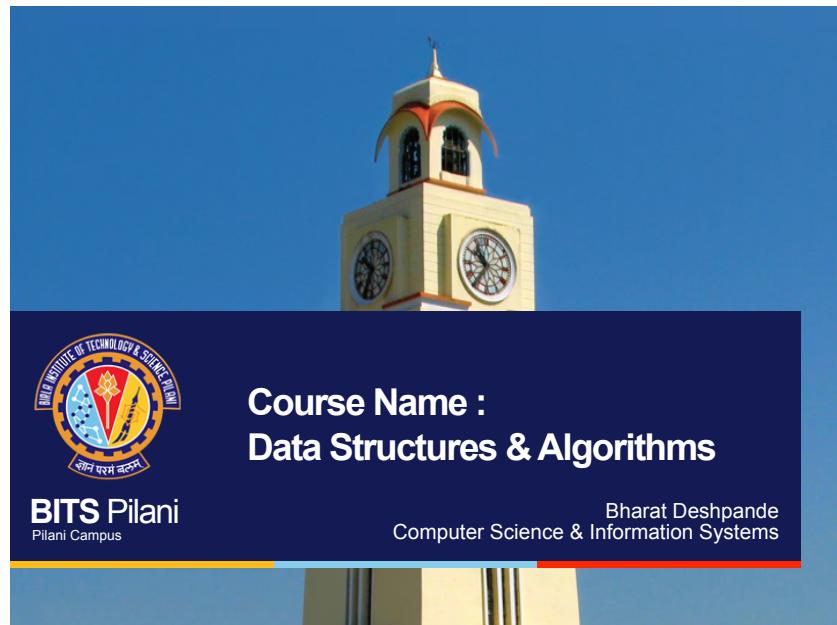
Bellman-Ford Example



- Bellman-Ford running time:

$$-(|V|-1)|E| + |E| = \Theta(VE)$$

BITS Pilani, Pilani Campus



The Greedy Method

- **The greedy method** is a general algorithm design paradigm, built on the following elements:
 - **configurations**: different choices, collections, or values to find
 - **objective function**: a score assigned to configurations, which we want to either maximize or minimize
- It works best when applied to problems with the **greedy-choice** property:
 - a globally-optimal solution can always be found by a series of local improvements from a starting configuration.

BITS Pilani, Pilani Campus



The Fractional Knapsack Problem

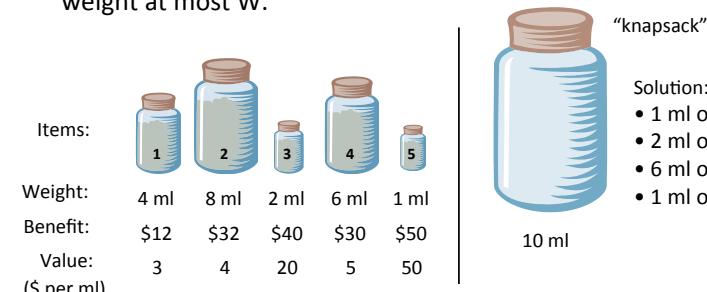
- Given: A set S of n items, with each item i having
 - b_i - a positive benefit
 - w_i - a positive weight
- Goal:** Choose items with maximum total benefit but with weight at most W .
- If we are allowed to take fractional amounts, then this is the **fractional knapsack problem**.
 - In this case, we let x_i denote the amount we take of item i

– Objective: maximize $\sum_{i \in S} b_i(x_i / w_i)$

– Constraint: $\sum_{i \in S} x_i \leq W$

BITS Pilani, Pilani Campus

- Given: A set S of n items, with each item i having
 - b_i - a positive benefit
 - w_i - a positive weight
- Goal: Choose items with maximum total benefit but with weight at most W .



BITS Pilani, Pilani Campus



The Fractional Knapsack Algorithm

- Greedy choice:** Keep taking item with highest **value** (benefit to weight ratio)
 - Since $\sum b_i(x_i / w_i) = \sum (b_i / w_i)x_i$
 - Run time:** $O(n \log n)$.
- Correctness:** Suppose there is a better solution
 - there are two items i & j such that $x_i < w_i$, $x_j > 0$ and $v_i < v_j$
 - If we substitute some i with j , we get a better solution
 - How much of i : $\min\{w_i - x_i, x_j\}$
 - Thus, there is no better solution than the greedy one

```

Algorithm fractionalKnapsack(S, W)
  Input: set  $S$  of items w/ benefit  $b_i$  and weight  $w_i$ ; max. weight  $W$ 
  Output: amount  $x_i$  of each item  $i$  to maximize benefit w/ weight at most  $W$ 
  for each item i in S
     $x_i \leftarrow 0$ 
     $v_i \leftarrow b_i / w_i$  {value}
     $w \leftarrow 0$  {total weight}
    while  $w < W$ 
      remove item i w/ highest  $v_i$ 
       $x_i \leftarrow \min\{w_i, W - w\}$ 
       $w \leftarrow w + \min\{w_i, W - w\}$ 
  
```

BITS Pilani, Pilani Campus



Task Scheduling

- Given: a set T of n tasks, each having:
 - A start time, s_i
 - A finish time, f_i (where $s_i < f_i$)
- Goal:** Perform all the tasks using a minimum number of "machines."
- Tasks i & j are non-conflicting if $f_i \leq s_j$ or $f_j \leq s_i$
- Two tasks can be scheduled on the same machine if they are non-conflicting

BITS Pilani, Pilani Campus



Task Scheduling Algorithm

- **Greedy choice:** consider tasks by their start time and use as few machines as possible with this order.
 - **Run time:** $O(n \log n)$.
- **Correctness:** Suppose there is a better schedule.
 - We can use $k-1$ machines
 - The algorithm uses k
 - Let i be first task scheduled on machine k
 - Task i must conflict with $k-1$ other tasks
 - All these tasks also conflict with each other
 - But that means there is no non-conflicting schedule using $k-1$ machines

Algorithm *taskSchedule(T)*

```

Input: set  $T$  of tasks w/ start time  $s_i$  and finish time  $f_i$ 
Output: non-conflicting schedule with minimum number of machines
 $m \leftarrow 0$  {no. of machines}
while  $T$  is not empty
  remove task  $i$  w/ smallest  $s_i$ 
  if there's a machine  $j$  for  $i$  then
    schedule  $i$  on machine  $j$ 
  else
     $m \leftarrow m + 1$ 
    schedule  $i$  on machine  $m$ 

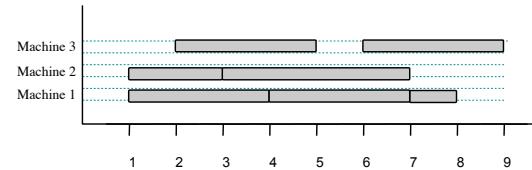
```

BITS Pilani, Pilani Campus



Example

- Given: a set T of n tasks, each having:
 - A start time, s_i
 - A finish time, f_i (where $s_i < f_i$)
 - $[1,4], [1,3], [2,5], [3,7], [4,7], [6,9], [7,8]$ (ordered by start)
- Goal: Perform all tasks on min. number of machines



BITS Pilani, Pilani Campus



Text Compression

- Given a string X , efficiently encode X into a smaller bit string Y .

Easy Solution

Since there are only 26 characters in English Language & there are 32 bit string of length 5, Each alphabet can be represented by a bit string of length 5.

- This is called a **fixed length code**.

BITS Pilani, Pilani Campus



Text Compression

Question

Is it possible to find coding scheme, in which fewer bits are used.

Answer – Variable length codes

- Alphabets that occur more frequently should be encoded using short bit strings & rarely occurring alphabets should be encoded using long bit strings

BITS Pilani, Pilani Campus

Example

	a	b	c	d	e	f
Frequency	45	13	12	16	9	5
Fixed length	000	001	010	011	100	101
Variable length	0	101	100	111	1101	1100

A file with 100000 character contains only the alphabets a – f.

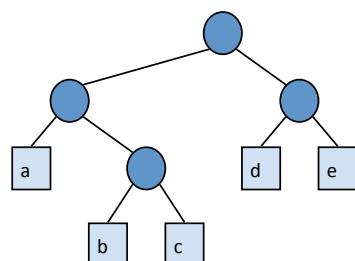
- Fixed length code requires 300000 bits to code the file.
- Variable length coding requires 224000 bits to code the file.

Encoding Tree

Key: A binary code can be represented by a binary tree.

- Each external node stores a character
- The code word of a character is given by the path from the root to the external node storing the character (0 for a left child and 1 for a right child)

00	010	011	10	11
a	b	c	d	e



Prefix Code

• Important

In variable length code, some method must be used to determine where the bits for each character start and end.

For Example: If e : 0, a : 1, t : 01

Then 0101 could correspond to eat, tea or eaea.

A **prefix code** is a binary code such that no code word is the prefix of another code-word

The Huffman Coding algorithm- History

- In 1951, David Huffman and his MIT information theory classmates given the choice of a term paper or a final exam
- Huffman hit upon the idea of using a frequency-sorted binary tree and quickly proved this method the most efficient.
- In doing so, the student outdid his professor, who had worked with information theory inventor Claude Shannon to develop a similar code.
- Huffman built the tree from the bottom up instead of from the top down



Huffman Code

- Greedy algorithm that constructs an optimal prefix code.
- Algorithm builds the tree in a bottom up manner.
- C be set of n alphabets and $f[c]$ be the frequency of $c \in C$.
- Algorithm begins with a set of $|C|$ leaves and performs a sequence of $|C| - 1$ merging operations to create the final tree.
- A priority queue Q keyed on f values, is used to identify the two least frequent objects to merge together.
- The result of the merger of two objects is a new object whose frequency is the sum of two frequencies of the two objects that are merged. The new object is inserted back into Q .

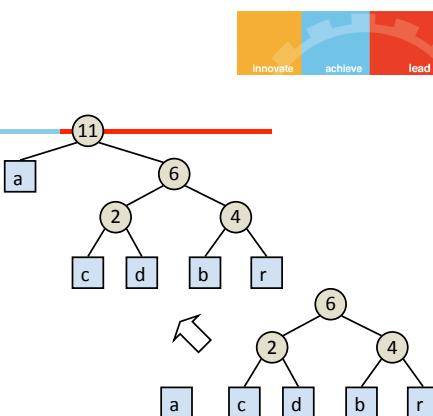
BITS Pilani, Pilani Campus

Example

$X = \text{abracadabra}$

Frequencies

a	b	c	d	r
5	2	1	1	2



a	b	c	d	r
5	2	1	1	2

a	b	c	d	r
5	2	1	1	2

BITS Pilani, Pilani Campus



Huffman's Algorithm

- It runs in time $O(n + d \log d)$, where n is the size of X and d is the number of distinct characters of X
- A heap-based priority queue is used as an auxiliary structure

Algorithm *HuffmanEncoding(X)*

```

Input string  $X$  of size  $n$ 
Output optimal encoding trie for  $X$ 
 $C \leftarrow \text{distinctCharacters}(X)$ 
 $\text{computeFrequencies}(C, X)$ 
 $Q \leftarrow \text{new empty heap}$ 
for all  $c \in C$ 
     $T \leftarrow \text{new single-node tree storing } c$ 
     $Q.\text{insert}(\text{getFrequency}(c), T)$ 
while  $Q.\text{size}() > 1$ 
     $f_1 \leftarrow Q.\text{minKey}()$ 
     $T_1 \leftarrow Q.\text{removeMin}()$ 
     $f_2 \leftarrow Q.\text{minKey}()$ 
     $T_2 \leftarrow Q.\text{removeMin}()$ 
     $T \leftarrow \text{join}(T_1, T_2)$ 
     $Q.\text{insert}(f_1 + f_2, T)$ 
return  $Q.\text{removeMin}()$ 

```

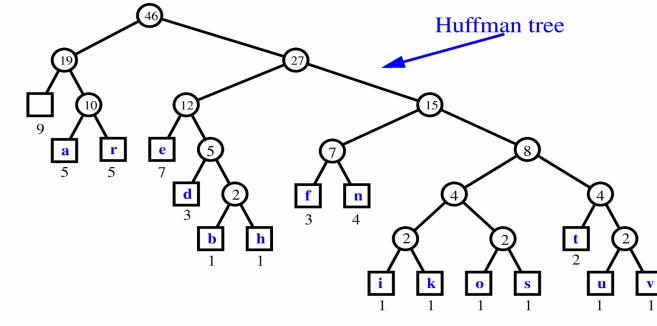
BITS Pilani, Pilani Campus



Huffman Tree Example

String: a fast runner need never be afraid of the dark

Character	a	b	d	e	f	h	i	k	n	o	r	s	t	u	v
Frequency	9	5	1	3	7	3	1	1	4	1	5	1	2	1	1



BITS Pilani, Pilani Campus

PH	
R3	Algorithm Design, Jon Kleinberg, Eva Tardos, First Ed., Pearson

Content Structure

1. Analyzing Algorithms [3 Hours]
 - 1.1. Theoretical Foundation
 - 1.1.1. Algorithms and its Specification
 - 1.1.2. Random Access Machine Model
 - 1.1.3. Counting Primitive Operations
 - 1.1.4. Notion of best case, average case and worst case
 - 1.2. Characterizing Run Time
 - 1.2.1. Use of asymptotic notation
 - 1.2.2. Big-Oh Notation, Little-Oh, Omega and Theta Notations
 - 1.3. Correctness of Algorithms
 - 1.4. Analyzing Recursive Algorithms
 - 1.4.1. Recurrence relations
 - 1.4.2. Specifying runtime of recursive algorithms
 - 1.4.3. Solving recurrence equations
 - 1.5. Case Study: Analyzing Algorithms
2. Elementary Data Structures [2 hours]
 - 2.1.
 - 2.2. Stacks
 - 2.2.1. Stack ADT and Implementation
 - 2.2.2. Applications
 - 2.3. Queues

Topics Covered



BITS Pilani, Pilani Campus

BIRLA INSTITUTE OF TECHNOLOGY & SCIENCE, PILANI WORK INTEGRATED LEARNING PROGRAMMES	
Digital Learning	
Part A: Course Design	
Course Title	Data Structures and Algorithms Design
Course No(s)	SS26519
Credit Units	
Credit Model	
Content Authors	BHARAT M DESHPANDE
Course Objectives	
CO1	Introduce mathematical and experimental techniques to analyze algorithms
CO2	Introduce linear and non-linear data structures and best practices to choose appropriate data structure for a given application
CO3	Exposes students to various sorting and searching techniques
CO4	Overview of algorithm design methods such as the greedy method, divide and conquer, dynamic programming, backtracking, and branch and bound
CO5	Introduce complexity classes and ways of classifying problem.
Text Books(s)	
T1	Algorithms: Design Foundations, Analysis and Internet Examples Michael T. Goodrich, Roberto Tamassia, 2006, Wiley (Students Edition)
R1	Data Structures, Algorithms and Applications in C++, Sanjay Suri, Second Ed, 2005, Universities Press
R2	Introduction to Algorithms, TH Cormen, CE Leiserson, RL Rivest, CL Stein, Third Ed, 2009,
Reference Book(s) & other resources	

4. Dictionaries [as Hash Tables and Search Trees] [3 hours]
 - 4.1. Unordered Dictionary
 - 4.1.1. ADT Specification

6.2.4.2.	Vertex Cover	4.1.2. Applications
6.2.4.3.	Clique and Set-Cover	4.2. Hash Tables
6.2.4.4.	Subset-Sum and Knapsack	4.2.1. Notion of Hashing and Collision (with a simple vector based hash table)
6.2.4.5.	Hamiltonian Cycle and TSP	4.2.2. Hash Functions
	Learning Outcomes:	
No	Learning Outcomes	
LO1	Analyzing algorithms using recurrence relations and expressing it using asymptotic notation.	4.2.2.1. Properties
LO2	Understanding different sorting and searching algorithms with the analysis of best, worst and average cases.	4.2.2.2. Simple hash functions
LO4	Understanding of graph algorithms for Reachability and Path Finding	4.2.2.3. Methods for Collision Handling
LO3	Students will be able to solve problems using appropriate data structures to implement solutions.	4.2.3.1. Separate Chaining
LO5	Ability to classify problems into complexity classes P and NP and to prove hardness of problems	4.2.3.2. Notion of Load Factor
		4.2.3.3. Rehashing
		4.2.3.4. Open Addressing [Linear & Quadratic Probing, Double Hash]
		4.3. Ordered Dictionary
		4.3.1. ADI Specification
		4.3.2. Applications
		4.4. Binary Search Tree
		4.4.1. Motivation with the task of Searching and Binary Search Algorithm
		4.4.2. Properties of BST
		4.4.3. Searching an element in BST
		4.4.4. Insertion and Removal of Elements
		4.4.5. Performance
		5. Algorithm Design Techniques [6 Hours]
		5.1. Greedy Method
		5.1.1. Design Principles and Strategy
		5.1.2. Fractional Knapsack Problem
		5.1.3. Task Scheduling Problem
		5.2. Divide and Conquer
		5.2.1. Design Principles and Strategy
		5.2.2. Analyzing Divide-and-Conquer Algorithms
		5.2.3. Integer Multiplication Problem
		5.2.4. Sorting Problem
		5.2.4.1. Merge Sort Algorithm
		5.2.4.2. Quick Sort Algorithm
		5.2.5. Searching Problem and Binary Search Algorithm [Ref to 4.4.1]
		5.3. Dynamic Programming
		5.3.1. Design Principles and Strategy
		5.3.2. Matrix Chain Product Problem
		5.3.3. 0/1 Knapsack Problem
		5.4. Graph Algorithms
		5.4.1. Introduction to Graphs
		5.4.2. Prim's and Kruskal's Algorithms [Greedy]
		5.4.3. Dijkstra's Algorithm [Greedy]
		5.4.4. Bellman-ford Shortest Path Algorithm [Greedy]
		5.4.5. All Pair Shortest Path Algorithm [Dynamic Programming]
		6. Complexity Classes [4 hours]
		6.1. Definition of P and NP classes and examples
		6.2. Understanding NP-Completeness
		6.2.1. NP-Hardness
		6.2.2. Polynomial time reducibility
		6.2.3. Cook-Levin theorem
		6.2.4. Problems in NP-Complete and using polynomial time reductions
		6.2.4.1. CNF-SAT, 3-SAT

Academic Term	Second Semester 2015-2016
Course Title	Data Structures and Algorithms Design
Course No	SS20519
Content Developer	BHARAT M DESHPANDE

Part B: Content Development Plan	
Academic Term	Second Semester 2015-2016

- Glossary of Terms:
1. Contact Hour (CH) stands for a hour long live session with students conducted either in a physical classroom or enabled through technology. In this model of instruction, instructor led sessions will be for 20 CH.
 - a. Pre CH = Self Learning done prior to a given contact hour
 - b. During CH = Content to be discussed during the contact hour by the course instructor
 - c. Post CH = Self Learning done post the contact hour
 2. RL stands for Recorded Lecture or Recorded Lesson. It is presented to the student through an online portal. A given RL unfolds as a sequence of video segments interleaved with exercises
 3. SS stands for Self-Study to be done as a study of relevant sections from textbooks and reference books. It could also include a study of external resources.
 4. LE stands for Lab Exercises
 5. HW stands for Home Work will consists could be a selection of problems from the text.

Post CH	LE4
Post CH	QZ4

Notes:

Contact Hour 5

Time	Type	Sequence	Content Reference
Pre CH	RL1.2		
During CH	CH5	CH5.1 = List: Notion of position in lists CH5.2 = List ADT and Implementation	T1: 2.2
Post CH	SS5		
Post CH	HW5	T1 - R2.1	
Post CH	LE5		
Post CH	QZ5		

Notes:

Contact Hour 6

Time	Type	Sequence	Content Reference
Pre CH	RL1.3		
During CH	CH6	CH6.1 = Trees: Terms and Definition, CH6.2= Tree ADT, Applications CH6.3= Binary Trees : Terms and Definition, CH6.4= Properties	T1: 2.3.1, 2.3.2
Post CH	SS6		
Post CH	HW6		

Notes:

Contact Hour 2

Time	Type	Sequence	Content Reference
Pre CH	RL1.1		
During CH		CH1.1 = Introduction and it's Specification, CH1.2 = Random Access Machine Model, CH1.3 = Counting Primitive Operations, CH1.4= Notion of best case, average case and worst case	T1: 1.1.1-1.1.3
Post CH	SS1		
Post CH	HW1		
Post CH	LE1		
Post CH	QZ1		

Notes:

Contact Hour 1

Time	Type	Sequence	Content Reference
Pre CH	RL1.1		
During CH		CH1.1 = Algorithms and it's Specification, CH1.2 = Random Access Machine Model, CH1.3 = Counting Primitive Operations, CH1.4= Notion of best case, average case and worst case	T1: 1.1.1-1.1.3
Post CH	SS1		
Post CH	HW1		
Post CH	LE1		
Post CH	QZ1		

Notes:

Contact Hour 7

Time	Type	Sequence	Content Reference
Pre CH	RL1.3		
During CH	CH7	CH7.1 = Binary Trees: Representations (Vector Based and Linked), CH7.2= Binary Tree traversal (In Order, Pre Order, Post Order), CH7.3= Applications	T1: 2.3.3
Post CH	SS7		
Post CH	HW7		
Post CH	LE7		
Post CH	QZ7		

Notes:

Contact Hour 3

Time	Type	Sequence	Content Reference
Pre CH	RL1.1		
During CH		CH2.1 = Use of asymptotic notation (Big-Oh, Little-Oh, Omega and Theta Notations) CH2.2 = Correctness of Algorithms	T1: 1.2
Post CH	SS2		
Post CH	HW2	T1 - R1.15, R1.19	
Post CH	LE2		

Notes:

Contact Hour 4

Time	Type	Sequence	Content Reference
Pre CH	RL1.2		
During CH		CH4.1=Stacks: ADT and Implementation CH4.2=Queues: Queue ADT and Implementation, CH4.3= Insertion and deletion of elements	T1: 2.1
Post CH	SS4		
Post CH	HW4		

Notes:

Contact Hour 8

		Searching and Binary Search Algorithm, CH12.1 = Properties of Insert, CH12.4 = Searching an element in Insert, insertion and Removal of Elements	
Post CH	SS12		
Post CH	HW12		
Post CH	LE12		
Post CH	QZ12		

Notes:

Contact Hour 13

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH13	CH13.1 =CH12.1=Greedy Method; Design Principles and Strategy, CH12.2 =Fractional Knapsack Problem	T1: 5.1.1
Post CH	SS13		
Post CH	HW13		
Post CH	LE13		
Post CH	QZ13		

Notes:

Contact Hour 14

Time	Type	Sequence	Content Reference
Pre CH			

Notes:

Contact Hour 9

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH9		CH19.1 =Heaps: Heap implementation of priority queue, CH19.2 = Heap sort
Post CH	SS9		T1 - Section 2.4.2 PQ-Sort
Post CH	HW9		
Post CH	LE9		
Post CH	QZ9		

Notes:

Contact Hour 10

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH10		CH10.1 =Unordered Dictionary ADT, Applications CH10.2 = Hash Tables: Notion of Hashing and Collision (with a simple vector based hash table) CH10.3 = Hash Functions: Properties, Simple hash functions
Post CH	SS10		T1: 2.5.1, 2.5.2, 2.5.3, 2.5.4
Post CH	HW10		
Post CH	LE10		
Post CH	QZ10		

Notes:

Contact Hour 11

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH11		CH11.1 =Methods for Collision Handling, CH11.2 = Notion of Load Factor, CH11.3 = Rehashing, CH11.4 = Open Addressing
Post CH	SS11		
Post CH	HW11		
Post CH	LE11		
Post CH	QZ11		

Notes:

Contact Hour 15

Time	Type	Sequence	Content Reference
Pre CH	RL2.1		
During CH	CH15	CH15.1 =Divide and Conquer: Design Principles and Strategy, CH15.2 = Analyzing Divide and Conquer Algorithms, CH15.3 = Integer Multiplication Problem	T1:5.2, 1, 5.2.2
Post CH	SS15		
Post CH	HW15		
Post CH	LE15		
Post CH	QZ15		

Notes:

Contact Hour 16

Time	Type	Sequence	Content Reference
Pre CH	RL2.1, RL2.2		
During CH	CH16	CH16.1 =Merge Sort CH16.2 =Quick Sort	T1:4.1, 4.3
Post CH	SS16		
Post CH	HW16		
Post CH	LE16		
Post CH	QZ16		

Notes:

Contact Hour 12

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH12	CH12.1 =Ordered Dictionary ADT, Applications CH12.2 = Binary Search tree: Motivation with the task of	T1: 3.1, 3.1.5
Post CH	SS12		
Post CH	HW12		
Post CH	LE12		
Post CH	QZ12		

Post CH	SS20		
Post CH	IW20		
Post CH	LE20		
Post CH	QZ20		

Notes:

Contact Hour 21

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH21	CH21.1 = Clique and Set-Cover, CH21.2 = Subset-Sum and Knapsack	T1:13.3
Post CH	SS21	Traveling Salesman Problem	
Post CH	IW21		
Post CH	LE21		
Post CH	QZ21		

Notes:

Contact Hour 22

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH22	CH22.1 = Course Review	
Post CH	SS22		
Post CH	IW22		
Post CH	LE22		

Notes:

Contact Hour 17

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH21	CH21.1 = Clique and Set-Cover, CH21.2 = Subset-Sum and Knapsack	T1:13.3
Post CH	SS21	Traveling Salesman Problem	
Post CH	IW21		
Post CH	LE21		
Post CH	QZ21		

Notes:

Contact Hour 18

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH17	CH17.1 = Dynamic Programming: Design Principles and Strategy, CH17.2 = Matrix Chain Product Problem CH17.3 = 0-1 Knapsack Problem	T1:5.3.1, 5.3.2
Post CH	SS17		
Post CH	IW17		
Post CH	LE17		
Post CH	QZ17		

Notes:

Contact Hour 19

Time	Type	Sequence	Content Reference
Post CH	SS18		T1:6.4
Post CH	IW18		
Post CH	LE18		
Post CH	QZ18		

Notes:

Contact Hour 19

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH18	CH18.1 = Graphs: Terms and Definitions, Properties, CH18.2 = Representations (Edge list, Adjacency list, Adjacency Matrix) CH18.3 = Graph Traversals	T1:6.1, 6.2, 6.3
Post CH	SS18		
Post CH	IW18		
Post CH	LE18		
Post CH	QZ18		

Notes:

Contact Hour 19

Time	Type	Sequence	Content Reference
Post CH	SS18		T1:6.4
Post CH	IW18		
Post CH	LE18		
Post CH	QZ18		

Notes:

Contact Hour 19

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH19	CH19.1 = Single Source Shortest Path algorithm: Dijkstra's Algorithm CH19.2 = Definition of P and NP classes and examples	T1:7.1.1
Post CH	SS19		T1:7.1.2, 7.2.1, 7.3.1, 7.3.2
Post CH	IW19		
Post CH	LE19		
Post CH	QZ10		

Notes:

Contact Hour 20

Time	Type	Sequence	Content Reference
Post CH	SS18		T1:6.4
Post CH	IW18		
Post CH	LE18		
Post CH	QZ18		

Notes:

Contact Hour 20

Time	Type	Sequence	Content Reference
Pre CH			
During CH	CH20	CH20.1 = Understanding NP-Completeness: CNF-SAT, Cook-Levin theorem, CH20.2 = Polynomial time reducibility: CNF-SAT and 3-SAT, Vertex Cover	T1:13.2, T13.3
Post CH	SS18		
Post CH	IW18		
Post CH	LE18		
Post CH	QZ18		

Notes:

Contact Hour 20